

# ECE 1508: Reinforcement Learning

## Chapter 2: Model-based RL

Ali Bereyhi

`ali.bereyhi@utoronto.ca`

Department of Electrical and Computer Engineering  
University of Toronto

Summer 2024

# Classical RL Methods: Recall

Ultimate goal in an RL problem is to find the *optimal policy*

As mentioned, we have two *major challenges* in this way

- 1 We need to compute *values* explicitly
- 2 We often deal with settings with *huge state spaces*?

In this part of the course, we are going to handle the first challenge

- This chapter  $\rightsquigarrow$  Model-based methods
- Next chapter  $\rightsquigarrow$  Model-free methods

# A Good Start Point: *Model-based RL*

In a nutshell, in **model-based** methods

*we are able to describe mathematically the behavior of environment*

This might come from the **nature of problem** or simply *postulated by us*

## Model-Based RL

Bellman Equation

value iteration

policy iteration

## Model-free RL

on-policy methods

temporal difference

Monte Carlo

SARSA

off-policy methods

Q-learning

## Complete State is *Markov Process*

When we formulated the RL framework, we stated that

*a complete state must describe a Markov process*

### Markov Process

Sequence  $S_1 \rightarrow S_2 \rightarrow \dots$  describe a Markov process if

$$\Pr \{S_{t+1} = s_{t+1} | S_t = s_t, \dots, S_1 = s_1\} = \Pr \{S_{t+1} = s_{t+1} | S_t = s_t\}$$

Following this fact, we introduced the concepts of

*rewarding and transition functions*

## Recall: *Transition and Rewarding*

Both these mappings **only** depend on **current state** and **action**

**Transition function** maps **state**  $S_t$  and **action**  $A_t$  to the next state  $S_{t+1}$

$$\mathcal{P}(\cdot) : \mathcal{S} \times \mathcal{A} \mapsto \mathcal{S}$$

**Rewarding function** maps **state**  $S_t$  and **action**  $A_t$  to reward  $R_{t+1}$

$$\mathcal{R}(\cdot) : \mathcal{S} \times \mathcal{A} \mapsto \{r^1, \dots, r^L\}$$

We said that *these mappings are in general random*

# Describing Markov Trajectory

Markovity of the state indicates that we observe the following trajectory

$$S_0, A_0 \rightarrow (R_1, S_1), A_1 \rightarrow \dots \rightarrow (R_t, S_t), A_t \rightarrow (R_{t+1}, S_{t+1})$$

This trajectory describes a **Markov process** with conditional distribution

$$\begin{aligned} p(r, \bar{s} | s, a) &= \Pr \{ R_{t+1} = r, S_{t+1} = \bar{s} | S_t = s, A_t = a \} \\ &= \Pr \{ R_t = r, S_t = \bar{s} | S_{t-1} = s, A_{t-1} = a \} \\ &\vdots \\ &= \Pr \{ R_1 = r, S_1 = \bar{s} | S_0 = s, A_0 = a \} \end{aligned}$$

*The above trajectory describes a **Markov Decision Process (MDP)***

# Finite MDPs

In this course, we focus on *finite* MDPs

## Finite MDP

*The Markov process*

$$S_0, A_0 \rightarrow (R_1, S_1), A_1 \rightarrow \dots \rightarrow (R_t, S_t), A_t$$

is a finite MDP if *rewards*, *actions* and *states* belong to a finite set, i.e.,

$$r \in \{r^1, \dots, r^L\} \quad a \in \{a^1, \dots, a^M\} \quad s \in \{s^1, \dots, s^N\}$$

MDPs are completely described by conditional distribution  $p(r, \bar{s} | s, a)$

We call  $p(r, \bar{s} | s, a)$  hereafter *rewarding-transition model*

# Model-based RL via MDP

- + What makes it now *model-based RL*?
- We assume that *rewarding-transition model*  $p(r, \bar{s} | s, a)$  is given to us
- + But you said for model-based RL, we should know the *transition and rewarding functions*!
- Well, we can describe them using  $p(r, \bar{s} | s, a)$ !

## Rewarding Model

Assume we are in state  $S_t = s$  and act  $A_t = a$ ; then,  $R_{t+1}$  is a *random variable* whose distribution is given by

$$p(r | s, a) = \sum_{n=1}^N p(r, s^n | s, a)$$

We call this distribution hereafter *rewarding model*



## Model-based RL via MDP

Similarly, we can describe the *transition function*

### Transition Model

Assume we are in state  $S_t = s$  and act  $A_t = a$ ; then, next state  $S_{t+1}$  is a *random variable* whose distribution is given by

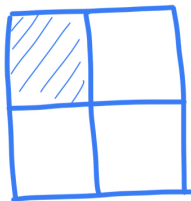
$$p(\bar{s}|s, a) = \sum_{\ell=1}^L p(r^\ell, \bar{s}|s, a)$$

We call this distribution hereafter *transition model*

## Example: Dummy Grid World

We have a grid board where at each cell we can move

$$\mathcal{A} = \{0 \equiv \text{left}, 1 \equiv \text{down}, 2 \equiv \text{right}, 3 \equiv \text{up}\}$$

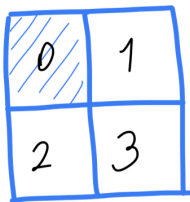


Our ultimate goal is to arrive at **top-left corner**  
through **shortest** path

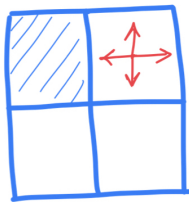
This problem describes an MDP with **deterministic rewarding-transition model**

- **State** is the number of cell
- **Action** is the direction we move
- **Reward** is  $-1$  each time we move until we get to destination
  - ↳ **Reward** is  $-0.5$  when we **hit the corners**

## Example: *Dummy Grid World*



$$\mathcal{S} = \{0, 1, 2, 3\}$$

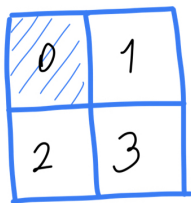


$$\mathcal{A} = \{0 \equiv \text{left}, 1 \equiv \text{down}, 2 \equiv \text{right}, 3 \equiv \text{up}\}$$

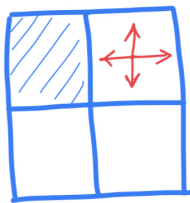
Let's write the *rewarding-transition model* down

$$p(r, \bar{s} | \mathbf{3}, \mathbf{3}) = \begin{cases} 1 & (r, \bar{s}) = (-1, 1) \\ 0 & (r, \bar{s}) \neq (-1, 1) \end{cases}$$

## Example: *Dummy Grid World*



$$\mathcal{S} = \{0, 1, 2, 3\}$$



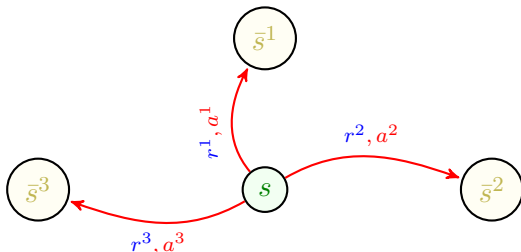
$$\mathcal{A} = \{0 \equiv \text{left}, 1 \equiv \text{down}, 2 \equiv \text{right}, 3 \equiv \text{up}\}$$

Let's write the *rewarding-transition model* down

$$p(r, \bar{s} | 0, a) = \begin{cases} 1 & (r, \bar{s}) = (0, 0) \\ 0 & (r, \bar{s}) \neq (0, 0) \end{cases} \rightsquigarrow s = 0 \text{ is terminal state}$$

# Transition Diagram

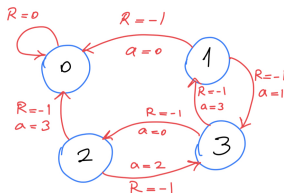
It is sometimes helpful to show **transition model** via a *transition diagram*



*This diagram describes a graph*

- Each node is a **possible state**: we have in total  $N$  nodes
- Node  $s$  is connected to  $\bar{s}$  if the probability of transition is **non-zero**
  - ↳ We could specify the **action** that **can** lead us to the **new state**
  - ↳ The graph could have **loops** including **self-loop**

# Transition Diagram: Dummy Grid World



In our dummy grid world, we have **four states**

- If we are in **terminal state** we always remain there with **no rewards**
- From **state**  $s = 1$  we can go to **states**  $\bar{s} = 0, 3$  depending on **action**
  - ↳ We can also remain in **state**  $s = 1$  and reward with  $-0.5$  if we hit corners
- From **state**  $s = 2$  we can go to **states**  $\bar{s} = 0, 3$  depending on **action**
  - ↳ We can also remain in **state**  $s = 1$  and reward with  $-0.5$  if we hit corners
- From **state**  $s = 3$  we can go to **states**  $\bar{s} = 1, 2$  depending on **action**
  - ↳ We can also remain in **state**  $s = 1$  and reward with  $-0.5$  if we hit corners

## Expected Action Reward

As we said, using *rewarding-transition model* we can describe the environment completely: for instance, let's see what would be the expected immediate reward that we get if in *state*  $s$  we act  $a$

$$\begin{aligned}
 \bar{\mathcal{R}}(s, a) &= \mathbb{E} \{ R_{t+1} | s, a \} \rightsquigarrow \text{we simplify notation } S_t = s \text{ to } s \\
 &= \sum_{\ell=1}^L r^{\ell} p \left( r^{\ell} | s, a \right) \\
 &= \sum_{\ell=1}^L r^{\ell} \sum_{n=1}^N p \left( r^{\ell}, s^n | s, a \right) \\
 &= \sum_{\ell=1}^L \sum_{n=1}^N r^{\ell} p \left( r^{\ell}, s^n | s, a \right) \rightsquigarrow \text{rewarding-transition model}
 \end{aligned}$$

## Expected Action Reward

$\bar{\mathcal{R}}(s, a)$  describes

*the reward we expect to see immediately after acting  $a$  in state  $s$*

We are going to see this expectation a lot, so maybe we could give it a name

### Expected Action Reward

*The expected reward for a state-action pair  $(s, a)$  is defined as*

$$\bar{\mathcal{R}}(s, a) = \mathbb{E} \{ R_{t+1} | s, a \} = \sum_{\ell=1}^L \sum_{n=1}^N r^{\ell} p \left( r^{\ell}, s^n | s, a \right)$$

Obviously,  $\bar{\mathcal{R}}(s, a)$  does **not** depend on **policy**



# Expected Policy Reward

- + Can we relate it also to *our policy*?
- Sure! We could *average over our policy*

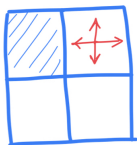
## Expected Policy Reward

The expected immediate reward of policy  $\pi$  at *state*  $s$  is defined as

$$\begin{aligned}\bar{\mathcal{R}}_{\pi}(s) &= \mathbb{E}_{\pi} \{R_{t+1} | s\} = \sum_{m=1}^M \mathbb{E} \{R_{t+1} | s, a^m\} \pi(a^m | s) \\ &= \sum_{m=1}^M \sum_{\ell=1}^L \sum_{n=1}^N r^{\ell} p(r^{\ell}, s^n | s, a) \pi(a^m | s)\end{aligned}$$

It describes reward we expect to see immediately after *state*  $s$  while playing  $\pi$

## Example: Dummy Grid World



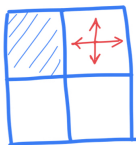
In our dummy grid world, we can easily compute the *expected immediate reward*

$$\bar{\mathcal{R}}(1, a) = \begin{cases} -1 & a \in \{0, 1\} \\ -0.5 & a \in \{2, 3\} \end{cases}$$

Obviously in *terminal state* we always get *zero expected reward*, e.g., for *all*  $a$

$$\bar{\mathcal{R}}(0, a) = 0$$

## Example: Dummy Grid World



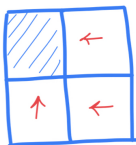
Now assume that we play *uniformly at random*, i.e., for all  $a$  and  $s$

$$\pi(a|s) = \frac{1}{4}$$

In this case the *expected policy reward* is

$$\bar{\mathcal{R}}_{\pi}(1) = \sum_{a=0}^3 \bar{\mathcal{R}}(1, a) \pi(a|1) = -0.75$$

## Example: Dummy Grid World



But if we change to *above deterministic* policy: the *expected reward* changes to

$$\bar{\mathcal{R}}_{\pi}(1) = \sum_{a=0}^3 \bar{\mathcal{R}}(1, a) \pi(a|1) = \bar{\mathcal{R}}(1, 0) = -1$$

and we can easily show that

$$\bar{\mathcal{R}}_{\pi}(0) = 0 \quad \bar{\mathcal{R}}_{\pi}(2) = -1 \quad \bar{\mathcal{R}}_{\pi}(3) = -1$$

## Computing Value Functions: *Naive Approach*

Now that we have a **concrete model** for our **environment**: we should go ahead and compute the **value function**, as we want to **optimize** it

*Let's start with direct computation*

$$\begin{aligned}v_{\pi}(s) &= \mathbb{E}_{\pi} \{G_t | s\} \\&= \mathbb{E}_{\pi} \{R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | s\} \\&= \mathbb{E}_{\pi} \{R_{t+1} | s\} + \gamma \mathbb{E}_{\pi} \{R_{t+2} | s\} + \gamma^2 \mathbb{E}_{\pi} \{R_{t+3} | s\} + \dots \\&= \bar{\mathcal{R}}_{\pi}(s) + \gamma \mathbb{E}_{\pi} \{R_{t+2} | s\} + \gamma^2 \mathbb{E}_{\pi} \{R_{t+3} | s\} + \dots\end{aligned}$$

- + How can we compute next terms?
- We could use the **rewarding-transition model of MDP**

# Computing Value Functions: *Naive Approach*

*Let's try the second term for example: we first define the notation*

$$\mathbb{E}_{\pi} \{R_{t+2} | s, s^n, a^m, a^j\} = \mathbb{E}_{\pi} \{R_{t+2} | S_t = s, S_{t+1} = s^n, A_t = a^m, A_{t+1} = a^j\}$$

*We can easily compute  $\mathbb{E}_{\pi} \{R_{t+2} | s, s^n, a^m, a^j\}$  as*

$$\begin{aligned} \mathbb{E}_{\pi} \{R_{t+2} | s, s^n, a^m, a^j\} &= \sum_{\ell=1}^L r^{\ell} p \left( r^{\ell} | s, s^n, a^m, a^j \right) \\ &= \sum_{\ell=1}^L r^{\ell} p \left( r^{\ell} | s^n, a^j \right) \end{aligned}$$

# Computing Value Functions: Naive Approach

We can then say that

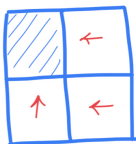
$$\mathbb{E}_{\pi} \{R_{t+2}|s\} = \sum_{n=1}^N \sum_{m=1}^M \sum_{j=1}^M \mathbb{E}_{\pi} \{R_{t+2}|s, s^n, a^m, a^j\} p(a^m, s^n, a^j|s)$$

and write down  $p(a^m, s^n, a^j|s)$  using chain rule

$$\begin{aligned} p(a^m, s^n, a^j|s) &= p(a^m|s) p(s^n|s, a^m) p(a^j|s, a^m, s^n) \\ &= \pi(a^m|s) \underbrace{p(s^n|s, a^m)}_{\text{transition model}} \pi(a^j|s^n) \end{aligned}$$

- + How can we compute the next term?
- We should repeat the same approach: there will be *more nested sums*

## Example: Dummy Grid World



Let's start with the *above policy*:  $\pi^1$

$$v_{\pi^1}(1) = \mathbb{E}_{\pi^1} \{ R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | 1 \}$$

Following policy at  $s = 1$  we end up at *terminal state* at next time

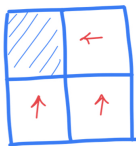
$$v_{\pi^1}(1) = \mathbb{E}_{\pi^1} \{ R_{t+1} + \gamma 0 + \gamma^2 0 + \dots | 1 \} = \bar{\mathcal{R}}_{\pi^1}(1) = -1$$

Same way, we can conclude that

$$v_{\pi^1}(0) = 0 \quad v_{\pi^1}(2) = -1 \quad v_{\pi^1}(3) = -2$$



## Example: *Dummy Grid World*

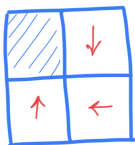


Let's change to *policy*  $\pi^2$ : we could follow same steps to show that

$$v_{\pi^2}(0) = 0 \quad v_{\pi^2}(1) = -1 \quad v_{\pi^2}(2) = -1 \quad v_{\pi^2}(3) = -2$$

We note that it returns the same values as *policy*  $\pi^1$

## Example: *Dummy Grid World*



Let's now look at *policy*  $\pi^3$ : we could follow same steps to show that

$$v_{\pi^3}(0) = 0 \quad v_{\pi^3}(1) = -3 \quad v_{\pi^3}(2) = -1 \quad v_{\pi^3}(3) = -2$$

We can see that

$$\pi^1 = \pi^2 \geq \pi^3$$

## Computing Value Functions: *Practical Approach*

- + *But, we should compute **infinite** terms **in general**!*
- *Well, if we are lucky: the sequence either **terminates** or **shows a pattern***
- + *What if that doesn't happen?*
- *Then, this approach really does **not** work!*

---

*This is why we called it the **naive approach**, since we **never** use this approach: in practice,*

*we always invoke a **Bellman equation***

*and find it via **dynamic programming***

## Future Return: *Recursive Property*

Even though **future return** looks **infinte**, it has a simple recursive property

$$\begin{aligned} G_t &= \sum_{i=0}^{\infty} \gamma^i R_{t+i+1} \\ &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \\ &= R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \dots) \\ &= R_{t+1} + \gamma G_{t+1} \end{aligned}$$

---

We can use this property to find a *fixed-point equation* for the value function!

## Value Function: Recursive Property

Say we are playing with policy  $\pi$ : we can write the value function as

$$\begin{aligned}
 v_{\pi}(s) &= \mathbb{E}_{\pi} \{G_t | s\} \\
 &= \mathbb{E}_{\pi} \{R_{t+1} + \gamma G_{t+1} | s\} \\
 &= \mathbb{E}_{\pi} \{R_{t+1} | s\} + \gamma \mathbb{E}_{\pi} \{G_{t+1} | s\} \\
 &= \bar{\mathcal{R}}_{\pi}(s) + \gamma \underbrace{\mathbb{E}_{\pi} \{G_{t+1} | s\}}_{?}
 \end{aligned}$$

- + Isn't that term again the value function at  $s$ ?
- Be careful! It's **not**

### Attention

The second term is not the value of **state**  $s$

$$\mathbb{E}_{\pi} \{G_{t+1} | s\} = \mathbb{E}_{\pi} \{G_{t+1} | S_t = s\} \neq \mathbb{E}_{\pi} \{G_t | S_t = s\} = v_{\pi}(s)$$

## Value Function: *Recursive Property*

Let's do some marginalization

$$\begin{aligned}\mathbb{E}_{\pi} \{G_{t+1}|s\} &= \sum_{n=1}^N \mathbb{E}_{\pi} \{G_{t+1}|S_t = s, S_{t+1} = s^n\} \Pr \{S_{t+1} = s^n|S_t = s\} \\ &= \sum_{n=1}^N \mathbb{E}_{\pi} \{G_{t+1}|s, s^n\} p(s^n|s)\end{aligned}$$

*Well, we need to specify the two terms in under summation, i.e.,*

- $\mathbb{E}_{\pi} \{G_{t+1}|s, s^n\}$
- $p(s^n|s) = \Pr \{S_{t+1} = s^n|S_t = s\}$

## Value Function: Recursive Property

Recall the trajectory

$$S_0, A_0 \rightarrow (R_1, S_1), A_1 \rightarrow \dots \rightarrow (R_{t+1}, S_{t+1}), A_{t+1} \rightarrow (R_{t+2}, S_{t+2})$$

If we know *state*  $S_{t+1}$  any reward after  $t + 1$  *only* depends on  $S_{t+1}$ , i.e.,

$$\mathbb{E}_{\pi} \{G_{t+1} | S_t = s, S_{t+1} = s^n\} = \mathbb{E}_{\pi} \{G_{t+1} | S_{t+1} = s^n\}$$

This indicates that

$$\mathbb{E}_{\pi} \{G_{t+1} | s, s^n\} = v_{\pi}(s^n)$$

I.e., the *value function* at state  $s^n$

## Value Function: *Recursive Property*

We can further find  $p(s^n|s)$  from **transition model** and **policy**

$$\begin{aligned} p_{\pi}(s^n|s) &= \sum_{m=1}^M p(s^n, a^m|s) \\ &= \sum_{m=1}^M p(a^m|s) p(s^n|a^m, s) \\ &= \sum_{m=1}^M \pi(a^m|s) p(s^n|s, a^m) \rightsquigarrow \text{depends on policy} \end{aligned}$$

We know have both terms in terms of **transition model** and **policy**



## Value Function: *Recursive Property*

Replacing into the equation, where we left we have

$$\begin{aligned}\mathbb{E}_{\pi} \{G_{t+1} | s\} &= \sum_{n=1}^N \mathbb{E}_{\pi} \{G_{t+1} | s, s^n\} p(s^n | s) \\ &= \sum_{n=1}^N v_{\pi}(s^n) p_{\pi}(s^n | s) \\ &= \sum_{n=1}^N \sum_{m=1}^M v_{\pi}(s^n) p(s^n | s, a^m) \pi(a^m | s)\end{aligned}$$

We can also present it by shorter notation as

$$\mathbb{E}_{\pi} \{G_{t+1} | s\} = \mathbb{E}_{\pi} \{v_{\pi}(S_{t+1}) | s\}$$

## Value Function: *Recursive Property*

Back to computation of **value function**, we have

$$\begin{aligned}v_{\pi}(s) &= \bar{\mathcal{R}}_{\pi}(s) + \gamma \mathbb{E}_{\pi} \{G_{t+1} | s\} \\&= \bar{\mathcal{R}}_{\pi}(s) + \gamma \mathbb{E}_{\pi} \{v_{\pi}(S_{t+1}) | s\} \\&= \bar{\mathcal{R}}_{\pi}(s) + \gamma \sum_{n=1}^N v_{\pi}(s^n) p_{\pi}(s^n | s)\end{aligned}$$

This is a **recursive equation** that relates value of one state to other values  
*which is a **Bellman equation***

# Bellman Equation: Value

## Bellman Equation for Value Function

For any policy  $\pi$  the value function at each **state**  $s$  satisfies

$$v_{\pi}(s) = \bar{\mathcal{R}}_{\pi}(s) + \gamma \sum_{n=1}^N v_{\pi}(s^n) p_{\pi}(s^n | s)$$

- + Well! What is the **use** of **Bellman equation**?
- It describes a **fixed-point** equation that can be solved for  $v_{\pi}(s)$ !

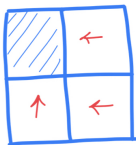
# Bellman Equation: *Breaking Down*

$$v_{\pi}(s) = \bar{\mathcal{R}}_{\pi}(s) + \gamma \sum_{n=1}^N v_{\pi}(s^n) p_{\pi}(s^n|s)$$

In general, we have  $N$  possible state  $\rightsquigarrow$  we have  $N$  possible values

- Bellman equation relates each value to other  $N - 1$  values
  - ↳ For each  $s$ , Bellman equation has  $N$  unknowns  $v_{\pi}(s^1), \dots, v_{\pi}(s^N)$
- We can write the Bellman equation for all  $N$  states
  - ↳ We have  $N$  equations each with  $N$  unknowns
- We solve this system of equations for unknowns  $v_{\pi}(s^1), \dots, v_{\pi}(s^N)$

## Example: Dummy Grid World



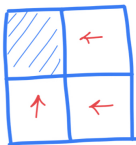
Let's try with our dummy grid world: we saw that

$$\bar{\mathcal{R}}_{\pi}(0) = 0 \quad \bar{\mathcal{R}}_{\pi}(1) = -1 \quad \bar{\mathcal{R}}_{\pi}(2) = -1 \quad \bar{\mathcal{R}}_{\pi}(3) = -1$$

Now let's consider the values **unknown**

$$v_{\pi}(0), v_{\pi}(1), v_{\pi}(2), v_{\pi}(3)$$

## Example: *Dummy Grid World*



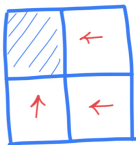
We set  $\gamma = 1$  and start with state  $s = 0$

$$v_{\pi}(0) = \bar{\mathcal{R}}_{\pi}(0) + \sum_{\bar{s}=0}^3 v_{\pi}(\bar{s}) p_{\pi}(\bar{s}|0)$$

We know that

$$p_{\pi}(\bar{s}|0) = \begin{cases} 1 & \bar{s} = 0 \\ 0 & \bar{s} \neq 0 \end{cases}$$

## Example: *Dummy Grid World*

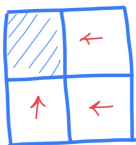


This concludes that at state  $s = 0$ , *Bellman equation reads*

$$v_{\pi}(0) = 0 + v_{\pi}(0)$$

*which is an obvious equation; let's try  $s = 1$*

## Example: *Dummy Grid World*



At state  $s = 0$ , we have

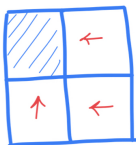
$$v_{\pi}(1) = \bar{\mathcal{R}}_{\pi}(1) + \sum_{\bar{s}=0}^3 v_{\pi}(\bar{s}) p_{\pi}(\bar{s}|1)$$

Again we can easily say based on the *policy* that

$$p_{\pi}(\bar{s}|1) = \begin{cases} 1 & \bar{s} = 0 \\ 0 & \bar{s} \neq 0 \end{cases}$$



## Example: Dummy Grid World



This concludes that at state  $s = 1$ , Bellman equation reads

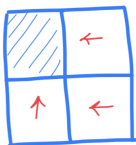
$$v_{\pi}(1) = -1 + v_{\pi}(0)$$

which relates  $v_{\pi}(1)$  to  $v_{\pi}(0)$ . If we keep repeating we get further

$$v_{\pi}(2) = -1 + v_{\pi}(0)$$

$$v_{\pi}(3) = -1 + v_{\pi}(2)$$

## Example: *Dummy Grid World*



We now have the system of equations

$$v_{\pi}(1) = -1 + v_{\pi}(0)$$

$$v_{\pi}(2) = -1 + v_{\pi}(0)$$

$$v_{\pi}(3) = -1 + v_{\pi}(2)$$

We also know that  $s = 0$  is a **terminal state**, and thus  $v_{\pi}(0) = 0$ : so, we get

$$v_{\pi}(1) = -1 \quad v_{\pi}(2) = -1 \quad v_{\pi}(3) = -2$$

## Bellman Equation: Action-Value

We can find a **Bellman equation** for Action-value function as well: say we play with policy  $\pi$

$$\begin{aligned}
 q_{\pi}(s, a) &= \mathbb{E}_{\pi} \{G_t | s, a\} \\
 &= \mathbb{E}_{\pi} \{R_{t+1} + \gamma G_{t+1} | s, a\} \\
 &= \mathbb{E} \{R_{t+1} | s, a\} + \gamma \mathbb{E}_{\pi} \{G_{t+1} | s, a\} \\
 &= \bar{\mathcal{R}}(s, a) + \gamma \underbrace{\mathbb{E}_{\pi} \{G_{t+1} | s, a\}}_{?}
 \end{aligned}$$

We need to compute

$$\mathbb{E}_{\pi} \{G_{t+1} | s, a\}$$

in terms of the **rewarding-transition model** and **policy**

## Action-Value: Recursive Property

We apply the marginalization trick

$$\mathbb{E}_{\pi} \{G_{t+1} | s, a\} = \sum_{n=1}^N \mathbb{E}_{\pi} \{G_{t+1} | S_t = s, S_{t+1} = s^n, A_t = a\} p(s^n | s, a)$$

### Attention

Recalling the trajectory of the MDP, we should note that

$$q_{\pi}(s^n, a) \neq \mathbb{E}_{\pi} \{G_{t+1} | S_t = s, S_{t+1} = s^n, A_t = a\} = v_{\pi}(s^n)$$

In fact, once we know  $S_{t+1}$ , the **previous action** does not contain any extra information! We only gain information, if we observe  $A_{t+1}$ , i.e.,

$$\mathbb{E}_{\pi} \{G_{t+1} | S_t = s, S_{t+1} = s^n, A_{t+1} = a\} = q_{\pi}(s^n, a)$$

## Action-Value: *Recursive Property*

So, we can replace it into original equation to get

$$\begin{aligned}\mathbb{E}_{\pi} \{G_{t+1} | s, a\} &= \sum_{n=1}^N \mathbb{E}_{\pi} \{G_{t+1} | S_t = s, S_{t+1} = s^n, A_t = a\} p(s^n | s, a) \\ &= \sum_{n=1}^N v_{\pi}(s^n) p(s^n | s, a)\end{aligned}$$

This implies that

$$\begin{aligned}q_{\pi}(s, a) &= \bar{\mathcal{R}}(s, a) + \gamma \sum_{n=1}^N v_{\pi}(s^n) p(s^n | s, a) \\ &= \bar{\mathcal{R}}(s, a) + \gamma \mathbb{E} \{v_{\pi}(S_{t+1}) | s, a\}\end{aligned}$$

# Bellman Equation: Action-Value

## Bellman Equation I for Action-Value Function

*For any policy  $\pi$  the action-value function at each pair  $(s, a)$  satisfies*

$$q_{\pi}(s, a) = \bar{\mathcal{R}}(s, a) + \gamma \sum_{n=1}^N v_{\pi}(s^n) p(s^n | s, a)$$

*After doing Assignment 1, you will immediately conclude the following extension*

## Bellman Equation II for Action-Value Function

*For any policy  $\pi$  the action-value function at each pair  $(s, a)$  satisfies*

$$q_{\pi}(s, a) = \bar{\mathcal{R}}(s, a) + \gamma \sum_{n=1}^N \sum_{m=1}^M q_{\pi}(s^n, a^m) \pi(a^m | s^n) p(s^n | s, a)$$

# Computing Action-Value via Bellman Equation

We can again use the recursive equation

$$q_{\pi}(s, a) = \bar{\mathcal{R}}(s, a) + \gamma \sum_{n=1}^N \sum_{m=1}^M q_{\pi}(s^n, a^m) \pi(a^m | s^n) p(s^n | s, a)$$

to find the **action-value function**: we have in this case  $NM$  possible values

- **Bellman equation** relates **each action-value** to other **action-values**  
 ↳ For **each  $s$  and  $a$** , Bellman equation has  $NM$  unknowns  $q_{\pi}(s^n, a^m)$
- We can write the **Bellman equation** for all  $NM$  cases
- We solve this system of equations for **unknowns**  $q_{\pi}(s^n, a^m)$

# Bellman Equation: Backup Diagram

Bellman equation gives an

interesting *visualization* for *values* and *action-values*

which can be shown in the s-called *backup diagram*

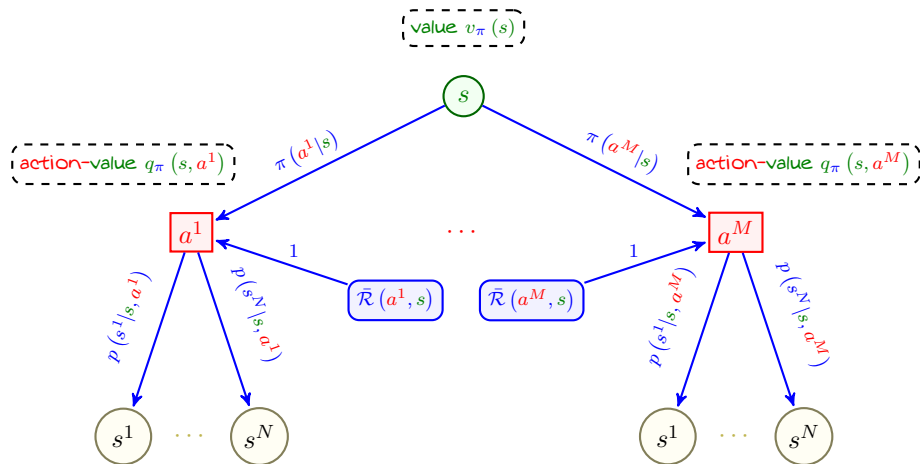
---

For simplicity, we consider  $\gamma = 1$  in the *backup diagram*

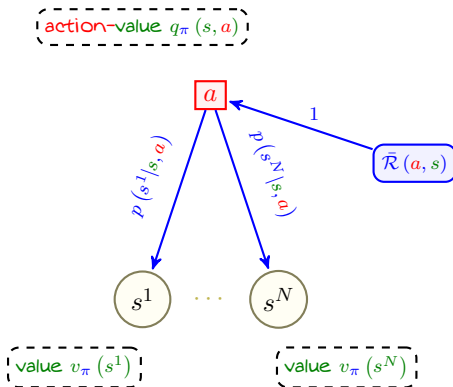
- Each circle node is a *state* and carries the *value of the state*
- Each square node is an *action* and carries the *action-value of the pair*
- Each *edge* is a transition and carries *a probability*
- As we pass from *leaves* to *root*
  - Value of each node multiplies to its probability on the edge
  - They add up when they meet at a parent node
    - ↳ This makes the value of the parent node



# Backup Diagram: For Given Policy



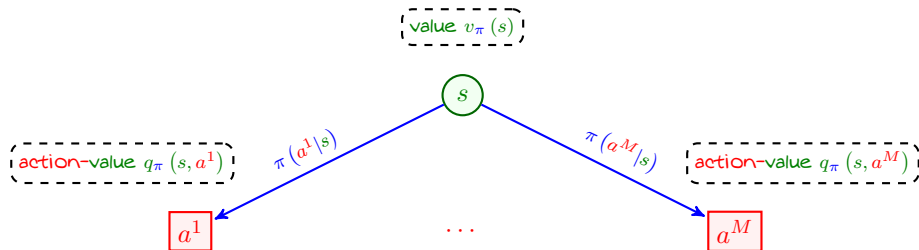
# Backup Diagram: For Given Policy



Let's look at it part by part: *first we pass from leaves to **action** parent*

$$q_\pi(s, a) = \bar{\mathcal{R}}(s, a) + \sum_{n=1}^N v_\pi(s^n) p(s^n|s, a)$$

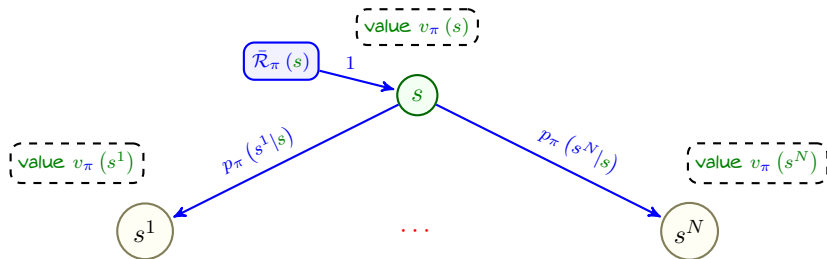
# Backup Diagram: For Given Policy



Then, we pass from **action** parents to the **root state**

$$v_\pi(s) = \sum_{m=1}^M \pi(a^m|s) q_\pi(s, a^m)$$

# Backup Diagram: For Given Policy



We could also have its alternative form **expected** over **actions**

$$v_{\pi}(s) = \bar{\mathcal{R}}_{\pi}(s) + \sum_{n=1}^N p_{\pi}(s^n|s) v_{\pi}(s^n)$$

# Finding Optimal Values

- + *Well! Bellman lets us compute **value** of a **given** policy. But, how can we find the optimal value? It doesn't seem to solve this problem!*
- We can in fact use it to directly find the **optimal values**!
- + *That sounds a bit **weird**!*
- Once we know the **optimality constraint**, it doesn't anymore

## Optimal Value: Optimality Constraint

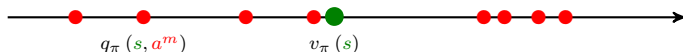
In Assignment 1, you show that for *any* **state** we have

$$v_{\pi}(s) = \sum_{m=1}^M q_{\pi}(s, a^m) \pi(a^m | s)$$

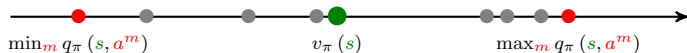
Now, recall that *policy* is a conditional **distribution** meaning that

$$0 \leq \pi(a^m | s) \leq 1$$

We can think of it as



# Optimal Value: Optimality Constraint



*It is hence obvious that*

$$\min_m q_\pi(s, a^m) \leq v_\pi(s) \leq \max_m q_\pi(s, a^m)$$

*We can use this simple fact to find a constraint on **optimal values***

## Optimal Value: Optimality Constraint

If our policy is the *optimal policy*; then, we should have

$$v_{\star}(s) = \text{maximum possible value} = \max_m q_{\star}(s, a^m)$$

- + But, can we guarantee that we can achieve such value?
- Sure! We can set an *optimal policy* to

$$\pi^{\star}(a^m | s) = \begin{cases} 1 & m = \operatorname{argmax}_m q_{\star}(s, a^m) \\ 0 & m \neq \operatorname{argmax}_m q_{\star}(s, a^m) \end{cases}$$

- + But, they are both in terms of  $q_{\star}(s, a^m)$ ! We don't have the *optimal action-values*!
- Sure! But, we could say that *optimal values must* satisfy this constraint: *if not, they cannot be optimal*



# Optimal Value: Optimality Constraint

## Optimality Constraint

Optimal value at each state  $s$  satisfies the following identity

$$v_{\star}(s) = \max_m q_{\star}(s, a^m)$$

and is achieved if we set the policy to

$$\pi^{\star}(a^m | s) = \begin{cases} 1 & m = \operatorname{argmax}_m q_{\star}(s, a^m) \\ 0 & m \neq \operatorname{argmax}_m q_{\star}(s, a^m) \end{cases}$$

which is an **optimal** policy

- + But, how can we relate this constraint to **Bellman equation**?
- Let's see!

## Optimal Value: *Bellman Equation*

We know from Bellman equation II for **action-value** function that

$$q_{\pi}(s, a) = \bar{\mathcal{R}}(s, a) + \gamma \sum_{n=1}^N v_{\pi}(s^n) p(s^n | s, a)$$

If we play with **optimal policy**: we are going to have *same identity*

$$q_{\star}(s, a) = \bar{\mathcal{R}}(s, a) + \gamma \sum_{n=1}^N v_{\star}(s^n) p(s^n | s, a)$$

We now substitute it in *optimality constraint*

$$v_{\star}(s) = \max_m \bar{\mathcal{R}}(s, a^m) + \gamma \sum_{n=1}^N v_{\star}(s^n) p(s^n | s, a^m)$$

## Optimal Value: Bellman Equation

This is again a *recursive equation* that

does **not** depend on any policy!

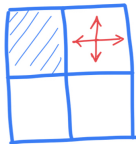
### Bellman Optimality Equation

The optimal value function  $v_{\star}(s)$  satisfies

$$v_{\star}(s) = \max_m \bar{\mathcal{R}}(s, a^m) + \gamma \sum_{n=1}^N v_{\star}(s^n) p(s^n | s, a^m)$$

We can again treat it as a fixed-point equation and solve it for  $v_{\star}(s)$

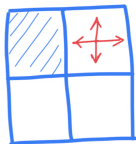
## Example: Dummy Grid World



Let's find optimal values for our dummy grid world: we first find  $\bar{\mathcal{R}}(s, a)$

$$\begin{array}{llll}
 \bar{\mathcal{R}}(0, a) = 0 & \bar{\mathcal{R}}(1, 0) = -1 & \bar{\mathcal{R}}(2, 0) = -0.5 & \bar{\mathcal{R}}(3, 0) = -1 \\
 \bar{\mathcal{R}}(1, 1) = -1 & \bar{\mathcal{R}}(2, 1) = -0.5 & \bar{\mathcal{R}}(3, 1) = -0.5 & \\
 \bar{\mathcal{R}}(1, 2) = -0.5 & \bar{\mathcal{R}}(2, 2) = -1 & \bar{\mathcal{R}}(3, 2) = -0.5 & \\
 \bar{\mathcal{R}}(1, 3) = -0.5 & \bar{\mathcal{R}}(2, 3) = -1 & \bar{\mathcal{R}}(3, 3) = -1 & 
 \end{array}$$

## Example: Dummy Grid World

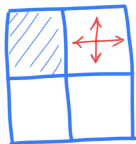


We next write down Bellman equations

- ① Since  $s = 0$  is a **terminal state** we know that  $v_{\star}(0) = 0$
- ② Now, let's consider  $s = 1$

$$\begin{aligned}
 p(0|1, \textcolor{red}{0}) &= 1 \\
 p(1|1, \textcolor{red}{0}) &= 0 \\
 p(2|1, \textcolor{red}{0}) &= 0 \\
 p(3|1, \textcolor{red}{0}) &= 0
 \end{aligned}
 \rightsquigarrow \sum_{\bar{s}=0}^4 v_{\star}(\bar{s}) p(\bar{s}|\textcolor{green}{1}, \textcolor{red}{0}) = v_{\star}(0) = 0$$

## Example: Dummy Grid World

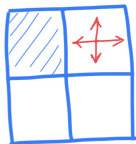


We next write down Bellman equations

- ① Since  $s = 0$  is a *terminal state* we know that  $v_{\star}(0) = 0$
- ② Now, let's consider  $s = 1$

$$\begin{aligned}
 p(0|1, \textcolor{red}{1}) &= 0 \\
 p(1|1, \textcolor{red}{1}) &= 0 \\
 p(2|1, \textcolor{red}{1}) &= 0 \\
 p(3|1, \textcolor{red}{1}) &= 1
 \end{aligned}
 \rightsquigarrow \sum_{\bar{s}=0}^4 v_{\star}(\bar{s}) p(\bar{s}|1, \textcolor{red}{1}) = v_{\star}(\textcolor{blue}{3})$$

## Example: Dummy Grid World

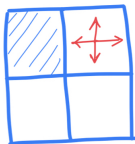


We next write down Bellman equations

- ① Since  $s = 0$  is a *terminal state* we know that  $v_{\star}(0) = 0$
- ② Now, let's consider  $s = 1$

$$\begin{aligned}
 p(0|1, \textcolor{red}{2}) &= 0 \\
 p(\textcolor{blue}{1}|1, \textcolor{red}{2}) &= 1 \\
 p(\textcolor{blue}{2}|1, \textcolor{red}{2}) &= 0 \\
 p(\textcolor{blue}{3}|1, \textcolor{red}{2}) &= 0
 \end{aligned}
 \rightsquigarrow \sum_{\bar{s}=0}^4 v_{\star}(\bar{s}) p(\bar{s}|\textcolor{blue}{1}, \textcolor{red}{2}) = v_{\star}(\textcolor{blue}{1})$$

## Example: Dummy Grid World



We next write down Bellman equations

- ① Since  $s = 0$  is a *terminal state* we know that  $v_{\star}(0) = 0$
- ② Now, let's consider  $s = 1$

$$\begin{aligned}
 p(0|1, \textcolor{red}{3}) &= 0 \\
 p(\textcolor{blue}{1}|1, \textcolor{red}{3}) &= 1 \\
 p(2|1, \textcolor{red}{3}) &= 0 \\
 p(\textcolor{blue}{3}|1, \textcolor{red}{3}) &= 0
 \end{aligned}
 \rightsquigarrow \sum_{\bar{s}=0}^4 v_{\star}(\bar{s}) p(\bar{s}|\textcolor{blue}{1}, \textcolor{red}{3}) = v_{\star}(\textcolor{blue}{1})$$



## Example: Dummy Grid World



We next write down Bellman equations

- ① Since  $s = 0$  is a **terminal state** we know that  $v_{\star}(0) = 0$
- ② Now, let's consider  $s = 1$

$$\begin{aligned}
 v_{\star}(1) &= \max_m \bar{\mathcal{R}}(1, a^m) + \sum_{\bar{s}=0}^4 v_{\star}(\bar{s}) p(\bar{s}|1, a^m) \\
 &= \max \{-1, -1 + v_{\star}(3), -0.5 + v_{\star}(1), -0.5 + v_{\star}(1)\}
 \end{aligned}$$

## Example: Dummy Grid World



We next write down Bellman equations

- 1 Since  $s = 0$  is a **terminal state** we know that  $v_{\star}(0) = 0$
- 2 Now, let's consider  $s = 1$

$$v_{\star}(1) = \max \{-1, -1 + v_{\star}(3), -0.5 + v_{\star}(1), -0.5 + v_{\star}(1)\}$$

- 3 Similarly, we have for  $s = 2$

$$v_{\star}(2) = \max \{-0.5 + v_{\star}(2), -0.5 + v_{\star}(2), -1 + v_{\star}(3), -1\}$$

## Example: Dummy Grid World



We next write down Bellman equations

- ① Since  $s = 0$  is a **terminal state** we know that  $v_{\star}(0) = 0$
- ② Now, let's consider  $s = 1$

$$v_{\star}(1) = \max \{-1, -1 + v_{\star}(3), -0.5 + v_{\star}(1), -0.5 + v_{\star}(1)\}$$

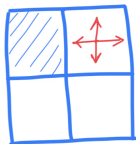
- ③ Similarly, we have for  $s = 2$

$$v_{\star}(2) = \max \{-0.5 + v_{\star}(2), -0.5 + v_{\star}(2), -1 + v_{\star}(3), -1\}$$

- ④ Finally for  $s = 3$ , we have

$$v_{\star}(3) = \max \{-1 + v_{\star}(2), -0.5 + v_{\star}(3), -0.5 + v_{\star}(3), -1 + v_{\star}(1)\}$$

## Example: *Dummy Grid World*



*After sorting out the Bellman equations, we get*

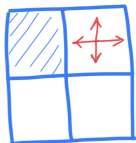
$$v_{\star}(1) = \max \{-1, -1 + v_{\star}(3), -0.5 + v_{\star}(1)\}$$

$$v_{\star}(2) = \max \{-1, -1 + v_{\star}(3), -0.5 + v_{\star}(2)\}$$

$$v_{\star}(3) = \max \{-1 + v_{\star}(2), -0.5 + v_{\star}(3), -1 + v_{\star}(1)\}$$

*We should now solve this system of equations*

## Example: *Dummy Grid World*



We first note that

$$\max \{-1, -1 + v_{\star}(3), -0.5 + v_{\star}(1)\} \neq -0.5 + v_{\star}(1)$$

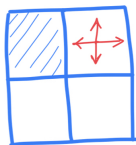
**Proof:** Assume that

$$\max \{-1, -1 + v_{\star}(3), -0.5 + v_{\star}(1)\} = -0.5 + v_{\star}(1)$$

Then, we have

$$v_{\star}(1) - 0.5 + v_{\star}(1) \rightsquigarrow 0 = -0.5 \quad \text{impossible!}$$

## Example: *Dummy Grid World*



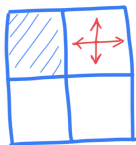
*For the same reason, we have*

$$\begin{aligned}\max \{-1, -1 + v_{\star}(3), -0.5 + v_{\star}(2)\} &\neq -0.5 + v_{\star}(2) \\ \max \{-1 + v_{\star}(2), -0.5 + v_{\star}(3), -1 + v_{\star}(1)\} &\neq -0.5 + v_{\star}(3)\end{aligned}$$

*So the equations reduce to*

$$\begin{aligned}v_{\star}(1) &= \max \{-1, -1 + v_{\star}(3)\} = v_{\star}(2) \\ v_{\star}(2) &= \max \{-1, -1 + v_{\star}(3)\} = v_{\star}(1) \\ v_{\star}(3) &= \max \{-1 + v_{\star}(2), -1 + v_{\star}(1)\} = -1 + v_{\star}(1)\end{aligned}$$

## Example: *Dummy Grid World*



Thus, we should only solve

$$v_{\star}(1) = \max \{-1, -1 + v_{\star}(3)\}$$

$$v_{\star}(3) = -1 + v_{\star}(1)$$

It is *again easy to see* that  $\max \{-1, -1 + v_{\star}(3)\} \neq -1 + v_{\star}(3)$ ; therefore,

$$v_{\star}(1) = v_{\star}(2) = -1 \rightsquigarrow v_{\star}(3) = -2$$

Well! This is what we expected!

# From Optimal Values to *Optimal Policy*

- + *What is the benefit then? It only finds **optimal value**, but we are looking for **optimal policy**!*
- We can actually back-track **optimal policy**, once we have **optimal value**

*The idea is quite simple:*

- 1 We can find optimal values from **Bellman optimality equations**
- 2 We could then find the **optimal action**-values
- 3 We finally get the **optimal policy** from **optimal action**-values



# Finding Optimal Policy: *Back-Tracking from Optimal Values*

We could summarize this approach algorithmically as follows

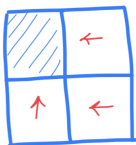
OptimBackTrack():

- 1: **for**  $n = 1 : N$  **do**
- 2:   Solve Bellman equation  $v_\star(s^n) = \max_m \bar{\mathcal{R}}(s^n, a^m) + \gamma \mathbb{E} \{v_\star(\bar{S}) | s^n, a^m\}$
- 3: **end for**
- 4: **for**  $n = 1 : N$  **do**
- 5:   **for**  $m = 1 : M$  **do**
- 6:     Compute action-value  $q_\star(s^n, a^m) = \bar{\mathcal{R}}(s^n, a^m) + \gamma \mathbb{E} \{v_\star(\bar{S}) | s^n, a^m\}$
- 7:   **end for**
- 8:   Compute optimal policy via optimality constraint

$$\pi^\star(a^m | s) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_\star(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_\star(s, a^m) \end{cases}$$

- 9: **end for**

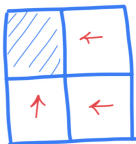
## Example: Dummy Grid World



Let's find optimal policy at **state**  $s = 1$  in our dummy grid world: first, we write

$$\begin{bmatrix} q_{\star}(1,0) \\ q_{\star}(1,1) \\ q_{\star}(1,2) \\ q_{\star}(1,3) \end{bmatrix} = \begin{bmatrix} \bar{\mathcal{R}}(1,0) + \sum_{\bar{s}} v_{\star}(\bar{s}) p(\bar{s}|1,0) \\ \bar{\mathcal{R}}(1,1) + \sum_{\bar{s}} v_{\star}(\bar{s}) p(\bar{s}|1,1) \\ \bar{\mathcal{R}}(1,2) + \sum_{\bar{s}} v_{\star}(\bar{s}) p(\bar{s}|1,2) \\ \bar{\mathcal{R}}(1,3) + \sum_{\bar{s}} v_{\star}(\bar{s}) p(\bar{s}|1,3) \end{bmatrix} = \begin{bmatrix} -1+0 \\ -1-2 \\ -0.5-1 \\ -0.5-1 \end{bmatrix} = \begin{bmatrix} -1 \\ -3 \\ -1.5 \\ -1.5 \end{bmatrix}$$

## Example: *Dummy Grid World*



The optimal policy at *state*  $s = 1$  is then given by

$$\pi^*(a|1) = \begin{cases} 1 & a = \operatorname{argmax}_a q_\star(1, a) \\ 0 & a \neq \operatorname{argmax}_a q_\star(1, a) \end{cases} = \begin{cases} 1 & a = 0 \\ 0 & a \neq 0 \end{cases}$$

Well! We know that this is *optimal* in this problem!

## Finding Optimal Policy: *Back-Tracking from Optimal Values*

- + Wait a moment! Does that mean that our optimal policy is always *deterministic*? But, you said it could be also *random*!
- Well! In some cases we could find *random optimal policies* as well!

---

If  $q_{\star}(s, a^m)$  has a single maximizer; then,

optimal policy  $\pi^*(a^m|s)$  is *deterministic*

But, if it has *multiple* maximizers

optimal policy  $\pi^*(a^m|s)$  can also be *random*

# Finding Optimal Policy: *General Form*

## Generic Optimal Policy

Assume that  $m^1, \dots, m^J$  are all maximizers of  $q_\star(s, a^m)$ ; then, policy

$$\pi^\star(a^m|s) = \begin{cases} p_1 & m = m^1 \\ \vdots & \\ p_J & m = m^J \\ 0 & m \notin \{m^1, \dots, m^J\} \end{cases}$$

for any  $p_1, \dots, p_J$  that satisfy

$$\sum_{j=1}^J p_j = 1$$

is an *optimal* policy

# Finding Optimal Policy

- + But, why are all such policies *optimal*?
- Well! We could look back at the *optimality constraint*

With any policy  $\pi^\star(a|s)$  of the form given in the last slide, we have

$$\begin{aligned} v_{\pi^\star}(s) &= \sum_{m=1}^M \pi^\star(a^m|s) q_{\pi^\star}(s, a^m) = \sum_{j=1}^J p_j q_{\pi^\star}(s, a^{mj}) + 0 \\ &= \sum_{j=1}^J p_j \max_m q_{\pi^\star}(s, a^m) = \max_m q_{\pi^\star}(s, a^m) \sum_{j=1}^J p_j = \max_m q_{\pi^\star}(s, a^m) \end{aligned}$$

which is the *optimality constraint*! It's intuitive, because

If we have *multiple options* for *next action* that give us *same maximal value*; then, we could *randomly* pick any of them

# Finding Optimal Policy

- + But, still we could have a *deterministic optimal* policy in such cases! Right?!
- Sure! We could *always* have a *deterministic optimal* policy!

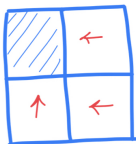
## Deterministic Optimal Policy

With known MDP for the environment, there exists *at least one deterministic optimal policy*

In the nutshell: if we know the *complete state* and its *transition model*

- We *always* can find a *deterministic optimal policy*
- We might have *multiple deterministic optimal policies*
  - ↳ In that case, we are going to have *also random optimal policies*

## Example: Dummy Grid World

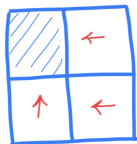


Let's find optimal policy at **state**  $s = 3$  in our dummy grid world: first, we write

$$\begin{bmatrix} q_{\star}(3,0) \\ q_{\star}(3,1) \\ q_{\star}(3,2) \\ q_{\star}(3,3) \end{bmatrix} = \begin{bmatrix} \bar{\mathcal{R}}(3,0) + \sum_{\bar{s}} v_{\star}(\bar{s}) p(\bar{s}|3,0) \\ \bar{\mathcal{R}}(3,1) + \sum_{\bar{s}} v_{\star}(\bar{s}) p(\bar{s}|3,1) \\ \bar{\mathcal{R}}(3,2) + \sum_{\bar{s}} v_{\star}(\bar{s}) p(\bar{s}|3,2) \\ \bar{\mathcal{R}}(3,3) + \sum_{\bar{s}} v_{\star}(\bar{s}) p(\bar{s}|3,3) \end{bmatrix} = \begin{bmatrix} -1 & -1 \\ -0.5 & -2 \\ -0.5 & -2 \\ -1 & -1 \end{bmatrix} = \begin{bmatrix} -2 \\ -2.5 \\ -2.5 \\ -2 \end{bmatrix}$$



## Example: *Dummy Grid World*

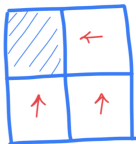


The optimal policy at *state*  $s = 3$  is then given by

$$\pi^*(a|3) = \begin{cases} 1 & a = \underset{a}{\operatorname{argmax}} q_\star(3, a) \\ 0 & a \neq \underset{a}{\operatorname{argmax}} q_\star(3, a) \end{cases} = \begin{cases} 1 & a = 0 \\ 0 & a \neq 0 \end{cases}$$

This is obviously *optimal* in this problem!

## Example: *Dummy Grid World*

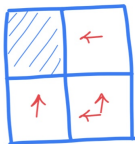


The optimal policy at *state*  $s = 3$  is then given by

$$\pi^*(a|3) = \begin{cases} 1 & a = \operatorname{argmax}_a q_*(3, a) \\ 0 & a \neq \operatorname{argmax}_a q_*(3, a) \end{cases} = \begin{cases} 1 & a = 3 \\ 0 & a \neq 3 \end{cases}$$

This is obviously *optimal* in this problem!

## Example: *Dummy Grid World*

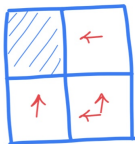


The optimal policy at *state*  $s = 3$  is then given by

$$\pi^*(a|3) = \begin{cases} 0.5 & a = 0 \\ 0 & a = 1, 2 \\ 0.5 & a = 3 \end{cases}$$

This is *also optimal* in this problem!

## Example: *Dummy Grid World*

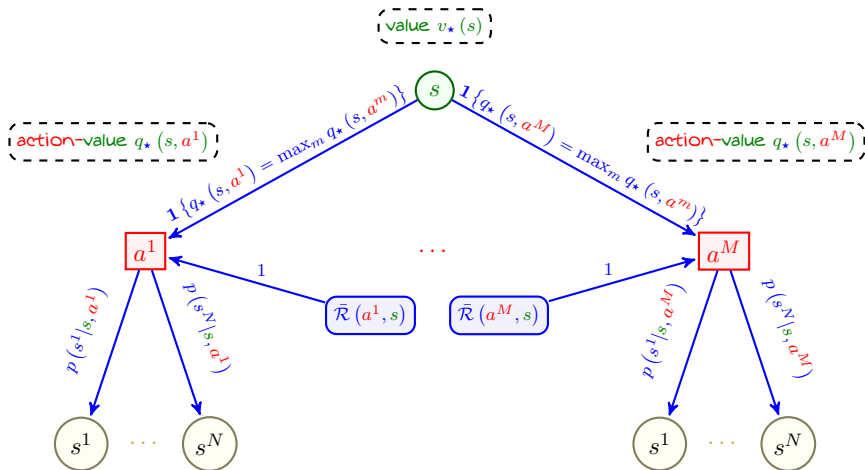


The optimal policy at *state*  $s = 3$  is then given by

$$\pi^*(a|3) = \begin{cases} 0.2 & a = 0 \\ 0 & a = 1, 2 \\ 0.8 & a = 3 \end{cases}$$

This is *also optimal* in this problem!

# Backup Diagram: For Optimal Policy



Here, we assume  $q_*(s, a^m)$  has one maximizer  $\equiv$  *optimal policy is deterministic*

## Last Piece: *Dynamic Programming*

Right now, we know what to do when we know *MDP of environment*

- ① We can find optimal values from *Bellman optimality equations*
- ② We could then find the *optimal action*-values
- ③ We finally get the *optimal policy* from *optimal action*-values

The only remaining challenge is to find

an algorithmic approach to solve *Bellman optimality equations*

We complete this last piece using

*Dynamic Programming*  $\equiv$  *DP*

# Dynamic Programming: *Basic Idea*

Assume, we want to solve the problem of

$$x = f(x)$$

for some function  $f(x)$

We could solve it via *direct approach*:

- ① Rewrite is as  $f(x) - x = 0$
- ② Solve it via classic algorithms
  - ↳ Reduce it to a *known form*, e.g., a *polynomial*
  - ↳ Solve it via an *iterative method*, e.g., *Newton-Raphson* or *method of intervals*

# Dynamic Programming: *Basic Idea*

Assume, we want to solve the problem of

$$x = f(x)$$

for some function  $f(x)$

We could also solve it by *recursion*:

- ① Start with an  $x^0$  and set  $x^1 = f(x^0)$
- ② Until  $x^{k+1} \approx x^k$ , we do
  - ↳ Update *recursively* as  $x^{k+1} = f(x^k)$
  - ↳ Set  $k \leftarrow k + 1$

Under *some conditions on  $f(\cdot)$* , this approach can *converge*



# Dynamic Programming: Example

We want to solve

$$x = \frac{-1}{2 + x}$$

- ① Start with an  $x^0 = 0$
- ② We now get into the recursion loop

$$\hookrightarrow x^1 = f(x^0) = -\frac{1}{2}$$

$$\hookrightarrow x^2 = f(x^1) = -\frac{2}{3}$$

$$\hookrightarrow x^3 = f(x^2) = -\frac{3}{4}$$

$$\hookrightarrow \dots$$

$$\hookrightarrow x^k = f(x^{k-1}) = -\frac{k}{k+1}$$

We asymptotically **converge** to  $x^\infty = -1$  which is the **solution**

$\hookrightarrow$  Note that we **always** converge **no matter** which point we start

# Dynamic Programming: Example

Now, let's write the *same equation* in a *different recursive form*

$$x = \frac{-1 - x^2}{2}$$

① Start with an  $x^0 = 0$

② We get into recursion loop

$$\hookrightarrow x^1 = f(x^0) = -0.5$$

$$\hookrightarrow x^2 = f(x^1) = -0.625$$

$$\hookrightarrow \dots$$

$$\hookrightarrow x^\infty = -1$$

① Start with an  $x^0 = 5$

② We get into recursion loop

$$\hookrightarrow x^1 = f(x^0) = -13$$

$$\hookrightarrow x^2 = f(x^1) = -85$$

$$\hookrightarrow \dots$$

$$\hookrightarrow x^\infty = -\infty$$

We can now diverge if we start with a wrong initial point!

*Not all recursive forms are always converging!*

# Dynamic Programming: *Applications to Our Problem*

Our problem has a similar form: we need to solve *Bellman equations*  
which are *recursive equations*

So, we could use **DP** to find the solution

---

There are two major **DP** approaches

- Policy Iteration that uses recursion to iterate between
  - ↳ Policy *Evaluation*
  - ↳ Policy *Improvement*
- Value Iteration which applies recursion on *optimal Bellman* equation

Let's look at these two approaches in detail

# Policy Evaluation: Step 1

The first step is *policy evaluation*: we can formulate this problem as follows

## Ultimate Goal of Policy Evaluation

Given a *policy*  $\pi$ , we intend to *evaluate* values of *all states* by *recursion*

Before we start, let's recap a few definitions: recall *expected policy reward*

$$\bar{\mathcal{R}}_{\pi}(s) = \sum_{m=1}^M \bar{\mathcal{R}}(a^m, s) \pi(a^m | s)$$

For sake of compactness, we use the following notation

$$\bar{\mathcal{R}}_{\pi}(s) = \mathbb{E}_{\pi} \{ \bar{\mathcal{R}}(A, s) | s \}$$

# Policy Evaluation: Step 1

Similarly, we define the notation

$$\mathbb{E}_{\pi} \{v_{\pi}(\bar{S}) | s, a\} = \sum_{n=1}^N v_{\pi}(s^n) p(s^n | s, a)$$

and also denote its expected form over the **action set** by

$$\begin{aligned} \mathbb{E}_{\pi} \{v_{\pi}(\bar{S}) | s\} &= \sum_{n=1}^N v_{\pi}(s^n) p_{\pi}(s^n | s) \\ &= \sum_{m=1}^M \underbrace{\sum_{n=1}^N v_{\pi}(s^n) p(s^n | s, a^m)}_{\mathbb{E}_{\pi} \{v_{\pi}(\bar{S}) | s, a^m\}} \pi(a^m | s) \\ &= \sum_{m=1}^M \mathbb{E}_{\pi} \{v_{\pi}(\bar{S}) | s, a^m\} \pi(a^m | s) \end{aligned}$$

## Policy Evaluation: Step 1

We can then write the *Bellman equations* compactly as

$$v_{\pi}(s) = \bar{\mathcal{R}}_{\pi}(s) + \gamma \mathbb{E}_{\pi} \{ v_{\pi}(\bar{S}) | s \}$$

for *value function* and also as

$$q_{\pi}(s, a) = \bar{\mathcal{R}}(s, a) + \gamma \mathbb{E}_{\pi} \{ v_{\pi}(\bar{S}) | s, a \}$$

for *action-value function*

Now, we are ready to *evaluate a policy by recursion*

# Policy Evaluation: *Value Computation via Recursion*

Recall our perspective on value computation:

values are  *$N$  unknowns* that we want to compute from *Bellman equations*

Now, if someone claims that *the values*

$$v_{\pi}(s^n) = v_n$$

for  $n = 1 : N$  are values of policy  $\pi$ , can we confirm it?

- + Shouldn't we simply use *Bellman Equation*?!
  - Exactly!

# Policy Evaluation: Value Computation via Recursion

We could confirm

$$v_{\pi}(s^n) = v_n$$

by writing first finding for every state  $s$

$$\begin{aligned}
 \mathbb{E}_{\pi} \{ v_{\pi}(\bar{S}) | s \} &= \sum_{n=1}^N v_{\pi}(s^n) p_{\pi}(s^n | s) \\
 &= \sum_{n=1}^N \sum_{m=1}^M v_{\pi}(s^n) p(s^n | s, a^m) \pi(a^m | s) \\
 &= \sum_{n=1}^N \sum_{m=1}^M \underbrace{v_n}_{\text{claimed value}} \underbrace{p(s^n | s, a^m)}_{\text{transition model}} \underbrace{\pi(a^m | s)}_{\text{policy}}
 \end{aligned}$$



# Policy Evaluation: Value Computation via Recursion

We could confirm

$$v_{\pi}(s^n) = v_n$$

by writing first finding for every state  $s$

$$\mathbb{E}_{\pi} \{v_{\pi}(\bar{S}) | s\} = \text{computed from } v_n\text{'s} := F(\{v_1, \dots, v_N\}, s)$$

and then checking if

$$\begin{aligned} v_{\pi}(s^n) &= v_n = \bar{\mathcal{R}}_{\pi}(s^n) + \gamma \mathbb{E}_{\pi} \{v_{\pi}(\bar{S}) | s^n\} \\ &= \bar{\mathcal{R}}_{\pi}(s^n) + \gamma F(\{v_1, \dots, v_N\}, s) \end{aligned}$$

holds for all  $n = 1 : N$

## Policy Evaluation: Value Computation via Recursion

If it happens that the claimed  $v_\pi(\cdot)$  is **not** a **valid claim**; then, we get out of Bellman equation

$$\bar{v}_\pi(s^n) = \bar{v}_n = \bar{\mathcal{R}}_\pi(s^n) + \gamma \mathbb{E}_\pi \{v_\pi(\bar{S}) | s^n\}$$

which is different from the claimed  $v_\pi(\cdot)$ , i.e.,  $v_n \neq \bar{v}_n$

### Policy Evaluation

We iterate this procedure until we can confirm, i.e., we

- 1 set  $v_\pi(\cdot) \leftarrow \bar{v}_\pi(\cdot)$
- 2 repeat the same procedure and compute **new**  $\bar{v}_\pi(\cdot)$

We **stop** when  $v_\pi(\cdot) = \bar{v}_\pi(\cdot)$ , or at least it happens approximately

# Policy Evaluation

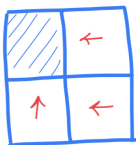
```

PolicyEval( $\pi, v_\pi^0$ ):
1: Initiate values with  $v_\pi^0$  and set  $k = 0$ 
2: Make sure that  $v_\pi^0(s) = 0$  for terminal states  $s$ 
3: Choose a small threshold  $\epsilon$  and initiate  $\Delta = +\infty$  # stopping criteria
4: for  $n = 1 : N$  do
5:   Compute  $\bar{\mathcal{R}}_\pi(s^n) = \mathbb{E}_\pi\{\bar{\mathcal{R}}(s^n, a)\}$  # average rewards
6: end for
7: while  $\Delta > \epsilon$  do
8:   for  $n = 1 : N$  do
9:     Update  $v_\pi^{k+1}(s^n) = \bar{\mathcal{R}}_\pi(s^n) + \gamma \mathbb{E}_\pi\{v_\pi^k(\bar{S}) | s^n\}$  # DP update
10:   end for
11:    $\Delta = \max_n |v_\pi^{k+1}(s^n) - v_\pi^k(s^n)|$  # check convergence
12:   Update  $k \leftarrow k + 1$ 
13: end while # Recursion Loop
  
```

## Attention

We should make sure that **terminal states** are all initiated with **zero** value

## Example: Dummy Grid World



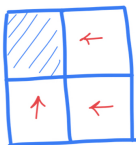
Let's try with our dummy grid world: we saw that

$$\bar{\mathcal{R}}_{\pi}(0) = 0 \quad \bar{\mathcal{R}}_{\pi}(1) = -1 \quad \bar{\mathcal{R}}_{\pi}(2) = -1 \quad \bar{\mathcal{R}}_{\pi}(3) = -1$$

Now let's *evaluate its values* by *recursion*: we first note that, if we have

$$\begin{aligned} \mathbb{E}_{\pi} \left\{ v_{\pi}^k(\bar{S}) \mid 0 \right\} &= v_{\pi}^k(0) & \mathbb{E}_{\pi} \left\{ v_{\pi}^k(\bar{S}) \mid 1 \right\} &= v_{\pi}^k(0) \\ \mathbb{E}_{\pi} \left\{ v_{\pi}^k(\bar{S}) \mid 2 \right\} &= v_{\pi}^k(0) & \mathbb{E}_{\pi} \left\{ v_{\pi}^k(\bar{S}) \mid 3 \right\} &= v_{\pi}^k(2) \end{aligned}$$

# Example: Dummy Grid World



PolicyEval( $\pi, v_\pi^0$ ):

- 1: Initiate values with  $v_\pi^0(1)$ ,  $v_\pi^0(2)$  and  $v_\pi^0(3)$  *at random* and *set*  $v_\pi^0(0) = 0$
- 2: Set  $\epsilon = 0.001$ , and initiate  $\Delta = 1000$  # stopping criteria
- 3: **while**  $\Delta > \epsilon$  **do**
- 4:   Update  $v_\pi^{k+1}(1) = -1 + v_\pi^k(0)$  # DP update
- 5:   Update  $v_\pi^{k+1}(2) = -1 + v_\pi^k(0)$  # DP update
- 6:   Update  $v_\pi^{k+1}(3) = -1 + v_\pi^k(2)$  # DP update
- 7:    $\Delta = \max_{s \in \{1,2,3\}} |v_\pi^{k+1}(s) - v_\pi^k(s)|$  # check convergence
- 8:   Update  $k \leftarrow k + 1$
- 9: **end while**

It converges after only *one* recursion!

# Policy Improvement

Let us now recall **optimality constraint**: with optimal policy, we have

$$v_{\star}(s) = \max_m q_{\star}(s, a^m)$$

which can be achieved by policy

$$\pi^{\star}(a^m | s) = \begin{cases} 1 & m = \operatorname{argmax}_m q_{\star}(s, a^m) \\ 0 & m \neq \operatorname{argmax}_m q_{\star}(s, a^m) \end{cases}$$

This means that if  $\pi$  is **not** optimal, we would have

$$\pi(a^m | s) \neq \begin{cases} 1 & m = \operatorname{argmax}_m q_{\pi}(s, a^m) \\ 0 & m \neq \operatorname{argmax}_m q_{\pi}(s, a^m) \end{cases}$$

# Policy Improvement

In other words, if we change our policy to

$$\bar{\pi}(a^m | s) = \begin{cases} 1 & m = \operatorname{argmax}_m q_{\pi}(s, a^m) \\ 0 & m \neq \operatorname{argmax}_m q_{\pi}(s, a^m) \end{cases}$$

Then, it should give us better values, i.e.,  $\bar{\pi} \geq \pi$ !

- + *Are you sure?! I don't see it immediately*
- We can actually show it!

This is what we call *policy improvement theorem*

# Policy Improvement

## Policy Improvement

Given (deterministic) policy  $\pi^k$ , we can always design a *better* policy  $\pi^{k+1}$  by setting it to

$$\pi^{k+1}(a^m | s) = \begin{cases} 1 & m = \operatorname{argmax}_m q_{\pi^k}(s, a^m) \\ 0 & m \neq \operatorname{argmax}_m q_{\pi^k}(s, a^m) \end{cases}$$



# Policy Improvement

PolicyImprov( $v_\pi$ ):

```

1: for  $n = 1 : N$  do
2:   for  $m = 1 : M$  do
3:     Compute  $\bar{R}(s^n, a^m)$ 
4:      $q_\pi(s^n, a^m) = \bar{R}(s^n, a^m) + \gamma \mathbb{E}_\pi \{v_\pi(\bar{S}) | s^n, a^m\}$    # action-values
5:   end for
6:   Compute an improved policy as                                     # policy improvement

```

$$\bar{\pi}(a^m | s^n) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_\pi(s^n, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_\pi(s^n, a^m) \end{cases}$$

```

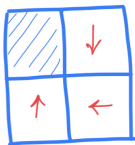
7: end for

```

## Attention

Here, we do **no recursion**

## Example: Dummy Grid World



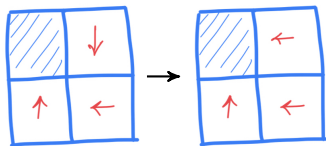
Let's try dummy grid world with above non-optimal policy: here, we have

$$v_{\pi}(0) = 0 \quad v_{\pi}(1) = -3 \quad v_{\pi}(2) = -1 \quad v_{\pi}(3) = -2$$

We now look at *action-values* at the *problematic state*  $s = 1$

$$\begin{aligned} q_{\pi}(1, 0) &= -1 \\ q_{\pi}(1, 1) &= -3 \\ q_{\pi}(1, 2) &= -3.5 \\ q_{\pi}(1, 3) &= -3.5 \end{aligned} \quad \rightsquigarrow \quad -3 = v_{\pi}(1) \neq \max_a q_{\pi}(1, a) = -1$$

## Example: *Dummy Grid World*



Now if we improve the policy, we get

$$\bar{\pi}(a|1) = \begin{cases} 1 & m = \operatorname{argmax}_m q_{\pi}(1, a) \\ 0 & m \neq \operatorname{argmax}_m q_{\pi}(1, a) \end{cases} = \begin{cases} 1 & a = 0 \\ 0 & a \neq 0 \end{cases}$$

which is actually *optimal*

# Policy Iteration: Improving Policy by Recursion

Looking at the **policy improvement** theorem, we see

If we plug in  $\pi^k = \pi^*$  into algorithm; then, after policy improvement

↳ we get  $\pi^{k+1} = \pi^*$

↳ say we **evaluate** values for  $\pi^{k+1} = \pi^*$  and plug back to algorithm

↳ we get  $\pi^{k+2} = \pi^*$

↳ say we **evaluate** values for  $\pi^{k+2} = \pi^*$  and plug back to algorithm

↳ ...

So, **optimal policy** is a **fixed-point** for this **recursion**

## Policy Iteration

We can start with an **arbitrary policy**  $\pi^0$  and keep doing the above recursion until we see that  $\pi^{k+1} = \pi^k$  which indicates that we reached **optimal policy**

# Policy Iteration

PolicyIter():

1: Initiate with **random**  $v_\pi(s)$  for all **non-terminal** states  $s$

2: Set  $v_\pi(s) = 0$  for **terminal** states  $s$

3: Initiate two random policies  $\pi$  and  $\bar{\pi}$

4: **while**  $\pi \neq \bar{\pi}$  **do**

5:  $v_\pi = \text{PolicyEval}(\pi, v_\pi)$  and  $\pi \leftarrow \bar{\pi}$  **Recursion**

6:  $\bar{\pi} = \text{PolicyImprov}(v_\pi)$

7: **end while**

Recursion

Note that this is a **nested** recursive computation

- There is a loop for recursion inside the algorithm in which
  - ↳ at each iteration we **evaluate the policy recursively**
- But, we initiate each **policy evaluation** loop with the values of last iteration
  - ↳ this can improve the convergence speed

# Back-Tracking by Recursion

- + *But wait a Moment! We already talked about back-tracking optimal policy from **Bellman optimality equation**! Don't we implement that?!*
- Sure! We can do the same thing by recursion

---

*We follow the same idea but we use recursion*

- 1 We can find optimal values from **Bellman optimality equations**
  - ↳ This is where we use **recursion**
- 2 We could then find the **optimal action**-values
- 3 We finally get the **optimal policy** from **optimal action**-values

## Recall: Back-Tracking from Optimal Values

OptimBackTrack():

```

1: Solve Bellman equations                                     # we use recursion
2: for  $n = 1 : N$  do
3:   for  $m = 1 : M$  do
4:     Set  $q_*(s^n, a^m) = \bar{\mathcal{R}}(s^n, a^m) + \gamma \mathbb{E} \{v_*(\bar{S}) | s^n, a^m\}$  # action-values
5:   end for
6:   Compute optimal policy via optimality constraint

```

$$\pi^*(a^m | s) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_*(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_*(s, a^m) \end{cases}$$

7: end for

# Recursion with Bellman Optimality

Recall Bellman optimality equation

$$v_{\star}(s) = \max_m \left( \bar{\mathcal{R}}(s, a^m) + \gamma \mathbb{E} \{ v_{\star}(\bar{S}) | s, a^m \} \right)$$

We can again solve it by recursion: we start with some  $v_{\star}^0(\cdot)$  and then for every state  $s$  and action  $a^m$ , we compute

$$\mathbb{E} \left\{ v_{\star}^k(\bar{S}) | s, a^m \right\} = \sum_{n=1}^N \underbrace{v_{\star}^k(s^n)}_{\text{last computed value}} \underbrace{p(s^n | s, a^m)}_{\text{transition model}}$$

We then update the optimal value function as

$$v_{\star}^{k+1}(s) = \max_m \left( \bar{\mathcal{R}}(s, a^m) + \gamma \mathbb{E} \left\{ v_{\star}^k(\bar{S}) | s, a^m \right\} \right)$$



## Value Iteration vs Policy Iteration

Before we complete the value iteration algorithm: *it is interesting to put its recursion next to the one used for policy evaluation*

*With optimality equation, we iterate as*

$$v_{\star}^{k+1}(s) = \max_m \left( \bar{\mathcal{R}}(s, a^m) + \gamma \mathbb{E} \left\{ v_{\star}^k(\bar{S}) \mid s, a^m \right\} \right)$$

*With Bellman equation for a given policy  $\pi$ , we iterate as*

$$\begin{aligned} v_{\pi}^{k+1}(s) &= \bar{\mathcal{R}}_{\pi}(s) + \gamma \mathbb{E}_{\pi} \left\{ v_{\pi}^k(\bar{S}) \mid s \right\} \\ &= \sum_{m=1}^M \left( \bar{\mathcal{R}}(s, a^m) + \gamma \mathbb{E} \left\{ v_{\pi}^k(\bar{S}) \mid s, a^m \right\} \right) \pi(a^m \mid s) \end{aligned}$$

# Value Iteration vs Policy Iteration

With optimality equation, we iterate as

$$v_{\star}^{k+1}(s) = \max_m \left( \bar{\mathcal{R}}(s, a^m) + \gamma \mathbb{E} \left\{ v_{\star}^k(\bar{S}) \mid s, a^m \right\} \right)$$

With Bellman equation for a given policy  $\pi$ , we iterate as

$$v_{\pi}^{k+1}(s) = \sum_{m=1}^M \left( \bar{\mathcal{R}}(s, a^m) + \gamma \mathbb{E} \left\{ v_{\pi}^k(\bar{S}) \mid s, a^m \right\} \right) \pi(a^m | s)$$

This indicates that for both recursive loops

- we compute  $M$  values *per iteration per state*
  - ↳ in policy iteration, we compute the *average* of these  $M$  via  $\pi$
  - ↳ in value iteration, we take the *largest* among these  $M$  values

# Value Iteration

ValueItr():

1: Initiate with **random**  $v_{\star}^0(s)$  for all states, and set  $v_{\star}^0(s) = 0$  for **terminal** states

2: Choose a **small** threshold  $\epsilon$ , initiate  $\Delta = +\infty$  and  $k = 0$

3: **while**  $\Delta > \epsilon$  **do**

4:   **for**  $n = 1 : N$  **do**

5:     **for**  $m = 1 : M$  **do**

6:       Compute  $q_{\star}(s^n, a^m) = \bar{\mathcal{R}}(s^n, a^m) + \gamma \mathbb{E} \{ v_{\star}^k(\bar{S}) | s^n, a^m \}$

7:     **end for**

8:     Update  $v_{\pi}^{k+1}(s^n) = \max_m q_{\star}(s^n, a^m)$  # DP update

9:   **end for**

10:   Set  $\Delta = \max_n |v_{\pi}^{k+1}(s^n) - v_{\pi}^k(s^n)|$  and  $k \leftarrow k + 1$

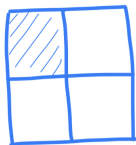
11: **end while**

Recursion

12: Compute an optimal policy as

$$\bar{\pi}(a^m | s) = \begin{cases} 1 & m = \operatorname{argmax}_m q_{\star}(s, a^m) \\ 0 & m \neq \operatorname{argmax}_m q_{\star}(s, a^m) \end{cases}$$

## Example: *Dummy Grid World*



*You may try policy and value iteration for this problem at home!*

*Easy as Pie* 😊

## Example: A Bit Larger Grid World<sup>1</sup>

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

Board  $\equiv$  states



moves  $\equiv$  actions

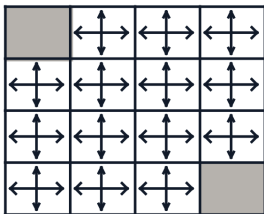
Let's do a bit of more serious example: we are now in a  $4 \times 4$  grid world

- We have two **terminal states** shown in gray
- Each move we do gets a -1 reward
  - ↳ We also get -1 reward if we hit a corner
  - ↳ We get zero reward at terminal state

In simple words: we are looking for **shortest path** to the **corners**

<sup>1</sup>This example is taken from Sutton and Barto's Book; Example 4.1 in Chapter 4

## Example: A Bit Larger Grid World



initial policy

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

initial values

Let's first try *policy iteration*: we start with

- a uniform random policy  $\pi^0$
- all values being zero, i.e.,  $v_{\pi^0}^0(s) = 0$  for all  $s$

## Example: A Bit Larger Grid World

Recall policy iteration:

PolicyIter():

1: Initiate with **random**  $v_\pi(s)$  for all **non-terminal** states  $s$

2: Set  $v_\pi(s) = 0$  for **terminal** states  $s$

3: Initiate two random policies  $\pi$  and  $\bar{\pi}$

4: **while**  $\pi \neq \bar{\pi}$  **do**

5:  $v_\pi = \text{PolicyEval}(\pi, v_\pi)$  and  $\pi \leftarrow \bar{\pi}$  **Recursion**

6:  $\bar{\pi} = \text{PolicyImprov}(v_\pi)$

7: **end while**

Recursion

We should start with  $v_{\pi^0}^0(\cdot)$  and do the **red recursion** first

- at the end of this recursion we have evaluated the random policy

## Example: A Bit Larger Grid World

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

 $v_{\pi^0}^0$ 

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

 $v_{\pi^0}^1$ 

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

 $v_{\pi^0}^2$ 

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

 $v_{\pi^0}^3$ 

...

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

 $v_{\pi^0}^{\infty}$ 

We now have *evaluated* the value of random policy  $v_{\pi^0} = v_{\pi^0}^{\infty}$



## Example: A Bit Larger Grid World

Recall policy iteration:

PolicyIter():

1: Initiate with **random**  $v_\pi(s)$  for all **non-terminal** states  $s$

2: Set  $v_\pi(s) = 0$  for **terminal** states  $s$

3: Initiate two random policies  $\pi$  and  $\bar{\pi}$

4: **while**  $\pi \neq \bar{\pi}$  **do**

5:  $v_\pi = \text{PolicyEval}(\pi, v_\pi)$  and  $\pi \leftarrow \bar{\pi}$  **Recursion**

6:  $\bar{\pi} = \text{PolicyImprov}(v_\pi)$

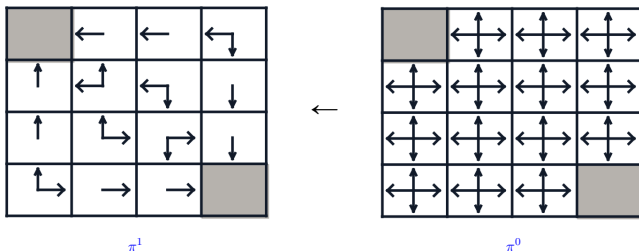
7: **end while**

Recursion

Next, we do the **outer recursion** recursion, i.e.,

- we **improve** the policy

## Example: A Bit Larger Grid World



We improve policy by taking actions with **maximal action-values**

- if we have multiple **maximal action-values** we can behave **randomly**

## Example: A Bit Larger Grid World

Recall policy iteration:

PolicyIter():

1: Initiate with **random**  $v_{\pi}(s)$  for all **non-terminal** states  $s$

2: Set  $v_{\pi}(s) = 0$  for **terminal** states  $s$

3: Initiate two random policies  $\pi$  and  $\bar{\pi}$

4: **while**  $\pi \neq \bar{\pi}$  **do**

5:  $v_{\pi} = \text{PolicyEval}(\pi, v_{\pi})$  and  $\pi \leftarrow \bar{\pi}$  **Recursion**

6:  $\bar{\pi} = \text{PolicyImprov}(v_{\pi})$

7: **end while**

Recursion

We now  $v_{\pi^1}^0 = v_{\pi^0} = v_{\pi^0}^{\infty}$  and do the **red recursion** again

- at the end of this recursion we have evaluated the **new policy**  $\pi^1$

## Example: A Bit Larger Grid World

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

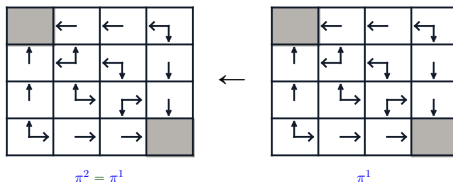
 $v_{\pi 1}^0$ 

• • •

0.0	-1.0	-2.0	-3.0
-1.0	-2.0	-3.0	-2.0
-2.0	-3.0	-2.0	-1.0
-3.0	-2.0	-1.0	0.0

 $v_{\pi^1}^{+\infty}$ 

After **evaluating** policy  $\pi^1$  as  $v_{\pi^1} = v_{\pi^1}^\infty$ , we do the next improvement



Well  $\pi^2 = \pi^1$  and we should **stop!**

## Example: A Bit Larger Grid World

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

initial values

Now we try *value iteration*: for start, we only need an initial value, so we set

- all values being zero, i.e.,  $v_{\star}^0(s) = 0$  for all  $s$

We keep recursion until we find the *optimal values*

## Example: A Bit Larger Grid World

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

 $v_{\star}^0$ 

...

0.0	-1.0	-2.0	-3.0
-1.0	-2.0	-3.0	-2.0
-2.0	-3.0	-2.0	-1.0
-3.0	-2.0	-1.0	0.0

 $v_{\star}^{+\infty}$ 

Now, we back-track the optimal policy  $\pi^{\star}$

0.0	-1.0	-2.0	-3.0
-1.0	-2.0	-3.0	-2.0
-2.0	-3.0	-2.0	-1.0
-3.0	-2.0	-1.0	0.0

 $v_{\star}$ 

action-values



	←	←	↙
↑	↖	↙	↓
↑	↗	↘	↓
↘	→	→	

 $\pi^{\star}$

# Complexity of Policy and Value Iteration

- + It seems that value iteration has *less complexity*!
- Well, it is not in order, but yes! It usually converge faster

In our example with *policy iteration*, we had to *evaluate two policies*

- once for  $\pi^0$  and once for  $\pi^1$
- say the first recursion took  $K_1$  iterations and the second took  $K_2$ 
  - ↳ the *total number* of iterations is then  $K_1 + K_2$
  - ↳ in practice, it often happens that  $K_2 \ll K_1$ 
    - ↳ because we already start from *good values* with  $v_{\pi^1}^0 = v_{\pi^0}^{+\infty}$

With *value iteration*, we had to *only evaluate optimal policy*

- say it takes  $K_*$  iterations: there is *no reason* that  $K_*$  be same as  $K_1$  or  $K_2$ 
  - ↳ each evaluation has a *different initial* and *converging* point
- ↳ in practice, it often happens that  $K_* > K_1$  and  $K_* \gg K_2$ 
  - ↳ so it *might be* that  $K_* \approx K_1 + K_2$
  - ↳ but usually with *multiple policy improvements*, we see  $K_* < K_1 + K_2 + \dots$

# Complexity of Policy and Value Iteration

- + If so, why should we use *policy iteration*?!
  - Well, not all problems are like a dummy grid world

In practice, it might be computationally *hard* to get *very close to optimal values*

- in this case, we take non-converged values
  - ↳ we consider them *estimates* of optimal values
- in value iteration we *approximate* optimal policy with on these *estimates*
  - ↳ this might be a *loose* estimate

If we do the same *approximative* computation with policy iteration

- we often end up with a *better policy*

## Moral of Story

While *value iteration* typically show *faster convergence*, *policy iteration* can give *better policies* after convergence



# Generalized Policy Iteration

*In practice, we can terminate or change the order of computation in policy iteration to reduce its complexity: for instance, we could have*

GenPolicyItr():

1: Initiate with **random**  $v_\pi(s)$  for all **non-terminal** states  $s$

2: Set  $v_\pi(s) = 0$  for **terminal** states  $s$

3: Initiate two random policies  $\pi$  and  $\bar{\pi}$

4: **while**  $\pi \neq \bar{\pi}$  **do**

5:  $v_\pi = \text{TerminPolicyEval}(\pi, v_\pi)$  and  $\pi \leftarrow$  **changed**

6:  $\bar{\pi} = \text{PolicyImprov}(v_\pi)$

7: **end while**

where  $\text{TerminPolicyEval}(\pi, v_\pi)$  evaluates **policy**  $\pi$  from starting value function  $v_\pi$  with a **terminating** recursion loop

# Generalized Policy Iteration: *Terminating Evaluation*

TerminPolicyEval( $\pi, v_\pi^0$ ):

```

1: Initiate values with  $v_\pi^0$  and set  $k = 0$ 
2: Make sure that  $v_\pi^0(s) = 0$  for terminal states  $s$ 
3: Choose a small threshold  $\epsilon$  and initiate  $\Delta = +\infty$  # stopping criteria
4: for  $n = 1 : N$  do
5:   Compute  $\bar{\mathcal{R}}_\pi(s^n) = \mathbb{E}_\pi \{ \bar{\mathcal{R}}(s^n, a) \}$  # average response
6: end for
7: while  $\Delta > \epsilon$  and  $k < K$  do changed
8:   for  $n = 1 : N$  do
9:     Update  $v_\pi^{k+1}(s^n) = \bar{\mathcal{R}}_\pi(s^n) + \gamma \mathbb{E}_\pi \{ v_\pi^k(\bar{S}) | s^n \}$  # DP update
10:   end for
11:    $\Delta = \max_n |v_\pi^{k+1}(s^n) - v_\pi^k(s^n)|$  # check convergence
12:   Update  $k \leftarrow k + 1$ 
13: end while

```

Obviously, TerminPolicyEval( $\pi, v_\pi$ ) does **not** return the **exact values** of the policy  $\pi$ , but only an **estimate** of them

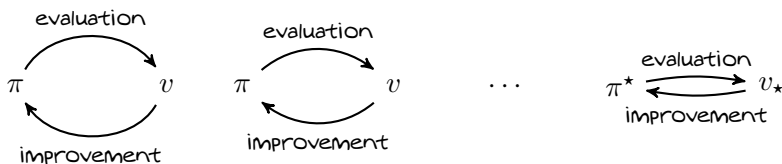
# Generalized Policy Iteration

We can come up with various such ideas: *these variants are often called*

*Generalized Policy Iteration  $\equiv$  GPI*

*These approaches all rely on*

*back-and-forth computation of policies and values*



*If designed properly, they all converge to optimal policy and optimal values*

## Some Final Remarks

- + We know the algorithms now, but how can we **guarantee** that they **converge**? You showed us an simple example that recursion could simply **diverge**!
- Well, we can show that what we discussed in this chapter converge: it comes from the nice properties of **Bellman equations**
  - ↳ There are several proofs; for instance see [a proof in Tom Mitchell's notes](#)

When it comes to practice, most known algorithms are proved to **converge to optimal policy and optimal values**; however, note that

- **Convergence guarantee** is different from the **speed of convergence**
  - ↳ An algorithm might **converge**, but **very slow**
- If you deal with an **unknown** algorithm; then, you should make sure that it **converges to optimal policy and optimal values**