

ECE 1508: Reinforcement Learning

Chapter 3: Model-free RL

Ali Bereyhi

`ali.bereyhi@utoronto.ca`

Department of Electrical and Computer Engineering
University of Toronto

Summer 2024

Classical RL Methods: Recall

Ultimate goal in an RL problem is to find the *optimal policy*

As mentioned, we have two *major challenges* in this way

- 1 We need to compute *values* explicitly
- 2 We often deal with settings with *huge state spaces*?

In this part of the course, we are going to handle the first challenge

- *Previous chapter \rightsquigarrow Model-based methods*
- *This chapter \rightsquigarrow Model-free methods*

Finally We Got Serious: *Model-free RL*

In *model-free* methods

we *do not have* an *analytic model* for the behavior of *environment*

We intend to compute *values* from *real data* collected from *environment*

Model-Based RL

Bellman Equation

value iteration

policy iteration

Model-free RL

on-policy methods

temporal difference

Monte Carlo

SARSA

off-policy methods

Q-learning

Model-free RL in Nutshell

- + If this is the *typical case* in RL problems, why did we spend so much time on learning *MDPs* and finding *optimal policy* there?
- Well! We need all those things, since we are going to do *the same thing* here *only without explicit model*

In a nutshell, we are going to find a way to apply

Generalized Policy Iteration \equiv GPI

But, now *without* knowing the *transition-rewarding function*

Let's take a look back at GPI

Generalized Policy Iteration

We wrote the pseudo-code for GPI as below

```
GenPolicyItr():  
  1: Initiate two random policies  $\pi$  and  $\bar{\pi}$   
  2: while  $\pi \neq \bar{\pi}$  do  
  3:    $v_\pi = \text{GenPolicyEval}(\pi)$  and  $\pi \leftarrow \bar{\pi}$   
  4:    $\bar{\pi} = \text{PolicyImprov}(v_\pi)$   
  5: end while
```

Let's recall where we had to use **environment's model**

- 1 In **policy evaluation** phase when we compute values via **Bellman equations**
- 2 In **policy improvement** when we compute action-values out of **values**

How can we do these tasks **without** knowing **transition-rewarding model**?

Computing Statistics from Data

Let's start with a very simple problem: assume we have an **unknown** signal generator which returns signals at **random**; this generator is connected to a device and we can **only see** the output of this device, i.e., we see

$$Y = f(X)$$

where X is the **random** signal and $f(\cdot)$ denotes transform by the device

We want to know the expected output of our device, i.e.,

$$\mu_Y = \mathbb{E}\{Y\} = \mathbb{E}\{f(X)\}$$

If we knew the model of the **generator's model**, we could write

$$\mu_Y = \mathbb{E}\{f(X)\} = \sum_{\substack{x \in \mathbb{X} \\ \text{all outcomes}}} f(x) \underbrace{p(x)}_{\text{model}}$$

Monte-Carlo Method

Now what can we do if we *don't know* the *model*

- + Well! Shouldn't we evaluate it by a *simple* numerical simulation?
- Exactly! This is what we call it *Monte-Carlo method*

In *Monte-Carlo method*, we sample our device K times *independently* as

$$Y_1, Y_2, \dots, Y_K$$

Then we *estimate* the *expected* value as

$$\hat{\mu}_Y = \frac{1}{K} \sum_{k=1}^K Y_k$$

Monte-Carlo Method

- + Why does Monte-Carlo work?
- Simply because of **central limit theorem**

Since the sequence Y_1, Y_2, \dots, Y_K contains **independent** samples of **identical process**, we could say that

$$\hat{\mu}_Y \sim \mathcal{N}\left(\mu_Y, \frac{\sigma^2}{K}\right)$$

when K is **large enough**: so we could think of it as

$$\hat{\mu}_Y \approx \mu_Y + \frac{\varepsilon}{\sqrt{K}}$$

for some random **error** term ε : this error vanishes as K goes **large**

Computing Values via Monte-Carlo

- + But, how can we apply this idea to RL? I don't see any connection!
- Well! Think of **rewards and transitions** as random signal and **value function** as device! We only need to take **enough** samples from the **environment**

Let's start with a very simple task: we want to compute the value of **state s** for **policy π** in an **episodic** environment. Monte-Carlo suggest that


- 1 We start at **state s** and play with **policy π** until we meet **terminal state**: say it happens at **time T**
- 2 We compute the **sample** return as $G[1] = R_1 + \gamma R_2 + \dots + \gamma^{T-1} R_T$
- 3 We repeat this for K episodes and each episode, we collect $G[k]$

Then, we could estimate the value of **state s** as

$$\hat{v}_{\pi}(s) = \frac{1}{K} \sum_{k=1}^K G[k]$$


Values via Monte-Carlo: *Trajectory Sampling*

We can look at this approach as estimating values from *sample trajectories*: with known model, we can compute values by averaging over *possible trajectories*




Values via Monte-Carlo: *Trajectory Sampling*

We can look at this approach as estimating values from *sample trajectories*: *without* known model, we can *sample* them and estimate values from them



Values via Monte-Carlo: *Trajectory Sampling*

We can look at this approach as estimating values from *sample trajectories*: *without* known model, we can *sample* them and estimate values from them



Computing Values via Monte-Carlo: Algorithm 1

Let's put our estimation approach into an algorithm

MC_verI(π, s):

- 1: Initiate estimator of value as $\hat{v}_{\pi}(s) = 0$
- 2: **for** episode = 1 : K **do**
- 3: Initiate with **state** $S_0 = s$ and act via policy $\pi(a|s)$
- 4: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminal}$$

- 5: Compute sample return $G = R_1 + \gamma R_2 + \dots + \gamma^{T-1} R_T$
- 6: Update estimate of value as $\hat{v}_{\pi}(s) \leftarrow \hat{v}_{\pi}(s) + G/K$
- 7: **end for**

Computing Values via Monte-Carlo

- + *But, doesn't that take too long to compute a single value?*
- Yes! This is in general a problem; however, in our **naive** algorithm it is too much delayed!

*In our algorithm, we need to wait **till very end of K episodes** to access an estimate, but we rather prefer to have a **bad** estimate which **gradually improves** over episodes*

*We could use the idea of **online** averaging \equiv **incremental** averaging*

Let's find out what it is!

Online Averaging

Say, we want to compute the average of K samples: we could write

$$\begin{aligned}
 \eta_K &= \frac{1}{K} \sum_{k=1}^K G_k = \frac{1}{K} \left(\sum_{k=1}^{K-1} G_k + G_K \right) = \frac{1}{K} ((K-1) \eta_{K-1} + G_K) \\
 &= \left(1 - \frac{1}{K} \right) \eta_{K-1} + \frac{G_K}{K} \\
 &= \eta_{K-1} + \frac{1}{K} (G_K - \eta_{K-1})
 \end{aligned}$$

But, we can define the previous average as

$$\eta_{K-1} = \frac{1}{K-1} \sum_{k=1}^{K-1} G_k \rightsquigarrow \sum_{k=1}^{K-1} G_k = (K-1) \eta_{K-1}$$

Online Averaging: Geometric Weights

Online Averaging

We can update the average in *online fashion* as

$$\eta_K = \eta_{K-1} + \frac{1}{K} \Delta_K$$

where $\Delta_K = G_K - \eta_{K-1}$ is the deviation in K -th episode

The above expression is given for *uniform averaging weights*, i.e., all samples have same weights: in more general form, we usually update

$$\eta_K = \eta_{K-1} + \alpha \Delta_K$$

for some $0 < \alpha \leq 1$ that can be *fixed* or *scaled with K*

- if it is *fixed* \equiv computing weighted average with *geometric* weights
- if it is *scaled linearly with K* \equiv computing *linear* averaging

Computing Values via Monte-Carlo: Algorithm II

Let's modify our earlier algorithm with online averaging

MC_verII(π, s):

1: Initiate estimator of value as $\hat{v}_{\pi}(s) = 0$

2: **for** episode = 1 : K **do**

3: Initiate with **state** $S_0 = s$ and act via policy $\pi(a|s)$

4: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminal}$$

5: Compute sample return $G = R_1 + \gamma R_2 + \dots + \gamma^{T-1} R_T$

6: Update estimate of value as $\hat{v}_{\pi}(s) \leftarrow \hat{v}_{\pi}(s) + \alpha(G - \hat{v}_{\pi}(s))$

7: **end for**

*Now after each episode, we have an estimate of value function at **state** s*

Computing Values via Monte-Carlo: *Improve Efficiency*

In our algorithm: we go through the whole trajectory to compute the value on the state we started with! This does not sound *sample efficient*!

- + Well! What can we do *more*?! It seems to be the case!
- **Not really**! We can estimate values of *other states* down the *trajectory*!

In the following sample trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

it's not only S_0 whose sample return can be computed! We can also compute sample returns of S_1, \dots, S_{T-1}

All-Visit Monte-Carlo: Algorithm III

This concludes a policy evaluation algorithm based on [Monte-Carlo](#)

MC_Eval(π):

1: Initiate estimator of value as $\hat{v}_\pi(s^n) = 0$ for $n = 1 : N$

2: **for** episode = 1 : K **do**

3: Initiate with a **random state** S_0 and act via policy $\pi(a|s)$

4: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminal}$$

5: Initiate with $G = 0$

6: **for** $t = T - 1 : 0$ **do**

7: Update current return $G \leftarrow R_{t+1} + \gamma G$

8: Update estimate of value as $\hat{v}_\pi(S_t) \leftarrow \hat{v}_\pi(S_t) + \alpha(G - \hat{v}_\pi(S_t))$

9: **end for**

10: **end for**

All-Visit Monte-Carlo: Convergence

It's *intuitive* to say this algorithm converges to true values after *lots of episodes*

Asymptotic Convergence of Monte-Carlo

Let $\mathcal{C}_K(s)$ denote number of visits at *state* s during K Monte-Carlo episodes. Assume that the random state initialization is distributed such that $\mathcal{C}_K(s^n)$ grows large as K increases for $n = 1 : N$, i.e.,


$$\lim_{K \rightarrow \infty} \mathcal{C}_K(s^n) = \infty$$

Then, as $K \rightarrow \infty$ the estimator of value function converges to its exact expression, i.e.,

$$\hat{v}_\pi(s) \xrightarrow{K \uparrow \infty} v_\pi(s)$$

for any *state* s

Example: Dummy Grid World with Random Walk




Let's get back to our *dummy world*: we now use *Monte-Carlo method* to compute the values for *uniform random policy*, i.e.,

$$\pi(a|s) = \frac{1}{4}$$

for all *actions* and *states*. From *Bellman equations*, we have

$$v_{\pi}(0) = 1 \quad v_{\pi}(1) = -4.5 \quad v_{\pi}(2) = -4.5 \quad v_{\pi}(3) = -6$$

Example: *Dummy Grid World with Random Walk*




Typical Behavior: *Variation Against Number of Episodes*

We can compute the error of our estimation in each *episode*

$$\text{RMS} = \sqrt{\sum_{n=1}^N |\hat{v}_{\pi}(s^n) - v_{\pi}(s^n)|^2}$$

if we know the true value function, e.g., our random walk example

If we plot it against K ; then, we see




Typical Behavior: Variation Against Averaging Coefficient

We can compute the error of our estimation in each *episode*

$$\text{RMS} = \sqrt{\sum_{n=1}^N |\hat{v}_{\pi}(s^n) - v_{\pi}(s^n)|^2}$$

if we know the true value function, e.g., our random walk example

If we plot it against α ; then, we could see a *minimum*



Monte-Carlo Method: Action-Values

- + Now that we have *Monte-Carlo algorithm*, can we use it in GPI?
- Not yet! Remember that we need *action-values* for *policy improvement*

In GPI, we used to use *Bellman equation* for this

$$\begin{aligned}
 q_{\pi}(s, a) &= \bar{\mathcal{R}}(s, a) + \gamma \mathbb{E} \{ v_{\pi}(\bar{S}) | s, a \} \\
 &= \mathbb{E} \{ R_{t+1} | S_t = s, A_t = a \} + \gamma \mathbb{E} \{ v_{\pi}(S_{t+1}) | S_t = s, A_t = a \} \\
 &= \sum_{\ell=1}^L \sum_{n=1}^N \left(r^{\ell} + \gamma v_{\pi}(s^n) \right) \underbrace{p(r^{\ell}, s^n | s, a)}_{\text{transition-rewarding model}}
 \end{aligned}$$

But, now we cannot use it anymore!

Maybe, we can use *Monte-Carlo* method to estimate *action-values* directly

All-Visit Monte-Carlo: Action-Values

MC_QEval(π):

- 1: Initiate estimator as $\hat{q}_\pi(s^n, a^m) = 0$ for $n = 1 : N$ and $m = 1 : M$
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a random state-action pair (S_0, A_0) and act via policy $\pi(a|s)$
- 4: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminal}$$

- 5: Initiate with $G = 0$
- 6: **for** $t = T - 1 : 0$ **do**
- 7: Update current return $G \leftarrow R_{t+1} + \gamma G$
- 8: Update $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(G - \hat{q}_\pi(S_t, A_t))$
- 9: **end for**
- 10: **end for**

We can now apply GPI using the Monte-Carlo method!

Policy Iteration with Monte-Carlo

We can use *Monte-Carlo method* to compute the *action-values*

- We then improve in each iteration by selecting best *action* for each *state*
 - ↳ This is what we typically call *greedy* improvement

```
MC_PolicyItr():
```

```
1: Initiate two random policies  $\pi$  and  $\bar{\pi}$ 
```

```
2: while  $\pi \neq \bar{\pi}$  do
```

```
3:    $\hat{q}_{\pi} = \text{MC\_QEval}(\pi)$  and  $\pi \leftarrow \bar{\pi}$ 
```

```
4:    $\bar{\pi} = \text{Greedy}(\hat{q}_{\pi})$ 
```

```
5: end while
```

Policy Iteration with Monte-Carlo

Algorithmically, we can write the greedy update as

Greedy(\hat{q}_π):

1: **for** $n = 1 : N$ **do**

2: Improve the by taking *deterministically* the *best action*

$$\bar{\pi}(a^m | s^n) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} \hat{q}_\pi(s^n, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} \hat{q}_\pi(s^n, a^m) \end{cases}$$

3: **end for**


This is however not the *best* we could do!

We are going to have a whole lecture about it


Stay tuned! We get back to this point in Section 4

GPI with Monte-Carlo

For any GPI, we said that we can think of




With **Monte-Carlo** evaluation, we can show this procedure as



GPI with Monte-Carlo

Another way to visualize this procedure is to think of following diagram



Non-Episodic Monte-Carlo: *Terminating Trajectory*

- + We only discussed *episodic* scenarios! Don't we use model-free RL in *non-episodic environment*?
- Sure we do! But, *Monte-Carlo* is not the best approach

A basic idea in this case is to *terminate sample trajectories*

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- With *long enough* T and $\gamma < 1$ the very later terms are *ineffective*
- But, we *cannot* use all the states in the trajectories
 - ↳ Sample returns of those who are *close to time* T are *not reliable*!

$$G_{T-1} = R_T + \underbrace{\gamma R_{T+1} + \dots}_{\text{we terminated them!}}$$

Terminating Monte-Carlo

TerminMC_Eval(π):

- 1: Initiate estimator of value as $\hat{v}_{\pi}(s^n) = 0$ for $n = 1 : N$
- 2: Choose **very large T and W** that satisfy $W < T$
- 3: **for** episode = $1 : K$ **do**
- 4: Initiate with a **random state S_0** and act via policy $\pi(a|s)$
- 5: Sample a trajectory and terminate after **T time steps**

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminated}$$

- 6: Initiate with $G = 0$
- 7: **for** $t = T - 1 : 0$ **do**
- 8: Update current return $G \leftarrow R_{t+1} + \gamma G$
- 9: **if** $t < T - W$ **then**
- 10: Update estimate of value as $\hat{v}_{\pi}(S_t) \leftarrow \hat{v}_{\pi}(S_t) + \alpha(G - \hat{v}_{\pi}(S_t))$
- 11: **end if**
- 12: **end for**
- 13: **end for**

Terminating Monte-Carlo: *Sample Inefficiency*

- + *But, isn't that as you said before **sample inefficient**?*
- Sure it is!

We are losing lots of states in each sample trajectory!

*A better solution is to use the **recursive property** of return and do*

bootstrapping

This is what we see next!

Testing Genie in the Box

Let's think again a bit science-fictional: assume a *genie* can tell us the value $v_\pi(s)$ for *each state* s , We want to test this *genie* via a numerical algorithm 😊

Bellman tells us that at any *state* s , we should see

$$v_\pi(s) = \mathbb{E} \{R_{t+1} | S_t = s\} + \gamma \mathbb{E} \{v_\pi(S_{t+1}) | S_t = s\}$$

Monte-Carlo tells us further that after K *sample trajectories* $S_0, A_0 \xrightarrow{R_1} S_1$ initiated at $S_0 = s$ and *terminated after only one step*, we have

$$\mathbb{E} \{R_{t+1} | S_t = s\} \approx \frac{1}{K} \sum_{k=1}^K R_1[k]$$

$$\mathbb{E} \{v_\pi(S_{t+1}) | S_t = s\} \approx \frac{1}{K} \sum_{k=1}^K v_\pi(S_1[k])$$

Testing Genie in the Box

We could hence find an estimator of $v_\pi(s)$ as

$$v_\pi(s) \approx \hat{v}_\pi(s) = \frac{1}{K} \sum_{k=1}^K R_1[k] + \gamma v_\pi(S_1[k])$$

And, of course we could simply evaluate this estimator **online** as

$$\hat{v}_\pi(s) \leftarrow \hat{v}_\pi(s) + \frac{1}{K} (R_1 + \gamma v_\pi(S_1) - \hat{v}_\pi(s))$$

if we need linear averaging or alternatively as

$$\hat{v}_\pi(s) \leftarrow \hat{v}_\pi(s) + \alpha (R_1 + \gamma v_\pi(S_1) - \hat{v}_\pi(s))$$

if we think of more general weighted averaging

Computing Values via Bootstrapping: Algorithm 1

We can write our genie-testing algorithm as

$\text{TD_verI}(\pi, s):$

- 1: Initiate estimator of value as $\hat{v}_\pi(s) = 0$
- 2: Ask genie $v_\pi(\bar{s})$ for all \bar{s} that can be followed after s
- 3: **for** episode = 1 : K **do**
- 4: Initiate with **state** $S_0 = s$ and act via policy $\pi(a|s)$
- 5: Sample a single-step **terminated** trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1$$

- 6: Update estimate of value as $\hat{v}_\pi(s) \leftarrow \hat{v}_\pi(s) + \alpha(R_1 + v_\pi(S_1) - \hat{v}_\pi(s))$
- 7: **end for**

Note that in this algorithm we **don't** need the environment to be **episodic**, as we use **recursive property** of value-function!

This explains the idea of **bootstrapping**

Bootstrapping: Using Value Estimates

- + But, in practice we don't have *genie*! And, say we find one, why should we compute values anymore?!
- Absolutely! But, we may use *this property*

We can replace those true values with their estimates

- They are initially *bad* estimates
 - ↳ and thus return in bad estimate for *the other state*
- They *gradually improve*
 - ↳ and therefore return *better estimate* for *the other state*

The key point is that we get rid of need for a *terminal state*!

Computing Values via Bootstrapping: Algorithm II

So, we could get rid of the *genie* finally

$\text{TD_verII}(\pi, s):$


- 1: Initiate estimator of value as $\hat{v}_{\pi}(s) = 0$
- 2: Use available $\hat{v}_{\pi}(\bar{s})$ for all \bar{s} that follow s
- 3: **for** episode = 1 : K **do**
- 4: Initiate with *state* $S_0 = s$ and act via policy $\pi(a|s)$
- 5: Sample a single-step *terminated* trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1$$

- 6: Update estimate of value as $\hat{v}_{\pi}(s) \leftarrow \hat{v}_{\pi}(s) + \alpha(R_1 + \hat{v}_{\pi}(S_1) - \hat{v}_{\pi}(s))$
- 7: **end for**


Backup Diagram: *Sampling with Bootstrapping*

Looking at the backup tree, we are now sampling *one-step trajectories*




Backup Diagram: *Sampling with Bootstrapping*

Looking at the backup tree, we are now sampling *one-step trajectories*




Backup Diagram: *Sampling with Bootstrapping*

Looking at the backup tree, we are now sampling *one-step trajectories*



Backup Diagram: *Sampling with Bootstrapping*

Looking at the backup tree, we are now sampling *one-step trajectories*



Temporal Difference

Bootstrapping can replace **Monte-Carlo** in our evaluation algorithms

- 1 We start with some initial value estimates
- 2 We sample a trajectory of finite length T

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- ↳ it can either end with a terminal state if we have any
- ↳ or could simply be terminated

- 3 We move over trajectory and update value of each state by **bootstrapping**

This idea of estimating values is called

Temporal Difference \equiv TD

Temporal Difference: TD-0

TD_Eval(π):

- 1: Initiate estimator of value as $\hat{v}_{\pi}(s^n) = 0$ for $n = 1 : N$
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a **random state** S_0 and act via policy $\pi(a|s)$
- 4: Sample a trajectory until either a **terminal** state or some **terminating** T

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - 1$ **do**
- 6: Update as $\hat{v}_{\pi}(S_t) \leftarrow \hat{v}_{\pi}(S_t) + \alpha(R_{t+1} + \gamma \hat{v}_{\pi}(S_{t+1}) - \hat{v}_{\pi}(S_t))$
- 7: **end for**
- 8: **end for**

Attention

With TD, we even don't need to wait till a trajectory is sampled!

Evaluating Action-Values via Bootstrapping

- + What about the **action-values**? Can we **bootstrap** again?
- Sure!

Recall **Bellman equation** of action-value function

$$q_{\pi}(s, a) = \mathbb{E} \{R_{t+1} | S_t = s, A_t = a\} + \gamma \mathbb{E} \{v_{\pi}(S_{t+1}) | S_t = s, A_t = a\}$$

So, we can use sample trajectory to estimate **action-values** as well: at time t , we can update estimate of pair (S_t, A_t) as

$$\hat{q}_{\pi}(S_t, A_t) \leftarrow \hat{q}_{\pi}(S_t, A_t) + \alpha(R_{t+1} + \gamma \hat{v}_{\pi}(S_{t+1}) - \hat{q}_{\pi}(S_t, A_t))$$

Temporal Difference: Action-Value


TD_QEval(π):

- 1: Initiate estimator of value as $\hat{q}_\pi(s^n, a^m) = 0$ for $n = 1 : N$ and $m = 1 : M$
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a random state-action pair (S_0, A_0) and act via policy $\pi(a|s)$
- 4: Sample a trajectory until either a terminal state or some terminating T

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - 1$ **do**
- 6: Set $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(R_{t+1} + \gamma \hat{v}_\pi(S_{t+1}) - \hat{q}_\pi(S_t, A_t))$
- 7: **end for**
- 8: **end for**

Example: Dummy Grid World with Random Walk



Let's get back to our *dummy world*: we now use *TD* to compute the values for *uniform random policy*, i.e.,


$$\pi(a|s) = \frac{1}{4}$$

for all *actions* and *states*


Try it at home 😊

Typical Behavior

We are going to see the same behavior also with TD: *against K we see*



and against α we have a *minimum*



Example: TD vs Monte-Carlo in Single-Button Game



Consider the following dummy game: we have got a single button to push; each time we push this button,

- we either get into **green** or **blue** mode allowing us to push the button again and returns a 0/1 reward
- or it gets into **red** mode which only returns a 0/1 reward and game is over

Obviously this game has

- Three states: **green**, **blue** and **red** which is **terminal**
- a single action, i.e., pushing the button

Example: TD vs Monte-Carlo in Single-Button Game



We play this game 6 episodes and get following *sample trajectories*

Blue $\xrightarrow{1}$ Blue $\xrightarrow{0}$ red

Blue $\xrightarrow{1}$ red

Blue $\xrightarrow{0}$ Blue $\xrightarrow{1}$ red

Blue $\xrightarrow{1}$ red

Blue $\xrightarrow{0}$ red

Green $\xrightarrow{0}$ Blue $\xrightarrow{0}$ red

Let's estimate $v(\text{blue})$ and $v(\text{green})$ by both TD-0 and Monte-Carlo

Example: TD vs Monte-Carlo in Single-Button Game

Blue $\xrightarrow{1}$ Blue $\xrightarrow{0}$ red
 Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{0}$ Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{0}$ red
 Green $\xrightarrow{0}$ Blue $\xrightarrow{0}$ red

With Monte-Carlo approach, we could say: we have 8 sample trajectories starting with **Blue** with 4 returning 1 and 4 returning 0; thus we have

$$\hat{v}(\text{blue}) = \frac{4}{8} = 0.5$$

We also have only one sample trajectory starting at **Green** with zero return; thus, we have

$$\hat{v}(\text{green}) = \frac{0}{1} = 0$$

Example: TD vs Monte-Carlo in Single-Button Game

Blue $\xrightarrow{1}$ Blue $\xrightarrow{0}$ red
 Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{0}$ Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{0}$ red
 Green $\xrightarrow{0}$ Blue $\xrightarrow{0}$ red

With TD-0, we could say: we have 8 Blue states followed by either Blue or red. If we bootstrap, we then get up to state Blue in the last trajectory


$$\hat{v}(\text{Blue}) \approx 0.5$$

We then get to the last trajectory which is the only one with state Green: since we have not yet updated, we yet have $\hat{v}(\text{Green}) = 0$; by bootstrapping we get

$$\hat{v}(\text{Green}) \leftarrow \underbrace{\hat{v}(\text{Green})}_0 + \left(\underbrace{R_{t+1}}_0 + \underbrace{\hat{v}(S_{t+1})}_{\hat{v}(\text{Blue}) \approx 0.5} - \underbrace{\hat{v}(\text{Green})}_0 \right) \approx 0.5$$

Temporal Difference vs Monte-Carlo: Note I

The observed difference follow a fundamental point: in *Monte-Carlo* we only look at *best approximation given data*, without looking into the *Markovity* of the state, whereas in *TD* we take into account the fact that we are dealing with a *Markov state*



Temporal Difference vs Monte-Carlo: Note I



This is common to see in the literature that people say

- TD finds *maximum-likelihood* estimate of the values
 - ↳ It uses this assumption that the state is a *Markov process*
 - ↳ It's the *better option*, if we are *sure* that we have access to the *complete (Markov) state*
- Monte-Carlo finds *least-squares* estimate of the values
 - ↳ It *ignores Markovity* of the state
 - ↳ Maybe *better option*, when we *cannot* access the *complete (Markov) state*

Back to Single-Button Game: *Side Note on Batch Updating*



What would happen, if we get *sample trajectories* in the following order

Green $\xrightarrow{0}$ Blue $\xrightarrow{0}$ red

Blue $\xrightarrow{1}$ Blue $\xrightarrow{0}$ red

Blue $\xrightarrow{1}$ red

Blue $\xrightarrow{0}$ Blue $\xrightarrow{1}$ red

Blue $\xrightarrow{1}$ red

Blue $\xrightarrow{0}$ red

Back to Single-Button Game: *Side Note on Batch Updating*

green $\xrightarrow{0}$ Blue $\xrightarrow{0}$ red
 Blue $\xrightarrow{1}$ Blue $\xrightarrow{0}$ red
 Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{0}$ Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{0}$ red

If we go *only once* over the *batch* of all episodes with TD, we get

$$\hat{v}(\text{green}) \approx 0$$

$$\hat{v}(\text{blue}) \approx 0.5$$

Back to Single-Button Game: *Side Note on Batch Updating*

green $\xrightarrow{0}$ Blue $\xrightarrow{0}$ red
 Blue $\xrightarrow{1}$ Blue $\xrightarrow{0}$ red
 Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{0}$ Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{1}$ red
 Blue $\xrightarrow{0}$ red

If we go *twice* over the *batch* of all episodes with TD, we get


$$\hat{v}(\text{green}) \approx 0.5$$

$$\hat{v}(\text{blue}) \approx 0.5$$

We get *better* if we go over the *batch of data* multiple times!

Temporal Difference vs Monte-Carlo: Note II

Let's see how Monte-Carlo performs against TD-0 algorithm for a bit larger example of *random walk* on a grid¹



¹This figure is taken from Chapter 6 of Sutton and Barto's book

Recall: *Bias and Variance of Estimator*

At this point, we need to have some clue about *bias* and *variance* of an estimator

If you need to recap, please look at the board

Temporal Difference vs Monte-Carlo: Note II

What has been seen in the diagram is a **general behavior**

- Monte-Carlo is good in sense of **bias** but bad in terms of **variance**

↳ It always returns an unbiased estimator of **value**, i.e.,

$$\mathbb{E} \{ \hat{v}_{\pi}(s) \} = v_{\pi}(s)$$

↳ It's estimation is however high **variance**, i.e.,

$$\mathbb{E} \left\{ (\hat{v}_{\pi}(s) - v_{\pi}(s))^2 \right\} \rightsquigarrow \text{large}$$

- TD-0 is good in sense of **variance** but can be bad in terms of **bias**

↳ It can return a biased estimator of **value**, i.e.,

$$\mathbb{E} \{ \hat{v}_{\pi}(s) \} \neq v_{\pi}(s)$$

↳ It's estimation is low **variance**, i.e.,

$$\mathbb{E} \left\{ (\hat{v}_{\pi}(s) - v_{\pi}(s))^2 \right\} \rightsquigarrow \text{small}$$

Policy Iteration with TD-0

We can use TD-0 to implement another variant of GPI

MC_PolicyItr():

- 1: Initiate two random policies π and $\bar{\pi}$
- 2: **while** $\pi \neq \bar{\pi}$ **do**
- 3: $\hat{q}_\pi = \text{TD_QEval}(\pi)$ and $\pi \leftarrow \bar{\pi}$
- 4: $\bar{\pi} = \text{Greedy}(\hat{q}_\pi)$
- 5: **end while**

And, recall that our greedy algorithm is

Greedy(q_π):


- 1: **for** $n = 1 : N$ **do**
- 2: Improve the by taking *deterministically* the *best action*

$$\bar{\pi}(a^m | s^n) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_\pi(s^n, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_\pi(s^n, a^m) \end{cases}$$

- 3: **end for**


GPI with TD-0: Visualization

We can plot the same figure again in this case




Monte-Carlo vs TD-0: *Spectrum*

Temporal Difference



Monte-Carlo




TD-0 and Monte-Carlo are two extreme sides of a *spectrum*

- In TD-0, we use sample trajectory only for **one step**
- In Monte-Carlo, we use the **complete sample trajectory** for each state

Can we draw a solution between these two extreme points?

First Solution: Bootstrapping with More Steps

A primary approach to find such a **balanced** solution is to extend the idea of *bootstrapping* to a larger number of steps



- + How can we do it?
- Well! We could simply **expand** the **recursive property** of value function

n -Bootstrapping

Looking at a sample return, we can simply write

$$\begin{aligned}
 G_t &= R_{t+1} + \gamma G_{t+1} \\
 &= R_{t+1} + \gamma R_{t+2} + \gamma^2 G_{t+2} \\
 &\quad \vdots \\
 &= R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^n R_{t+1+n} + \gamma^{n+1} G_{t+1+n}
 \end{aligned}$$

Now, assuming that *at time step t we are at state s , i.e.,*

$$S_t = s, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1} \xrightarrow{R_{t+2}} \cdots \xrightarrow{R_{t+1+n}} S_{t+1+n}$$

we can bootstrap over a longer part of sample trajectory

n-Bootstrapping

$$S_t = s, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1} \xrightarrow{R_{t+2}} \dots \xrightarrow{R_{t+1+n}} S_{t+1+n}$$

In this trajectory, we can expand **Bellman equation** with **deeper recursion**

$$v_{\pi}(s) = \sum_{i=0}^n \gamma^i \mathbb{E} \{R_{t+i+1} | s\} + \gamma^{n+1} \mathbb{E}_{\pi} \{v_{\pi}(S_{t+n+1}) | s\}$$

So, if we have K sample trajectory, we can estimate value of **state s** by **n steps of bootstrapping**, i.e.,

$$\hat{v}_{\pi}(s) = \frac{1}{K} \sum_{k=1}^K \underbrace{\left(\sum_{i=0}^n \gamma^i R_{t+i+1}[k] + \hat{v}_{\pi}(S_{t+n+1}[k]) \right)}_{\text{computed on sample } k}$$

n-Bootstrapping \equiv TD-n

Let's formulate this approach: for a given *sample trajectory* and value function estimator $\hat{v}_\pi(\cdot)$, we define *n-bootstrapping return* at time t as

$$G_t^n = \sum_{i=0}^n \gamma^i R_{t+i+1} + \gamma^{n+1} \hat{v}_\pi(S_{t+n+1})$$

Given that the *sample trajectory* was started at *state* $S_t = s$, we can use *online averaging* and update the value estimator of *state* s as


$$\hat{v}_\pi(s) \leftarrow \hat{v}_\pi(s) + \alpha(G_t^n - \hat{v}_\pi(s))$$

This is what we call TD-n method of learning

This is obviously *more general* than TD-0! In TD-0, we had simply $n = 0$


Backup Diagram: *Sampling with n -Bootstrapping*

With *n -bootstrapping* we sample $(n + 1)$ -step trajectories from action-state tree of *environment*




Backup Diagram: Sampling with n -Bootstrapping

With n -bootstrapping we sample $(n + 1)$ -step trajectories from action-state tree of environment



Backup Diagram: Sampling with n -Bootstrapping

With *n*-bootstrapping we sample $(n + 1)$ -step trajectories from action-state tree of environment



TD- n : Policy Evaluation

We can extend our TD-0 evaluation algorithm to TD- n

$\text{TDn_Eval}(\pi)$:

- 1: Initiate estimator of value as $\hat{v}_\pi(s) = 0$ for all states
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a random state S_0 and act via policy $\pi(a|s)$
- 4: Sample a trajectory until either a terminal state or some terminating T

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - n - 1$ **do**
- 6: Compute $G = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n+1} v_\pi(S_{t+n+1})$ ★
- 7: Update as $\hat{v}_\pi(S_t) \leftarrow \hat{v}_\pi(S_t) + \alpha (G - \hat{v}_\pi(S_t))$
- 8: **end for**
- 9: **end for**

TD- n : Policy Q-Evaluation

Same-wise, we can extend our algorithm for action-value computation

TD n _QEval(π):

- 1: Initiate estimator of value as $\hat{q}_\pi(s, a) = 0$ for all **states**
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a **random state-action pair** (S_0, A_0) and act via policy $\pi(a|s)$
- 4: Sample a trajectory until either a **terminal** state or some **terminating** T

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - n - 1$ **do**
- 6: Compute $G = R_{t+n+1} + \gamma R_{t+n+2} + \dots + \gamma^{n+1} v_\pi(S_{t+n+1})$ ★
- 7: Update as $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha (G - \hat{q}_\pi(S_t, A_t))$
- 8: **end for**
- 9: **end for**

TD- ∞ : Going Back to Monte-Carlo

- + You told us that we look for a solution between TD-0 and Monte-Carlo! I see TD-0 is TD- n with $n = 0$, but where does Monte-Carlo stand?
- Well! You may see already that Monte-Carlo is TD- ∞

Say the environment is episodic: we can say that if we bootstrap *very deep*; then, at some point we hit a *terminal state*, i.e.,

$$\lim_{n \rightarrow \infty} v_{\pi}(S_{t+n+1}) = 0$$

This concludes that at any time t in an episodic environment

$$G_t^{\infty} = \sum_{i=0}^{\infty} \gamma^i R_{t+i+1} = G_t$$

and hence TD- ∞ will update its estimator by


$$\hat{v}_{\pi}(s) \leftarrow \hat{v}_{\pi}(s) + \alpha(G_t - \hat{v}_{\pi}(s))$$

TD- n : A Discrete Spectrum

We can look at TD- n as a *discrete* spectrum between Monte-Carlo and TD-0




Example: *Dummy Grid World with Random Walk*



Once you got home, try to get back to our *dummy world* and use *TD- n method* for *multiple n* to compute the values for *uniform random policy* 😊

Typical Behavior: Variation Against Depth

If you do some practice with random walk, you will see following curves for different choices of n



We observe a **minimum** against n : this is a **typical** behavior! Any illustration?

Averaging Different Depth

- + How **deep** we should then bootstrap?
- Well! We could try to find the **best** one **for each setting**, or we could **average** the result of **multiple bootstrapping depths**

For example, we can

```
1: Initiate ...
2: for episode = 1 : K do
3:   ...
4:   for t = 0 : T - 3 do
5:     Compute  $G_0 = R_{t+1} + \gamma v_{\pi}(S_{t+1})$ 
6:     Compute  $G_2 = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 v_{\pi}(S_{t+3})$ 
7:     Set  $G = (G_0 + G_2) / 2$ 
8:     Update as  $\hat{v}_{\pi}(S_t) \leftarrow \hat{v}_{\pi}(S_t) + \alpha (G - \hat{v}_{\pi}(S_t))$ 
9:   end for
10: end for
```

Averaging Different Depth: λ -Return


- + Why should that be a *good idea*?
- Because if the *good* one is within the average; then, it could *dominate* and *improve* estimation
- + Then, *which ones* we should take? We still have no clue!
- Let's take *all of them*! We can do it by *geometric* weights!

λ -Return

For any $0 \leq \lambda \leq 1$, the λ -return at time t over L steps is defined as

$$G_t^{\lambda} = (1 - \lambda) \sum_{n=0}^{L-1} \lambda^n G_t^n + \lambda^L G_t^L$$

Averaging Different Depth: λ -Return



For each state in the *sample trajectory*, we can

- Compute 0-bootstrapping return
↳ Give it weight $1 - \lambda$
- Compute 1-bootstrapping return
↳ Give it weight $(1 - \lambda) \lambda$
- \vdots
- Compute $(L - 1)$ -bootstrapping return
↳ Give it weight $(1 - \lambda) \lambda^{L-1}$
- Compute L -bootstrapping return
↳ Give it weight λ^L

TD $_{\lambda}$: Policy Evaluation

We can evaluate policy by averaging its λ -returns over multiple episodes

TD_Eval $_{\lambda}(\pi)$:

- 1: Initiate estimator of value as $\hat{v}_{\pi}(s^n) = 0$ for $n = 1 : N$
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a **random state** S_0 and act via policy $\pi(a|s)$
- 4: Sample a trajectory until either a **terminal** state or some **terminating** T

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - 1$ **do**
- 6: Set $G \leftarrow G_t^{\lambda}$ towards end of trajectory
- 7: Update as $\hat{v}_{\pi}(S_t) \leftarrow \hat{v}_{\pi}(S_t) + \alpha(G - \hat{v}_{\pi}(S_t))$
- 8: **end for**
- 9: **end for**

TD- λ : Special Cases

It is easy to see that $\lambda = 0$ and $\lambda = 1$ are again two extreme cases: *assume we are dealing with an episodic environment*

- with $\lambda = 0$

↳ All weights are zero but that of G_t^0 which is weighted one

$$G_t^{\lambda} = R_{t+1} + \gamma G_{t+1}$$

↳ this is basic bootstrapping: TD $_0$ is hence simply TD-0

- with $\lambda = 1$


↳ All weights are zero but that of G_t^{T-t} which is weighted one

$$G_t^{\lambda} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-t-1} R_T$$

↳ this basic Monte-Carlo: TD $_1$ is Monte-Carlo


TD $_{\lambda}$: A Continuous Spectrum

We can look at TD $_{\lambda}$ as a *continuous* spectrum between Monte-Carlo and TD-0



Typical Behavior: Variation Against λ


If you do some practice with random walk, you will see following curves for different choices of λ



We observe a **minimum** against λ : this is **analog** to TD- n behavior!

TD_λ : Weighting Function

Let's look at the weights decay in λ -return²



²From Sutton and Barto's book in Chapter 12

Assigning Credit for Future by Weighting

- + What is the **intuition** behind computing the λ -return with **those weights**?
- This is a very **valid question!** Let's see!

Let's look back at TD_λ approach: at each time t in the trajectory, we compute

$$G_t^\lambda = \sum_{n=1}^{T-t-1} w_t^n G_t^n$$


and then update the value of the **current state** S_t based on these **future returns**

the more **far away** this **future** is $\leftarrow \rightsquigarrow$ the **less** weight its return gets

In fact, we are **assigning credit** to our **current state** which will impact our **future** update: the more we go **forward** in time, the **less** this impact will be

TD_λ : Forward View

We can imagine this time advancement as *building impact* towards *future*³



But this approach is hard to be implemented *online*

we need to wait till *the end of episode* to compute the λ -returns!

³This figure is taken from Sutton and Barto's book in Chapter 12

TD_λ : Backward View

- + *But, it does not seem to be another way for credit assignment?*
- Well maybe we could apply the same idea backward
- + *How can we do it?*
- We can invoke the idea of eligibility tracing

Eligibility Tracing in Nutshell

At each time, when we update the value of a state in a sample trajectory, we also update the value of previous states we already met, with a weight decaying as we go back in time

Let's make an algorithm for that!

Eligibility Tracing

$\text{ElgTrace}(S_t, E(\cdot)) :$

- 1: Eligibility tracing function has N components, i.e., $E(s)$ for all **states**
- 2: **for** $n = 1 : N$ **do**
- 3: Update $E(s^n) \leftarrow \gamma\lambda E(s^n) \leftarrow$ choosing $\gamma\lambda$ for equivalency to forward view
- 4: **end for**
- 5: Update $E(S_t) \leftarrow E(S_t) + 1$

Say we initiate $E(s) = 0$ for **all states** and get to the following **trajectory**

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

Assume that we set $\gamma\lambda = 0.1$; then we have

- ① At $t = 0$, we only change $E(S_0) = 1 \leftarrow$ fresh memory \equiv high impact
- ② At $t = 1$, we change $E(S_0) = 0.1$ and $E(S_1) = 1$

Eligibility Tracing

$\text{ElgTrace}(S_t, E(\cdot)) :$

- 1: Eligibility tracing function has N components, i.e., $E(s)$ for all **states**
- 2: **for** $n = 1 : N$ **do**
- 3: Update $E(s^n) \leftarrow \gamma\lambda E(s^n) \leftarrow$ choosing $\gamma\lambda$ for equivalency to forward view
- 4: **end for**
- 5: Update $E(S_t) \leftarrow E(S_t) + 1$

Say we initiate $E(s) = 0$ for **all states** and get to the following **trajectory**

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

Now say that we see at $t = 2$ the same state as in $t = 0$, i.e., $S_0 = S_2$

- ③ At $t = 2$, we change $E(S_1) = 0.1$ and $E(S_0) \leftarrow 0.1E(S_0) + 1 = 1.01$

Updating with Eligibility Tracing: *Intuition*

- + What is the *use* of this algorithm?
- We can simply update *all previous states* each time t weighted by *eligibility traces*

We can 0-bootstrap at each time, i.e., compute *error* as

$$\Delta_t = \underbrace{R_{t+1} + \gamma \hat{v}_\pi(S_{t+1})}_{G_t^0} - \hat{v}_\pi(S_t)$$

and update *any state* $s = S_0, \dots, S_t$ that has non-zero trace of eligibility as

$$\hat{v}_\pi(s) \leftarrow \hat{v}_\pi(s) + \alpha \Delta_t E_t(s)$$

with $E_t(s)$ denoting the *eligibility trace* that we have updated up to time t

TD $_{\lambda}$ vs Eligibility Tracing

- + Is there any *concrete reason* beside *simple intuition* that this is a good idea?
- We can actually see that this is an online form of basic TD $_{\lambda}$

In fact, we can show by telescopic sum that

$$\sum_{t=0}^{T-1} \Delta_t E_t(s) = \sum_{t=0}^{T-1} \left(G_t^{\lambda} - \hat{v}_{\pi}(S_t) \right) \mathbf{1}\{S_t = s\}$$

This means that at the end of episode, we are updating the same!

- If we set $\lambda = 0$ in the *eligibility tracing*
 - ↳ all *eligibility trace* remains 0 for all states: only we have $E(S_t) = 1$
 - ↳ we are back to TD-0 as it was with TD $_0$
- If we set $\lambda = 1$ in the *eligibility tracing*
 - ↳ we are back to *Monte-Carlo* approach as in TD $_1$

TD $_{\lambda}$ with Eligibility Tracing: Policy Evaluation

We can evaluate policy by averaging its λ -returns over multiple episodes

ElgTD_Eval $_{\lambda}(\pi)$:


- 1: Initiate value estimator and eligibility traces as $\hat{v}_{\pi}(s) = 0$ and $E(s) = 0$ for all s
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a **random state** S_0 and act via policy $\pi(a|s)$
- 4: Sample a trajectory until either a **terminal** state or some **terminating** T

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - 1$ **do**
- 6: $E(\cdot) \leftarrow \text{ElgTrace}(S_t, E(\cdot))$
- 7: Compute $G = R_{t+1} + \gamma \hat{v}_{\pi}(S_{t+1})$ and find error $\Delta = G - \hat{v}_{\pi}(S_t)$
- 8: **for** $n = 1 : N$ **do**
- 9: Update $\hat{v}_{\pi}(s^n) \leftarrow \hat{v}_{\pi} + \alpha E(s^n) \Delta$
- 10: **end for**
- 11: **end for**
- 12: **end for**

Eligibility Tracing: *Backward View*

We can imagine eligibility tracing as *backward assignment* of *credits*⁴



Now we can *update values* each time and don't need to wait till *end of episode*!

⁴This figure is taken from Sutton and Barto's book in Chapter 12

Control versus Prediction

What we have done up to now is *prediction*

Prediction

*We are given by a *policy* and intend to evaluate it by *sampling**

But, in most applications we deal with a *control* problem

Control

*We are looking to move towards *optimal policy* while we are *sampling**

- + *But, we already talked about *GPI with sampling*! Didn't we?*
- *Well! What we did makes sense if we learn *offline*!*

GPI in Control Loop: Offline vs Online Approach

In *offline* RL, we *sample* environment first and find the *optimal policy* later

- ① *Sample environment* with *sufficient* number of episodes
- ② *Evaluate policies and improve them from the available dataset*
 - ↳ Go over the *dataset* over and over if needed

We are however looking for an *online* RL approach, we *sample* environment and learn *optimal policy* gradually as we sample

- ① Take a single *sample* from *environment*, e.g., a single reward-state pair or a terminating trajectory
 - ② Estimate values and improve policy based on this single sample
 - ③ Take a new sample . . .
- + I am sure that we *always* thought about the *second case*! Isn't that right?!
- Yes! But, if we want to do *complete* evaluation in each GPI iteration, we will be *extremely slow*! Imagine *1000 episodes* for each iteration!

Direct GPI with Prediction

Recall a generic form of GPI with a prediction approach X

$X_PolicyItr()$:

1: Initiate two random policies π and $\bar{\pi}$

2: **while** $\pi \neq \bar{\pi}$ **do**

3: $\hat{q}_\pi = X_QEval(\pi)$ and $\pi \leftarrow \bar{\pi}$

4: $\bar{\pi} = Greedy(\hat{q}_\pi)$

5: **end while**

and $Greedy(\hat{q}_\pi)$ is the basic improvement strategy

$Greedy(\hat{q}_\pi)$:

1: **for** $n = 1 : N$ **do**

2: Improve the by taking **deterministically** the **best action**

$$\bar{\pi}(a^m | s^n) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} \hat{q}_\pi(s^n, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} \hat{q}_\pi(s^n, a^m) \end{cases}$$

3: **end for**

Direct GPI with Prediction

Recall a generic form of GPI with a prediction approach X

```

X_PolicyItr():
1: Initiate two random policies  $\pi$  and  $\bar{\pi}$ 
2: while  $\pi \neq \bar{\pi}$  do
3:    $\hat{q}_\pi = X\_QEval(\pi)$  and  $\pi \leftarrow \bar{\pi}$ 
4:    $\bar{\pi} = Greedy(\hat{q}_\pi)$ 
5: end while
  
```

If we want our evaluation to be **accurate enough**: we need to keep playing **each policy** for a **large** number of episodes

- This is **extremely sample inefficient**
 - ↳ Imagine how many times we should lose the game to update our strategy!
 - ↳ We as human do not really need so many losses
- In many practical settings is really **cost-inefficient**
 - ↳ Amount of losses we should pay in each iteration to evaluate the policy is not worth it!

Online Control Loop via GPI

- + But, how can we do anything about this?
- Maybe we could improve the policy after *each update* of *action-values*

```
X_Control():
```

```
1: Initiate two random policies  $\pi$  and  $\bar{\pi}$ 
```

```
2: while  $\pi \neq \bar{\pi}$  do
```

```
3:    $\hat{q}_\pi = \text{X\_QUpdate}(\pi)$  and  $\pi \leftarrow \bar{\pi}$ 
```

```
4:    $\bar{\pi} = \text{Greedy}(\hat{q}_\pi)$ 
```

```
5: end while
```

Here, $\text{X_QUpdate}(\pi)$ refers to one single update which is typically of the form

$$\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(G - \hat{q}_\pi(S_t, A_t))$$

for some G

Online Control Loop via GPI

`X_Control():`

1: *Initiate two random policies π and $\bar{\pi}$*

2: **while** $\pi \neq \bar{\pi}$ **do**

3: $\hat{q}_\pi = \text{X_QUpdate}(\pi)$ and $\pi \leftarrow \bar{\pi}$

4: $\bar{\pi} = \text{Greedy}(\hat{q}_\pi)$

5: **end while**

- + *It sounds like a loose approach! Why should that work?!*
- We see accurate illustrations about that, *but for the moment*
 - ↳ we could think of a single update as a *low-accuracy estimate*
 - ↳ we have already said the GPI is very *robust* against *estimation error*

Let's start by a Monte-Carlo control loop

First Try: Monte-Carlo Control Loop

We can build a Monte-Carlo control loop for **episodic environments**

MC_Control(π):

1: Initiate estimator as $\hat{q}_\pi(s, a) = 0$ for **all states** and **actions**

2: **for** episode = 1 : K or until π stops changing **do**

3: Initiate with a **random state-action pair** (S_0, A_0)

4: Act via $\pi = \text{Greedy}(\hat{q}_\pi)$

5: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminal}$$

6: Initiate with $G = 0$

7: **for** $t = T - 1 : 0$ **do**

8: Update current return $G \leftarrow R_{t+1} + \gamma G$

9: Update $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(G - \hat{q}_\pi(S_t, A_t))$

10: **end for**

11: **end for**

Monte-Carlo Control: *Updating Action-Values*


Comparing with $\text{MC_QEval}(\pi)$, there is only one difference, i.e.,

in-loop greedy improvement of the policy: line 4

- 1 Estimate *action-values* over a *sample trajectory*
- 2 *Improve* the policy using this estimate by *greedy approach*
- 3 *Sample* the next trajectory using the *improved policy*

-
- + Can we guarantee the convergence of this control loop?
 - In the current state, **not really!** Let's see an example!


Example: Our Multi-armed Bandit



Let's get back to our very first RL problem in which the robot is to decide for a company: say the robot follows Monte-Carlo control loop

- ① it starts with a random decision
 - ↳ Say it decides for Company B and receives **\$150 income**
 - ↳ Now, we have $\hat{q}_\pi(_, A) = 0$ and $\hat{q}_\pi(_, B) = 150$
- ② in the next episode, it would definitely chooses to work at Company B
 - ↳ Say it receives **\$250 income** this time
 - ↳ Now, we have $\hat{q}_\pi(_, A) = 0$ and $\hat{q}_\pi(_, B) = 200$
 - ⋮

Example: Our Multi-armed Bandit



The robot keeps working at **Company B!**

- + But, can we **guarantee** that company A is not paying better?!
- Well! **Not really!** In fact, even if we had worked there for a single day or so, we could still not guarantee!

Greedy Improvement: *Lack of Exploration*

- + *Why is this happening? Why it doesn't happen when we apply direct GPI via Monte-Carlo?*
- In the latter, we do *exploration*; but now, we are only *exploiting*!

*This is a general behavior of **greedy improvement***

Downside of Greedy Improvement

*In **greedy** improvement, we only **exploit** our knowledge, i.e.,*

*we always act **optimal** based on what we know **up to now***

*We thus **lack exploration**, i.e.,*

we remain unaware about states and actions that we have not explored

*We may **never** get the chance to explore them!*

Improving via ϵ -Greedy Improvement

A classical approach to handle this issue is to **improve** by ϵ -greedy approach

ϵ -Greedy Improvement

Choose a small $0 < \epsilon < 1$, and **improve** after each update of **action-values** by greedy approach: at beginning of each episode

- with probability $1 - \epsilon$ act by the **improved** policy
- with probability ϵ act randomly

We implemented this approach for multi-armed bandit in [Assignment 1](#): in this approach we render a trade-off between exploitation and exploration

- with probability $1 - \epsilon$ we **exploit** our **improved** policy
- with probability ϵ we **explore** the environment

It's hard to find any improvement approach that can beat ϵ -greedy!

ϵ -Greedy Algorithm

We can algorithmically specify ϵ -greedy as

ϵ -Greedy(\hat{q}_π):

1: **for** $n = 1 : N$ **do**

2: Take next step *randomly* as

$$\bar{\pi}(a^m | s^n) = \begin{cases} 1 - \epsilon + \frac{\epsilon}{M} & m = \underset{m}{\operatorname{argmax}} \hat{q}_\pi(s^n, a^m) \\ \frac{\epsilon}{M} & m \neq \underset{m}{\operatorname{argmax}} \hat{q}_\pi(s^n, a^m) \end{cases}$$

3: **end for**

- + That seems to solve exploration problem! But, is there any *guarantee* that $\bar{\pi}$ is going to be a *better* policy? For greedy approach, we could prove that we get always *better*!
- Yes! We can actually prove it!

ϵ -Greedy Algorithm

Let's assume we have policy π given after ϵ -greedy improvement, and we improved it again via the ϵ -greedy approach from its action-values: we can then write the value of new policy $\bar{\pi}$ as

$$\begin{aligned}
 v_{\bar{\pi}}(s) &= \sum_{m=1}^M \bar{\pi}(a^m|s) q_{\bar{\pi}}(s, a^m) \\
 &= \underbrace{\frac{\epsilon}{M} \sum_{m=1}^M q_{\pi}(s, a^m)}_{\text{exploration}} + \underbrace{(1 - \epsilon) q_{\pi}(s, a^*)}_{\text{exploitation}}
 \end{aligned}$$

We know that for any non-negative w_1, \dots, w_M that add up to one, we have

$$\sum_{m=1}^M w_m q_{\pi}(s, a^m) \leq q_{\pi}(s, a^*)$$

ϵ -Greedy Algorithm

We have the improved value in terms of the initial action-values as

$$v_{\bar{\pi}}(s) = \frac{\epsilon}{M} \sum_{m=1}^M q_{\pi}(s, a^m) + (1 - \epsilon) q_{\pi}(s, a^*)$$

Let's now define

$$w_m = \frac{\pi(a^m|s) - \epsilon/M}{1 - \epsilon}$$

We note that since π is an ϵ -greedy policy, we have $w_m \geq 0$ and

$$\sum_{m=1}^M w_m = \sum_{m=1}^M \frac{\pi(a^m|s) - \epsilon/M}{1 - \epsilon} = 1$$

ϵ -Greedy Algorithm

Now, let us replace this bound in the previous equation

$$\begin{aligned}
 v_{\bar{\pi}}(s) &= \frac{\epsilon}{M} \sum_{m=1}^M q_{\pi}(s, a^m) + (1 - \epsilon) q_{\pi}(s, a^*) \\
 &\geq \frac{\epsilon}{M} \sum_{m=1}^M q_{\pi}(s, a^m) + (1 - \epsilon) \sum_{m=1}^M \frac{\pi(a^m|s) - \epsilon/M}{1 - \epsilon} q_{\pi}(s, a^m) \\
 &= \sum_{m=1}^M \pi(a^m|s) q_{\pi}(s, a^m) = v_{\pi}(s)
 \end{aligned}$$

ϵ -Greedy Improvement Theorem

Let π and $\bar{\pi}$ be ϵ -greedy policies, i.e., computed from some action-value function using ϵ -greedy algorithm. Assume $\bar{\pi}$ is derived by ϵ -greedy improvement from $q_{\pi}(s, a)$; then, $\bar{\pi} \geq \pi$

Online Control Loop via GPI and ϵ -Greedy Improvement

We can now build our control loop via ϵ -greedy algorithm

```
X_Control():
```

```
1: Initiate two random policies  $\pi$  and  $\bar{\pi}$ 
```

```
2: while  $\pi \neq \bar{\pi}$  do
```

```
3:    $\hat{q}_\pi = \text{X\_QUpdate}(\pi)$  and  $\pi \leftarrow \bar{\pi}$ 
```

```
4:    $\bar{\pi} = \epsilon\text{-Greedy}(\hat{q}_\pi)$ 
```

```
5: end while
```

Attention

We are still using *single update* of action-values for *policy improvement*: this means that we may have *bad* estimates of action-values at initial iterations!

First Try: Monte-Carlo Control Loop

Monte-Carlo control loop for **episodic environments** is modified as

MC_Control(π):

1: Initiate estimator as $\hat{q}_\pi(s, a) = 0$ for **all states** and **actions**

2: **for** episode = 1 : K or until π stops changing **do**

3: Initiate with a **random state-action pair** (S_0, A_0)

4: Act via $\pi = \epsilon$ -Greedy(\hat{q}_π)

5: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminal}$$

6: Initiate with $G = 0$

7: **for** $t = T - 1 : 0$ **do**


8: Update current return $G \leftarrow R_{t+1} + \gamma G$

9: Update $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(G - \hat{q}_\pi(S_t, A_t))$

10: **end for**

11: **end for**

Example: Cliff Walking




We have seen the cliff walking example in Assignment 1: we want to

- get from S to G with **shortest possible path**
- avoid hitting the cliff \equiv gray squares
 - ↳ each time we hit the cliff, we get back to S with a **big negative reward**

Example: Cliff Walking


Say we use naive greedy policy: we **start** sampling trajectory and hit the cliff



We realize that our first action gave **bad** reward

Example: Cliff Walking


We now follow a *better* action, but we hit the cliff again



We realize that this action was even *worse*

Example: Cliff Walking

We get back to our first action, but now modify *next actions*



Say we are lucky and arrive at G

We will *never* go back to find the *optimal path*!

But with ϵ -greedy improvement, we get the chance to *explore* again: we may find the *optimal path*!

Greedy in Limit with Infinite Exploration Algorithms

- + *Sounds working! But can we guarantee that this approach will converge to optimal path?*
- Under some circumstances: Yes!

Recall that we said the following when we started Monte-Carlo

Asymptotic Convergence of Monte-Carlo

Let $\mathcal{C}_K(s, a)$ denote number of visits at *state* s followed by *action* a during K Monte-Carlo episodes. Assume random initialization is distributed such that

$$\lim_{K \rightarrow \infty} \mathcal{C}_K(s, a) = \infty$$

for any *state* s and *action* a ; then, we can guarantee $\hat{q}_\pi(s, a) \xrightarrow{K \uparrow \infty} q_\pi(s, a)$

Greedy in Limit with Infinite Exploration Algorithms

The main idea in this result was that

As long as we do *enough sampling*, so that we see *all states* and *actions* enough number of times, Monte-Carlo will converge

We can claim the same thing here

If we keep *playing enough*, we explore *all states* and *actions*; then, eventually we get *very sure* about *optimal values* and *actions*

But there is a small point here: if we keep on using ϵ -greedy policy even after we *got sure*, we can still perform sub-optimal

We should *stop exploring* once we have visited *all states* and *actions*

This is what we call

Greedy in Limit with Infinite Exploration \equiv GLIE

GLIE Algorithms

GLIE Algorithms

A GPI-type control loop is GLIE, if for any **state-action** pair (s, a) , we have the following asymptotic properties

- 1 The number of visits to all **state-action** pair grows large

$$\lim_{K \rightarrow \infty} C_K(s, a) = \infty$$

- 2 The **improved** policy in last episode converges to **greedy** policy

$$\lim_{K \rightarrow \infty} \pi_K(a^m | s) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_{\pi_K}(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_{\pi_K}(s, a^m) \end{cases}$$

GLIE control algorithms converge to **optimal policy**

GLIE Algorithms

- + It seems that they **contradict!** First one needs us to **explore** and the second to **exploit!**
- We could simply get rid of it by scaling ϵ

Say we choose ϵ to scale reversely by the number of episodes, e.g.,

$$\epsilon_k = \frac{1}{k}$$

Then, we have both the constraints satisfied

- 1 We keep exploring a lot in initial episodes
- 2 We focus more on exploiting in later episodes

This is what we do in practice!

ϵ -Greedy Monte-Carlo is GLIE

It is easy to show that **Monte-Carlo** with shrinking ϵ -greedy improvement is GLIE

MC_Control():

1: Initiate estimator as $\hat{q}_\pi(s, a) = 0$ for **all states** and **actions**

2: **for** episode = 1 : K or until π stops changing **do**

3: Initiate with a **random state-action pair** (S_0, A_0)

4: Set $\epsilon = 1/k$ and act via $\pi = \epsilon$ -Greedy(\hat{q}_π)

5: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminal}$$

6: Initiate with $G = 0$

7: **for** $t = T - 1 : 0$ **do**

8: Update current return $G \leftarrow R_{t+1} + \gamma G$

9: Update $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(G - \hat{q}_\pi(S_t, A_t))$

10: **end for**

11: **end for**

Control Loop via Temporal Difference

- + But still we are not fully **online**! We need to wait till **end of each episode**!
- Well! That's right! But, we could use TD!

Using TD in the control loop will make our algorithm **fully online**

- We update values after each **state-action** pair
- We then **improve** the policy

We should yet use **ϵ -greedy improvement** to keep exploration

SARSA: State-Action-Reward State-Action

SARSA \equiv State-Action Reward State-Action

SARSA algorithms use TD along with ϵ -greedy update for the control loop

In general, we can develop various forms of SARSA

- We may use TD-0 for updating *action-values*
 - ↳ This is the basic SARSA
- We may use TD- n for updating *action-values*
 - ↳ This is n -SARSA
- We may use TD- λ for updating *action-values*
 - ↳ This is SARSA(λ)

SARSA: First Try

Let's try to make a simple TD-based control loop

TD_Control():

- 1: Initiate estimator as $\hat{q}_\pi(s, a) = 0$ for *all states* and *actions*
- 2: **for** episode = 1 : K or until π stops changing **do**
- 3: Initiate with a *random state-action pair* (S_0, A_0)
- 4: **for** $t = 0 : T - 1$ that is either *terminal* or *terminated* **do**
- 5: Act A_t and observe

$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$
- 6: Update policy to $\pi \leftarrow \epsilon\text{-Greedy}(\hat{q}_\pi)$
- 7: Draw the new action A_{t+1} from $\pi(\cdot | S_{t+1})$
- 8: Compute $\hat{v}_\pi(S_{t+1})$ from $\hat{q}_\pi(S_{t+1}, a)$ and $\pi(\cdot | S_{t+1})$
- 9: Set $G \leftarrow R_{t+1} + \gamma \hat{v}_\pi(S_{t+1})$
- 10: Update $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(G - \hat{q}_\pi(S_t, A_t))$
- 11: **end for**
- 12: **end for**

SARSA: Going On-Policy

In line 8 of our control algorithm: we compute $\hat{v}_\pi(S_{t+1})$ as

$$\hat{v}_\pi(S_{t+1}) = \sum_{m=1}^M \pi(a^m | S_{t+1}) \hat{q}_\pi(S_{t+1}, a^m)$$

But, we do know that

- ① our estimates $\hat{q}_\pi(S_{t+1}, a^m)$ are **not** that good, and also
- ② our policy has led use to next action A_{t+1}

So, we could **move on our policy** and write

$$\pi(a | S_{t+1}) = \begin{cases} 1 & a = A_{t+1} \\ 0 & a \neq A_{t+1} \end{cases} \rightsquigarrow \hat{v}_\pi(S_{t+1}) = \hat{q}_\pi(S_{t+1}, A_{t+1})$$

We call this approach **on-policy**, since **move on our policy**

SARSA: Basic Algorithm

SARSA() :

- 1: Initiate estimator as $\hat{q}_\pi(s, a) = 0$ for *all states* and *actions*
- 2: **for** episode = 1 : K or until π stops changing **do**
- 3: Initiate with a *random state-action pair* (S_0, A_0)
- 4: **for** $t = 0 : T - 1$ that is either *terminal* or *terminated* **do**
- 5: Act A_t and observe

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$$

- 6: Update policy to $\pi \leftarrow \epsilon\text{-Greedy}(\hat{q}_\pi)$
- 7: Draw the new action A_{t+1} from $\pi(\cdot | S_{t+1})$ and move *on policy*

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1}$$

- 8: Set $G \leftarrow R_{t+1} + \gamma \hat{q}_\pi(S_{t+1}, A_{t+1})$
- 9: Update $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(G - \hat{q}_\pi(S_t, A_t))$
- 10: **end for**
- 11: **end for**

SARSA: Deeper Return Samples


We can use a longer trajectory while we learn on-policy, i.e.,

$$G^n = \sum_{i=0}^n R_{t+i+1} + \gamma \hat{q}_{\pi}(S_{t+n+1}, A_{t+n+1})$$

This will however add extra delay!

As a practice, you could

re-write the basic SARSA with *n*-return 😊



SARSA(λ): Tracing Eligibility of State-Action Pairs

We can extend SARSA to the case with λ -return: we have two options

- the case with forward-view
 - ↳ We know this is **not** practical! So, let's skip the details
- the case with backward-view and eligibility tracing
 - ↳ Let's look into this one

We first extend **eligibility tracing** to the case with **state-action pairs**

$\text{ElgTrace}(S_t, A_t, E(\cdot) \mid \lambda)$:

- 1: Eligibility tracing function has NM components, i.e., $E(s, a)$ for all **state-action** pairs
- 2: **for** all **state-action** pairs (s, a) **do**
- 3: Update $E(s, a) \leftarrow \gamma \lambda E(s, a)$
- 4: **end for**
- 5: Update $E(S_t, A_t) \leftarrow E(S_t, A_t) + 1$

SARSA: Alternative via TD- λ

SARSA(λ) :

- 1: Initiate $\hat{q}_\pi(s, a) = 0$ and $E(s, a) = 0$ for **all states** and **actions**
- 2: **for** episode = 1 : K or until π stops changing **do**
- 3: Initiate with a **random state-action pair** (S_0, A_0)
- 4: **for** $t = 0 : T - 1$ that is either **terminal** or **terminated** **do**
- 5: $E(\cdot) \leftarrow \text{ElgTrace}(S_t, A_t, E(\cdot) | \lambda)$
- 6: Act A_t and observe R_{t+1} and S_{t+1}
- 7: Update policy to $\pi \leftarrow \epsilon\text{-Greedy}(\hat{q}_\pi)$
- 8: Draw the new action A_{t+1} from $\pi(\cdot | S_{t+1})$ and move on policy

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1}$$

- 9: Set $\Delta \leftarrow R_{t+1} + \gamma \hat{q}_\pi(S_{t+1}, A_{t+1}) - \hat{q}_\pi(S_t, A_t)$
- 10: **for** all state-action pairs (s, a) **do**
- 11: Update $\hat{q}_\pi(s, a) \leftarrow \hat{q}_\pi(s, a) + \alpha \Delta E(s, a)$
- 12: **end for**
- 13: **end for**
- 14: **end for**

Going Off-Policy

Let's think about a fundamental question: *while sampling the environment with a specific policy π , can we estimate the values of another policy $\bar{\pi}$?*

- + Why should this be a *fundamental* question?
- Well! There are several reasons
 - ↳ Maybe we sampled *environment* with our *bad* policy: can't we use our sample again?
 - ↳ Maybe we are *looking at other players*: can't we learn something about the environment from their samples?
 - ↳ Maybe they are *good* players: can't we use this fact to improve our policy?
 - ↳ Maybe they are *bad* players: can't we use this fact to avoid doing mistakes?

This is the idea of *off-policy* control

Let's start with some basics

Importance Sampling

Consider following problem: we have random variable X drawn as $X \sim p(x)$ whose mean is

$$\mu_p = \mathbb{E}_p \{X\} = \sum_x p(x) x$$

We want to know how would be the expectation if we had $X \sim q(x)$: we write

$$\begin{aligned} \mu_q &= \mathbb{E}_q \{X\} = \sum_x q(x) x \\ &= \sum_x p(x) \frac{q(x)}{p(x)} x = \mathbb{E}_p \left\{ \frac{q(X)}{p(X)} X \right\} \end{aligned}$$

This gives us possibility to

estimate $\mathbb{E}_q \{X\}$ using samples drawn from $p(x)$

Importance Sampling

Say we have drawn K samples from $p(x)$, i.e., we have

$$X_1, X_2, \dots, X_K$$

We can use **Monte-Carlo** to estimate μ_p as

$$\hat{\mu}_p = \frac{1}{K} \sum_{k=1}^K X_k$$

We can **also** use **Monte-Carlo** to estimate μ_q as

$$\hat{\mu}_q = \frac{1}{K} \sum_{k=1}^K \frac{q(X_k)}{p(X_k)} X_k$$

We call this method **importance sampling**

Off-Policy Control via Importance Sampling

Now, let's get back to our problem: *assume we have played with policy π and collected K sample trajectories of length T all started at state $S_0 = s$, i.e.,*

$$s = S_0[k], A_0[k] \xrightarrow{R_1[k]} S_1[k], A_1[k] \xrightarrow{R_2[k]} \dots \xrightarrow{R_T[k]} S_T[k]$$

for $k = 1 : K$; then, we could write

$$\hat{v}_\pi(s) = \frac{1}{K} \sum_{k=1}^K G[k]$$

This is the basic *Monte-Carlo*

Off-Policy Control via Importance Sampling

But now, we want to use samples to evaluate *another policy* $\bar{\pi}$

$$s = S_0[k], A_0[k] \xrightarrow{R_1[k]} S_1[k], A_1[k] \xrightarrow{R_2[k]} \dots \xrightarrow{R_T[k]} S_T[k]$$

We could also use importance sampling to write

$$\begin{aligned} \hat{v}_{\bar{\pi}}(s) &= \frac{1}{K} \sum_{k=1}^K \frac{\Pr\{\text{same action sequence with } \bar{\pi}\}}{\Pr\{\text{same action sequence with } \pi\}} G[k] \\ &= \frac{1}{K} \sum_{k=1}^K \frac{\bar{\pi}(A_0[k]|S_0[k]) \cdots \bar{\pi}(A_{T-1}[k]|S_{T-1}[k])}{\pi(A_0[k]|S_0[k]) \cdots \pi(A_{T-1}[k]|S_{T-1}[k])} G[k] \\ &= \frac{1}{K} \sum_{k=1}^K \prod_{\ell=0}^{T-1} \frac{\bar{\pi}(A_{\ell}[k]|S_{\ell}[k])}{\pi(A_{\ell}[k]|S_{\ell}[k])} G[k] \end{aligned}$$

Off-Policy Control via Importance Sampling

We can further update the estimate in an *online fashion* from

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1} \xrightarrow{R_{t+2}} \dots \xrightarrow{R_T} S_T$$

by *online averaging* as

$$\hat{v}_{\bar{\pi}}(S_t) \leftarrow \hat{v}_{\bar{\pi}}(S_t) + \alpha \left(\prod_{\ell=t}^{T-1} \frac{\bar{\pi}(A_{\ell}|S_{\ell})}{\pi(A_{\ell}|S_{\ell})} G_t - \hat{v}_{\bar{\pi}}(S_t) \right)$$

So, we are evaluating $\bar{\pi}$ via *Monte-Carlo*

off our policy π

This is *off-policy control*

Off-Policy Control via Importance Sampling

We can further apply *off-policy control* via TD

$$\hat{v}_{\bar{\pi}}(S_t) \leftarrow \hat{v}_{\bar{\pi}}(S_t) + \alpha \left(\frac{\bar{\pi}(A_t|S_t)}{\pi(A_t|S_t)} (R_{t+1} + \gamma \hat{v}_{\bar{\pi}}(S_{t+1})) - \hat{v}_{\bar{\pi}}(S_t) \right)$$

Note that for action-values estimate

R_{t+1} does not depend any more on *policy* as we know *action* A_t

Therefore, we have for action-value update

$$\hat{q}_{\bar{\pi}}(S_t, A_t) \leftarrow \hat{q}_{\bar{\pi}}(S_t, A_t) + \alpha \left(R_{t+1} + \gamma \frac{\bar{\pi}(A_t|S_t)}{\pi(A_t|S_t)} \hat{v}_{\bar{\pi}}(S_{t+1}) - \hat{q}_{\bar{\pi}}(S_t, A_t) \right)$$

Q-Learning

Q-Learning

Q-learning is an **off-policy** TD **control** algorithm, where we sample with ϵ -greedy policy but update the action-values to evaluate greedy policy

This means in Q-learning π is ϵ -greedy policy and $\bar{\pi}$ is greedy. Let's consider basic TD evaluation: so, we can write

$$\hat{q}_{\bar{\pi}}(S_t, A_t) \leftarrow \hat{q}_{\bar{\pi}}(S_t, A_t) + \alpha(G - \hat{q}_{\bar{\pi}}(S_t, A_t))$$

where G should be

$$G = R_{t+1} + \gamma \hat{v}_{\bar{\pi}}(\bar{S}_{t+1})$$

Since we sample by ϵ -greedy policy π , we use **importance sampling** and write

$$G = R_{t+1} + \gamma \frac{\bar{\pi}(A_t | S_t)}{\pi(A_t | S_t)} \hat{v}_{\bar{\pi}}(S_{t+1})$$

Q-Learning

But, we really *don't need importance sampling*: we can simply observe that

$$\hat{v}_{\bar{\pi}}(S_{t+1}) = \sum_{m=1}^M \hat{q}_{\bar{\pi}}(S_{t+1}, a^m) \bar{\pi}(a^m | S_{t+1}) = \max_m \hat{q}_{\bar{\pi}}(S_{t+1}, a^m)$$

and we do know that

$$\frac{\bar{\pi}(A_t | S_t)}{\pi(A_t | S_t)} = \mathbf{1} \left\{ A_t = \operatorname{argmax}_a \hat{q}_{\bar{\pi}}(S_t, a) \right\}$$

So, we could directly update as

$$\hat{q}_{\bar{\pi}}(S_t, A_t) \leftarrow \hat{q}_{\bar{\pi}}(S_t, A_t) + \alpha \left(R_{t+1} + \gamma \max_m \hat{q}_{\bar{\pi}}(S_{t+1}, a^m) - \hat{q}_{\bar{\pi}}(S_t, A_t) \right)$$

This concludes *Q-learning algorithm*

Q-Learning: Basic Algorithm


Q-Learning():

- 1: Initiate estimator as $\hat{q}_\star(s, a) = 0$ for *all states* and *actions*
- 2: **for** episode = 1 : K or until π stops changing **do**
- 3: Initiate with a *random state* S_0
- 4: **for** $t = 0 : T - 1$ that is either *terminal* or *terminated* **do**
- 5: Update policy to $\pi \leftarrow \epsilon\text{-Greedy}(\hat{q}_\star)$
- 6: Draw action A_t from $\pi(\cdot | S_t)$ and observe

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$$


- 7: Set $G \leftarrow R_{t+1} + \gamma \max_m \hat{q}_\star(S_{t+1}, a^m)$
- 8: Update $\hat{q}_\star(S_t, A_t) \leftarrow \hat{q}_\star(S_t, A_t) + \alpha(G - \hat{q}_\star(S_t, A_t))$
- 9: **end for**
- 10: **end for**

Example: Cliff Walking



Let's compare SARSA to Q-Learning algorithm!

Example: Cliff Walking



Don't Mistake!

Q-learning collects less reward since it goes **off-policy**; however, it **estimates optimal** action-values: at some point it can start playing **optimally**

Convergence of SARSA

Recall: GLIE Algorithms

A GPI-type control loop is GLIE, if for any **state-action** pair (s, a) , we have the following asymptotic properties

- 1 The number of visits to all **state-action** pair grows large

$$\lim_{K \rightarrow \infty} C_K(s, a) = \infty$$

- 2 The **improved** policy in last episode converges to **greedy** policy

$$\lim_{K \rightarrow \infty} \pi_K(a^m | s) = \begin{cases} 1 & m = \operatorname{argmax}_m q_{\pi_K}(s, a^m) \\ 0 & m \neq \operatorname{argmax}_m q_{\pi_K}(s, a^m) \end{cases}$$

GLIE control algorithms converge to **optimal policy**

Convergence of SARSA

- + *But do we really have large number of episodes with SARSA?*
- Not necessarily! We may have only one **infinitely long** trajectory
- + *What should we do then?*
- We can simply treat it as a large number of episodes of length 1

In (basic) SARSA, we only need one step in the trajectory

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$$

We could hence think of it as one episode

- ① *each time step t we update the **action-values***
- ② *each time step we **improve the policy***

Convergence of SARSA

Modification: GLIE Algorithms

An online control loop is GLIE, if we have asymptotically in time t

- 1 The number of visits to all **state-action** pair grows large

$$\lim_{t \rightarrow \infty} \mathcal{C}_t(s, a) = \infty$$

- 2 The **improved** policy converges to **greedy** policy

$$\lim_{t \rightarrow \infty} \pi_t(a^m | s) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_{\pi_t}(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_{\pi_t}(s, a^m) \end{cases}$$

Convergence of SARSA: Make it GLIE

- + Can we guarantee that *both conditions* hold with SARSA?
- The *second one* is easy: we need to *scale ϵ down with t* , e.g., $\epsilon_t = 1/t$
- + What about the *first condition*?
- We should scale *the step-size α according to Robbins-Monro*

Robbins-Monro Sequence

Sequence α_t is Robbins-Monro if we have

$$\sum_{t=0}^{\infty} \alpha_t = \infty \quad \text{and} \quad \sum_{t=0}^{\infty} \alpha_t^2 < \infty$$

For instance, $\alpha_t = 1/t$ is a *Robbins-Monro sequence*

Convergence of SARSA

Convergence of SARSA

SARSA online control loop converges to the optimal action-values if

- 1 Step-size is scheduled by a *Robbins-Monro sequence*
- 2 Exploration factor ϵ *decays in time*

In practice however

- ϵ is a *hyperparameter*
 - ↳ We know that we should *schedule* it
 - ↳ How we should do the *scheduling*? This is *hyperparameter tuning*
- α is a *hyperparameter*: some people call it *learning rate*
 - ↳ Its *scheduling* is again *hyperparameter tuning*

Convergence of Q-Learning

Convergence of Q-Learning

Q-learning online control loop with exploration (non-zero ϵ) converges to the optimal action-values as $t \rightarrow \infty$

- + That's it?
- Yes!

Since we are evaluating *off-policy*, we *don't care* about *behaving policy*

Q-Learning vs SARSA

- + So! Does it mean that Q-learning is always better?
- Not always!

In general *Q-learning* has several benefits

- *Minimal* convergence requirements
- It converges faster to the *optimal policy*
 - ↳ If we want to make SARSA that fast, we may get to a *sub-optimal policy*
- It has more *flexibility* and *sample-efficiency*

But, *SARSA* also has some benefits

- It is better suited for *online control*
 - ↳ Our *behaving policy* is the one going towards optimal one
 - ↳ In Q-learning, the *behaving policy* is not the *optimal one*
- It has lower complexity
 - ↳ We just deal with one policy

End of Story!

Model-Based RL

Bellman Equation

value iteration

policy iteration

Model-free RL

on-policy methods

temporal difference

Monte Carlo

SARSA

off-policy methods

Q-learning

*I would strongly suggest to start with **programming part** of **Assignment 2!***

*There you solve Frozen Lake with **SARSA** and **Q-Learning***