

ECE 1508: Reinforcement Learning

Chapter 6: Actor Critic Methods

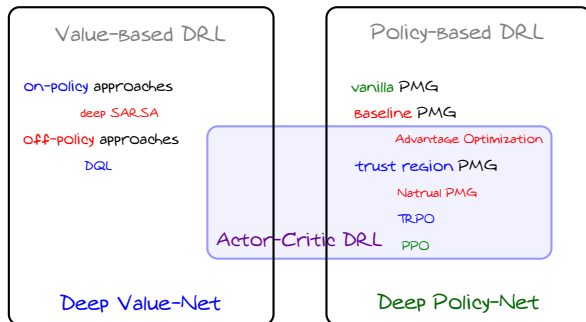
Ali Bereyhi

`ali.bereyhi@utoronto.ca`

Department of Electrical and Computer Engineering
University of Toronto

Summer 2024

Deep RL: Sort of Division



Deep RL: Sort of Division

In *actor-critic approaches* we have both networks

- an *actor* has a *policy network*
 - ↳ This network enables it to *act* at each *particular state*
- a *critic* has a *value network*
 - ↳ This network enables it to *evaluate* its *policy*
 - ↳ The *evaluation* will help *improving* the *policy*



Deep RL: Sort of Division

Attention

For many people *actor-critic* \equiv *PGM*: they usually argue that

- to implement a *PGM* we need to *estimate values*
- we should do it by a *value network*

So, any *PGM* is at the end *actor-critic*

That's practically true; however, in principle, we can

implement *PGMs* via basic *Monte Carlo*

So, we could also have a *pure PGM*, e.g., REINFORCE!

Implementing PGMs

Let's get back to **PGMs**: say we want to implement a **PGM**

- We usually use **sample advantages**, i.e.,

$$U_t = R_{t+1} + \gamma v_{\pi_{\theta}}(S_{t+1}) - v_{\pi_{\theta}}(S_t)$$

So, we need to know the value function $v_{\pi_{\theta}}(\cdot)$ of our **policy** π_{θ}

- + Well, why don't we **evaluate** it once and use it forever?
- **Attention!** We need this **evaluation each time** we update **policy** π_{θ} !
- + How exactly we do it then? You promised to tell us!
- Sure! Let's use what we have learned up to now

Advantage PGM: Implementing

Let's look at the classic *advantage optimization* PGM

AdvantagePGM():

```

1: Initiate with  $\theta$  and learning rate  $\alpha$ 
2: while interacting do
3:   Set  $\hat{\nabla} = 0$ 
4:   for mini-batch  $b = 1 : B$  do
5:     Sample  $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$  with policy  $\pi_\theta$ 
6:     for  $t = 0 : T - 1$  do
7:       Compute sample advantage  $U_t = R_{t+1} + \gamma v_{\pi_\theta}(S_{t+1}) - v_{\pi_\theta}(S_t)$ 
8:       Compute sample gradient  $\hat{\nabla} \leftarrow \hat{\nabla} + U_t \nabla \log \pi_\theta(A_t | S_t) / B$ 
9:     end for
10:  end for
11:  Update policy network  $\theta \leftarrow \theta + \alpha \hat{\nabla}$ 
12: end while

```

To implement, we need to estimate $v_{\pi_\theta}(S_t)$ for all trajectories in *mini-batch*

Estimating Values: Monte-Carlo

Say we are looking into one trajectory τ

$$\tau : S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

We know how to use this trajectory to compute value estimates: for each t

$$\hat{V}_t = \text{estimate of value for } S_t = G_t = \sum_{i=t}^T \gamma^i R_{i+1}$$

If we the same state happens multiple times in the trajectory: we count the number of times $S_t = S$ appears in the trajectory and average estimates, i.e.,

$$\hat{v}_{\pi_{\theta}}(S) = \frac{1}{\mathcal{N}(S \in \tau)} \sum_{t=0}^{T-1} \mathbf{1}\{S_t = S\} \hat{V}_t$$

where $\mathcal{N}(S \in \tau)$ is the number of times S has appeared in τ

Estimating Values: Monte-Carlo

If we have a *mini-batch* \mathbb{B} of trajectories

$$\tau : S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

We use the same approach

$$\hat{V}_t[\tau] = G_t[\tau] = \sum_{i=t}^T \gamma^i R_{i+1}[\tau]$$

We count the number of times $S_t = S$ appears in all trajectories and average the sample estimates, i.e.,

$$\hat{v}_{\pi_{\theta}}(S) = \frac{1}{\mathcal{N}(S \in \mathbb{B})} \sum_{\tau \in \mathbb{B}} \sum_{t=0}^{T-1} \mathbf{1}\{S_t[\tau] = S\} \hat{V}_t[\tau]$$

where $\mathcal{N}(S \in \mathbb{B})$ the number of times S has appeared in \mathbb{B}

Advantage PGM: With Value Estimates

EstAdvantagePGM() :

```

1: Initiate with  $\theta$  and learning rate  $\alpha$ 
2: while interacting do
3:   for mini-batch  $b = 1 : B$  do
4:     Sample  $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$  with policy  $\pi_\theta$ 
5:   end for
6:   Estimate value of all observed states in the mini-batch as  $\hat{v}_{\pi_\theta}(S_t)$ 
7:   Set  $\hat{\nabla} = \mathbf{0}$ 
8:   for  $b = 1 : B$  do
9:     for  $t = 0 : T - 1$  do
10:      Compute sample advantages  $U_t = R_{t+1} + \gamma \hat{v}_{\pi_\theta}(S_{t+1}) - \hat{v}_{\pi_\theta}(S_t)$ 
11:      Update sample gradient  $\hat{\nabla} \leftarrow \hat{\nabla} + U_t \nabla \log \pi_\theta(A_t | S_t) / B$ 
12:    end for
13:  end for
14:  Update policy network  $\theta \leftarrow \theta + \alpha \hat{\nabla}$ 
15: end while

```

Advantage PGM: With Value Estimates

We could guess that this algorithm is **not** going to perform very **impressive!**

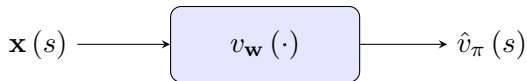
- + And why is that?!
- For the exact same reasons we said at the beginning of **Chapter 4**
 - ↳ We have **lots of states**
 - ↳ Many of them are **rarely** observed in a **small mini-batch**
 - ↳ The estimates can hence be very **high variance**
 - ↳ Also we need to wait for the **whole** mini-batch to be ready
 - ↳ ...
- + So, what is the solution?
- You tell me!
- + We go for function approximation via **value networks!**
- You got it right!

Recall: Value Network

Let's keep our trajectories here

$$\tau : S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

What we need is a simple *v-network*, as we only need the *state* values



In Chapter 4, we saw that we could train it via sample returns, i.e.,

$$\text{Dataset} = \left\{ (S_t[\tau], \hat{V}_t[\tau]) : \forall t \text{ and } \tau \right\}$$

and we train the network by minimizing the least-square loss

Value Network: Training

Let's keep our trajectories here

$$\tau : S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

This means that we compute the loss function

$$\mathcal{L}^v(\mathbf{w}) = \sum_{\tau} \sum_t \left(v_{\mathbf{w}}(S_t[\tau]) - \hat{V}_t[\tau] \right)^2$$

and update the weights of the *v-network* as

$$\mathbf{w} \leftarrow \underset{\mathbf{w}}{\operatorname{argmin}} \mathcal{L}^v(\mathbf{w})$$

which we approximately solve using gradient descent

Basic Actor-Critic

This is going to end us with a basic *actor-critic* algorithm:

AC_v1():

```

1: Initiate with  $\theta$  and  $\mathbf{w}$ , as well as a learning rate  $\alpha$ 
2: while interacting do
3:   for mini-batch  $b = 1 : B$  do
4:     Sample  $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$  with policy  $\pi_\theta$ 
5:     for  $t = 0 : T - 1$  do
6:       Compute value estimate  $\hat{V}_t$ 
7:       Compute sample advantages  $U_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_{t+1})$ 
8:       Update sample gradient  $\hat{\nabla} \leftarrow \hat{\nabla} + U_t \nabla \log \pi_\theta(A_t | S_t) / B$ 
9:     end for
10:  end for
11:  Update policy network  $\theta \leftarrow \theta + \alpha \hat{\nabla}$ 
12:  Update  $\mathbf{w}$  by SGD using value estimates  $\hat{V}_t$ 
13: end while

```

Training Value Network: TD Estimates

$$\tau : S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

But now that we have a *value network*, we could also use TD: at step t , we set

$$\hat{V}_t = \text{estimate of value for } S_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1})$$

- We can estimate the *advantage* using the current *value network*

$$U_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t)$$

- We then use least-squares update *value network* by TD sample estimates

$$\mathcal{L}^v(\mathbf{w}) = \sum_{\tau} \sum_{t=0}^{T-1} \left(v_{\mathbf{w}}(S_t[\tau]) - \hat{V}_t[\tau] \right)^2$$

Training Value Network: TD Estimates

Let's write the update rule of the **value network**: we compute the gradient of loss and move in that direction

$$\nabla \mathcal{L}^v(\mathbf{w}) = 2 \sum_{t=0}^{T-1} \underbrace{\left(v_{\mathbf{w}}(S_t) - \hat{V}_t \right)}_{-\Delta_t} \nabla v_{\mathbf{w}}(S_t)$$

We used to call Δ_t the TD **error**, and set the learning rate to some $\beta/2$ to get

$$\mathbf{w} \leftarrow \mathbf{w} + \beta \sum_{t=0}^{T-1} \Delta_t \nabla v_{\mathbf{w}}(S_t)$$

Training Value Network: TD Estimates

Let's look at **TD error**: recall that our **labels**, i.e., sample estimates of values, are

$$\hat{V}_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1})$$

So the **TD error** is given by

$$\begin{aligned}\Delta_t &= \hat{V}_t - v_{\mathbf{w}}(S_t) \\ &= R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t) \\ &= U_t\end{aligned}$$

This leads us to what we **observed** in Chapter 5

Recall: Advantage vs TD Error

Advantage is an estimator of **TD error**

A2C: Basic Version

A2C():

- 1: Initiate with θ and \mathbf{w} , as well as learning rates α and β
- 2: **while** interacting **do**
- 3: Start with zero gradients $\hat{\nabla}_{\mathbf{w}} = \hat{\nabla}_{\theta} = \mathbf{0}$
- 4: Sample $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$ with policy π_{θ}
- 5: Compute sample advantages $U_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t)$
- 6: **for** $t = 0 : T - 1$ **do**
- 7: Compute sample policy gradient $\hat{\nabla}_{\theta} \leftarrow \hat{\nabla}_{\theta} + U_t \nabla \log \pi_{\theta}(A_t | S_t)$
- 8: Compute sample value gradient $\hat{\nabla}_{\mathbf{w}} \leftarrow \hat{\nabla}_{\mathbf{w}} + U_t \nabla v_{\mathbf{w}}(S_t)$
- 9: **end for**
- 10: Update policy network $\theta \leftarrow \theta + \alpha \hat{\nabla}_{\theta}$
- 11: Update value network $\mathbf{w} \leftarrow \mathbf{w} + \beta \hat{\nabla}_{\mathbf{w}}$
- 12: **end while**

This is the *single-trajectory* form of

Advantage Actor Critic \equiv A2C

A2C: Online Version

Since we use TD, we can also update **online**, i.e., in each time step

OnlineA2C():

- 1: Initiate with θ and \mathbf{w} , a random state S_0 , $t = 0$ and learning rates α and β
- 2: **while** **interacting** **do**
- 3: Sample A_t from $\pi_\theta(\cdot|S_t)$
- 4: Sample single step $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$ from environment
- 5: Compute **sample advantage** $U_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t)$
- 6: Update policy network $\theta \leftarrow \theta + \alpha U_t \nabla \log \pi_\theta(A_t|S_t)$
- 7: Update value network $\mathbf{w} \leftarrow \mathbf{w} + \beta U_t \nabla v_{\mathbf{w}}(S_t)$
- 8: **if** S_{t+1} is terminal **then** draw a random S_{t+1}
- 9: Set $t \leftarrow t + 1$
- 10: **end while**

But, that would be too noisy and hence quite **unstable**

A2C: Mini-Batch Version

We can further extend to **mini-batch** learning

```

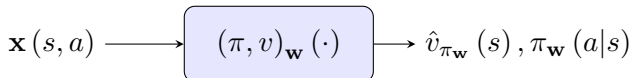
miniBatchA2C():
1: Initiate with  $\theta$  and  $\mathbf{w}$ , as well as learning rates  $\alpha$  and  $\beta$ 
2: while interacting do
3:   Start with zero gradients  $\hat{\nabla}_{\mathbf{w}} = \hat{\nabla}_{\theta} = \mathbf{0}$ 
4:   for mini-batch  $b = 1 : B$  do
5:     Sample  $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$  with policy  $\pi_{\theta}$ 
6:     Compute sample advantages  $U_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t)$ 
7:     for  $t = 0 : T - 1$  do
8:       Compute sample policy gradient  $\hat{\nabla}_{\theta} \leftarrow \hat{\nabla}_{\theta} + U_t \nabla \log \pi_{\theta}(A_t | S_t)$ 
9:       Compute sample value gradient  $\hat{\nabla}_{\mathbf{w}} \leftarrow \hat{\nabla}_{\mathbf{w}} + U_t \nabla v_{\mathbf{w}}(S_t)$ 
10:    end for
11:  end for
12:  Update policy network  $\theta \leftarrow \theta + \alpha \hat{\nabla}_{\theta}$ 
13:  Update value network  $\mathbf{w} \leftarrow \mathbf{w} + \beta \hat{\nabla}_{\mathbf{w}}$ 
14: end while
  
```

Actor-Critic via Shared-Network

There is one extra **obvious** fact: the **policy** and **values** that we learn are **very much** mutually related!

- + So, why don't we learn them **together**?!
- Actually we can!

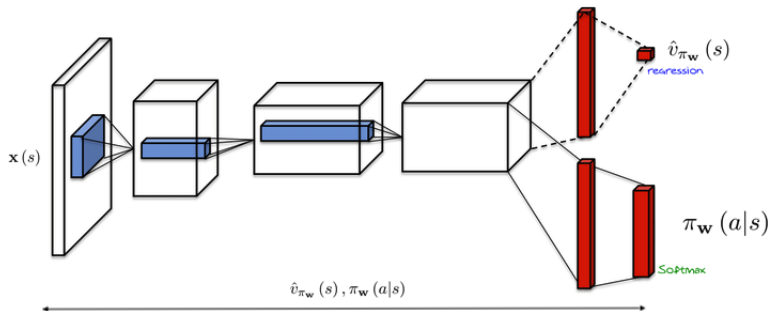
We can consider an **actor-critic** model, i.e.,



and train it all together!

- This model can be simply a DNN
- The DNN's output contains both **policy** and **value**

Actor-Critic via Shared-Network: Visualization



Here, *value* and *policy* share same layers except the few *last layers*

Actor-Critic via Shared-Network: Loss

- + But how can we train the loss in this network?
- We could let it to be proportional to **sum** of our both **objectives**

We could define a new loss as

$$\begin{aligned}\mathcal{L}(\mathbf{w}) &= -\mathcal{J}(\pi_{\mathbf{w}}) + \xi \mathcal{L}^v(\mathbf{w}) \\ &= \sum_{\tau} \sum_{t=0}^{T-1} -U_t[\tau] \log \pi_{\mathbf{w}}(A_t[\tau] | S_t[\tau]) + \xi \left(v_{\mathbf{w}}(S_t[\tau]) - \hat{V}_t[\tau] \right)^2\end{aligned}$$

for some hyperparameter ξ : it's easy to see that in this case

$$\nabla \mathcal{L}(\mathbf{w}) = - \sum_{\tau} \sum_{t=0}^{T-1} U_t[\tau] [\nabla \log \pi_{\mathbf{w}}(A_t[\tau] | S_t[\tau]) + \xi \nabla v_{\mathbf{w}}(S_t[\tau])]$$

A2C: Shared-Network Version

sharedNetA2C():

- 1: Initiate shared network $(\pi_{\mathbf{w}}, v_{\mathbf{w}})$ with \mathbf{w}
- 2: Choose potentially scheduled value-weight ξ and learning rate α
- 3: **while interacting do**
- 4: Start with zero gradients $\hat{\nabla}_{\mathbf{w}} = \mathbf{0}$
- 5: **for mini-batch $b = 1 : B$ do**
- 6: Sample $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$ with policy $\pi_{\mathbf{w}}$
- 7: Compute sample advantages $U_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t)$
- 8: **for $t = 0 : T - 1$ do**
- 9: Compute sample gradient $\hat{\nabla} \leftarrow \hat{\nabla} + U_t [\nabla \log \pi_{\mathbf{w}}(A_t | S_t) + \xi \nabla v_{\mathbf{w}}(S_t)]$
- 10: **end for**
- 11: **end for**
- 12: Update shared network $\mathbf{w} \leftarrow \mathbf{w} + \alpha \hat{\nabla}$
- 13: **end while**

Extension to Other PGMs

*In practice, all PGMs we studied in Chapter 5 can be implemented via the **actor-critic** idea through the following general framework*

Loop over the following three steps

- ① Specify **policy** and **value** networks
 - ↳ They could be either **separate** DNNs or DNNs with **shared layers**
- ② Set the policy and sample a batch of trajectories
- ③ Go over the batch for multiple epochs
 - ↳ After each **mini-batch** estimate the **policy** and **value** gradient
 - ↳ Update both **policy** and **value** networks after each **mini-batch**

Let's now look at the TRPO and PPO algorithms in **actor-critic** framework

TRPO: Actor-Critic

TRPO():

- 1: Initiate with θ and \mathbf{w} , as well as factor $\alpha < 1$ and learning rate β
- 2: **while** *interacting* **do**
- 3: Sample a *batch of trajectories* $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_T} S_T$ by policy π_θ
- 4: **for** multiple epochs **do**
- 5: **for** samples in each *mini-batch* **do**
- 6: Compute *sample advantage* using $v_{\mathbf{w}}(\cdot)$
- 7: Update policy *gradient* $\hat{\nabla}_\theta$ and *value gradient* $\hat{\nabla}_{\mathbf{w}}$
- 8: **end for**
- 9: Compute a *Hessian estimator* $\hat{\mathbf{H}}$ and solve $\hat{\mathbf{H}}\mathbf{y} = \hat{\nabla}$ for \mathbf{y}
- 10: *Backtrack on a line* to find minimum i satisfying $\bar{D}_{\text{KL}}(\pi_{\theta'} \parallel \pi_\theta) \leq d_{\max}$

$$\theta' \leftarrow \theta + \alpha^i \sqrt{\frac{2d_{\max}}{\mathbf{y}^\top \hat{\mathbf{H}} \mathbf{y}}} \mathbf{y}$$

- 11: Update $\theta \leftarrow \theta'$ and $\mathbf{w} \leftarrow \mathbf{w} + \beta \hat{\nabla}_{\mathbf{w}}$
- 12: **end for**
- 13: **end while**

Recall: Use Importance Sampling

Attention: Importance Sampling

It is important to *remember* that in each iteration of PGM

we estimate the *policy* gradient via *importance sampling*

Let's denote the policy of current mini-batch with π_{θ} : in next *mini-batch* we compute the *policy* gradient as


$$\hat{\nabla}_{\theta} \leftarrow \hat{\nabla}_{\theta} + \sum_{t=0}^T U_t \frac{\nabla \pi_{\mathbf{x}} (A_t | S_t) |_{\mathbf{x}=\theta}}{\pi_{\theta} (A_t | S_t)}$$

with U_t being the sample advantage of policy π_{θ}

Recall: Use Importance Sampling

- + You mentioned this before! What is new about this?!
- Well! We should also consider it in our value estimation!

Denote the policy that we sampled with in line 3 with $\pi_{\theta_{\text{old}}}$: after multiple **mini-batches** we have an updated **policy** gradient π_{θ}

 To update the value network in this mini-batch, we consider the **Bellman equation** which says

$$v_{\pi_{\theta}}(S_t) = \mathbb{E}_{\pi_{\theta}} \{R_{t+1} + \gamma v_{\pi_{\theta}}(S_{t+1})\}$$

and set the labels for **value network training** as

$$\hat{v}_{\pi_{\theta}}(S_t) = R_{t+1} + \gamma v_{\pi_{\theta}}(S_{t+1})$$

But this is only valid if we had sampled the trajectory by π_{θ} !

Recall: Use Importance Sampling

- + Shall we use *importance sampling* here as well?!
- Sure!

By *importance sampling* we could say

$$\begin{aligned} v_{\pi_{\theta}}(S_t) &= \mathbb{E}_{\pi_{\theta}} \{ R_{t+1} + \gamma v_{\pi_{\theta}}(S_{t+1}) \} \\ &= \mathbb{E}_{\pi_{\theta_{\text{old}}}} \left\{ (R_{t+1} + \gamma v_{\pi_{\theta}}(S_{t+1})) \frac{\pi_{\theta}(A_t | S_t)}{\pi_{\theta_{\text{old}}}(A_t | S_t)} \right\} \end{aligned}$$

and now we can compute value estimators from our sample trajectories as

$$\hat{v}_{\pi_{\theta}}(S_t) = \left(R_{t+1} + \gamma v_{\pi_{\theta_{\text{old}}}}(S_{t+1}) \right) \frac{\pi_{\theta}(A_t | S_t)}{\pi_{\theta_{\text{old}}}(A_t | S_t)}$$

Recall: Use Importance Sampling

This means that the advantage should be computed as

$$\begin{aligned}
 U_t &= \underbrace{\left(R_{t+1} + \gamma v_{\pi_{\theta_{\text{old}}}}(S_{t+1}) \right)}_{\text{samples in batch}} \underbrace{\frac{\pi_{\theta}(A_t|S_t)}{\pi_{\theta_{\text{old}}}(A_t|S_t)}}_{\text{importance sampling}} - v_{\mathbf{w}}(S_t) \\
 &\approx \left(R_{t+1} + \gamma v_{\pi_{\theta_{\text{old}}}}(S_{t+1}) - v_{\pi_{\theta_{\text{old}}}}(S_t) \right) \frac{\pi_{\theta}(A_t|S_t)}{\pi_{\theta_{\text{old}}}(A_t|S_t)} \\
 &= U_t^{\text{old}} \frac{\pi_{\theta}(A_t|S_t)}{\pi_{\theta_{\text{old}}}(A_t|S_t)}
 \end{aligned}$$

This is *the correct way* of estimating *advantage*: other implementations that ignore this will have *bias* especially with *too much epochs*

Recall: Use Importance Sampling

Consequently, the value gradient is updated as

$$\begin{aligned}\hat{\nabla}_{\mathbf{w}} &\leftarrow \hat{\nabla}_{\mathbf{w}} + \frac{1}{T} \sum_{t=0}^{T-1} U_t \nabla v_{\mathbf{w}}(S_t) \\ &\leftarrow \hat{\nabla}_{\mathbf{w}} + U_t^{\text{old}} \frac{\pi_{\boldsymbol{\theta}}(A_t | S_t)}{\pi_{\boldsymbol{\theta}_{\text{old}}}(A_t | S_t)} \nabla v_{\mathbf{w}}(S_t)\end{aligned}$$

at the end of each trajectory

- + Shall we also consider this when we *update the policy*?
- Of course!

Recall: Use Importance Sampling

In a given *mini-batch* we update the *policy* gradient as

$$\begin{aligned}\hat{\nabla}_{\theta} &\leftarrow \hat{\nabla}_{\theta} + \sum_{t=0}^T U_t \frac{\nabla \pi_{\mathbf{x}}(A_t | S_t) |_{\mathbf{x}=\theta}}{\pi_{\theta}(A_t | S_t)} \\ &\leftarrow \hat{\nabla}_{\theta} + \sum_{t=0}^T U_t^{\text{old}} \frac{\pi_{\theta}(A_t | S_t)}{\pi_{\theta_{\text{old}}}(A_t | S_t)} \frac{\nabla \pi_{\mathbf{x}}(A_t | S_t) |_{\mathbf{x}=\theta}}{\pi_{\theta}(A_t | S_t)} \\ &\leftarrow \hat{\nabla}_{\theta} + \sum_{t=0}^T U_t^{\text{old}} \frac{\nabla \pi_{\mathbf{x}}(A_t | S_t) |_{\mathbf{x}=\theta}}{\pi_{\theta_{\text{old}}}(A_t | S_t)}\end{aligned}$$

which is now consistent with *importance sampling*

TRPO: Actor-Critic

TRPO_AC() :

- 1: Initiate with $\theta = \theta_{\text{old}}$ and \mathbf{w} , as well as factor $\alpha < 1$ and learning rate β
- 2: **while** *interacting* **do**
- 3: Sample a *batch of trajectories* $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_T} S_T$ by policy $\pi_{\theta_{\text{old}}}$
- 4: Compute *sample advantage* U_t^{old} using $v_{\mathbf{w}}(\cdot)$
- 5: **for** multiple epochs in *each mini-batch* **do**
- 6: Compute gradient estimators $(\hat{\nabla}_{\theta}, \hat{\nabla}_{\mathbf{w}})$ from U_t^{old} via *importance sampling*
- 7: Compute a *Hessian estimator* $\hat{\mathbf{H}}$ and solve $\hat{\mathbf{H}}\mathbf{y} = \hat{\nabla}_{\theta}$ for \mathbf{y}
- 8: *Backtrack on a line* to find minimum i satisfying $\bar{D}_{\text{KL}}(\pi_{\theta'} \parallel \pi_{\theta}) \leq d_{\text{max}}$

$$\theta' \leftarrow \theta + \alpha^i \sqrt{\frac{2d_{\text{max}}}{\mathbf{y}^T \hat{\mathbf{H}} \mathbf{y}}} \mathbf{y}$$

- 9: Update $\theta \leftarrow \theta'$ and $\mathbf{w} \leftarrow \mathbf{w} + \beta \hat{\nabla}_{\mathbf{w}}$
- 10: **end for**
- 11: Update $\pi_{\theta} \leftarrow \pi_{\theta_{\text{old}}}$
- 12: **end while**

PPO: Actor-Critic

In PPO, we maximize in each iteration the *restricted surrogate*

$$\tilde{\mathcal{L}}(\pi_{\mathbf{x}}) = \mathbb{E}_{\pi_{\theta}} \left\{ \min \left\{ U_t \frac{\pi_{\mathbf{x}}(A_t|S_t)}{\pi_{\theta}(A_t|S_t)}, \ell_{\varepsilon}(U_t) \right\} \right\}$$

Similar to TRPO, we can estimate *restricted surrogate* via *importance sampling*

$$\begin{aligned} \tilde{\mathcal{L}}(\pi_{\mathbf{x}}) &= \text{mean} \left[\sum_t \min \left\{ U_t \frac{\pi_{\mathbf{x}}(A_t|S_t)}{\pi_{\theta}(A_t|S_t)}, \ell_{\varepsilon}(U_t) \right\} \right] \\ &= \text{mean} \left[\sum_t \min \left\{ U_t^{\text{old}} \frac{\pi_{\mathbf{x}}(A_t|S_t)}{\pi_{\theta_{\text{old}}}(A_t|S_t)}, \ell_{\varepsilon}(U_t) \right\} \right] \end{aligned}$$

where the *mean* is computed over the sample trajectories of a *mini-batch*

PPO Algorithm: Actor-Critic

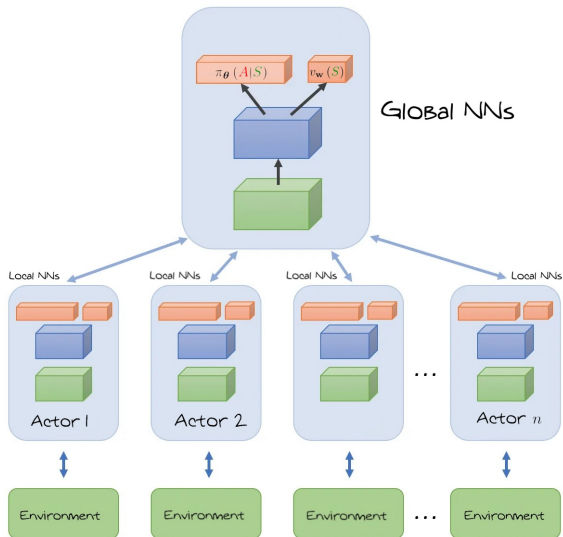
PPO_AC() :

- 1: Initiate with $\theta = \theta_{\text{old}}$ and \mathbf{w} , as well as learning rates α and β
- 2: **while** *interacting* **do**
- 3: Sample a *batch of trajectories* $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_T} S_T$ by policy $\pi_{\theta_{\text{old}}}$
- 4: Compute *sample advantage* U_t^{old} using $v_{\mathbf{w}}(\cdot)$
- 5: **for** multiple epochs in *each mini-batch* **do**
- 6: Compute *value* gradient estimator $\hat{\nabla}_{\mathbf{w}}$ from U_t^{old} via *importance sampling*
- 7: Compute the *restricted surrogate*

$$\tilde{\mathcal{L}}(\pi_{\mathbf{x}}) = \text{mean} \left[\sum_t \min \left\{ U_t^{\text{old}} \frac{\pi_{\mathbf{x}}(A_t|S_t)}{\pi_{\theta_{\text{old}}}(A_t|S_t)}, \ell_{\varepsilon}(U_t) \right\} \right]$$

- 8: Update $\theta \leftarrow \theta + \alpha \nabla \tilde{\mathcal{L}}(\pi_{\mathbf{x}})|_{\mathbf{x}=\theta}$ and $\mathbf{w} \leftarrow \mathbf{w} + \beta \hat{\nabla}_{\mathbf{w}}$
- 9: **end for**
- 10: Update $\pi_{\theta} \leftarrow \pi_{\theta_{\text{old}}}$
- 11: **end while**

Distributed Actor-Critic



Distributed Setting: Asynchronous vs Synchronous

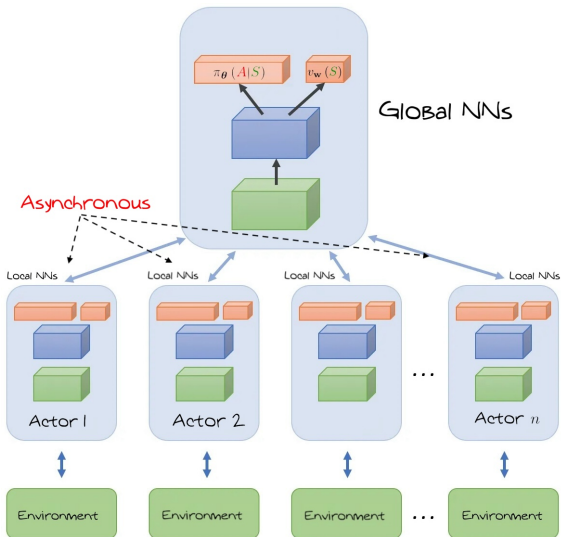
We can implement **actor-critic** approaches in a **distributed** fashion

- Multiple actors few samples with local policy and value network
 - ↳ Each actor focuses on a **specific part of environment**
- They share their gradient estimators with the **server**
- **Server** treats the collected estimators as of a large **mini-batch**
 - ↳ It updates its networks on this large **mini-batch**
- All actors update their local networks every time **server** shares its networks

We can implement this setting

- **Synchronous**
 - ↳ This is basically the **same** as what we do in A2C, TRPO or PPO
- **Asynchronous**
 - ↳ **Server** does **not** wait for all actors to send their estimators
 - ↳ It uses **what it has** every couple of rounds and remaining in next rounds

A3C: Asynchronous A2C



Some Final Remarks

It turns out that *asynchronous* update can negatively impact *convergence*

- A3C is hence not really extended to other PGMs
- In practice we usually implement *actor-critic* approaches in *synchronous distributed* form

-
- + Is that it? Are we free to go now?!
 - Pretty much Yes! Just we may further take a look at *deterministic policy gradient* approaches as well
 - + Is it a new set of approaches?!
 - **No!** It's a specific form of *actor-critic* methods that are *better compatible* with *continuous actions*

Learning Deterministic Policy

From model-based RL we know that: the *optimal* policy can be *deterministic*

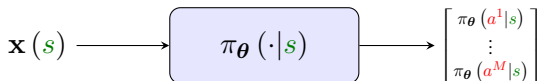
Why don't we train a *policy network* that learns a *deterministic* policy?

-
- + You are contradicting yourself! You said that *stochastic* policy is a general case that includes *deterministic* policies as well! Now you want to get back to a *deterministic* policy?!
 - Well! You're *right*! But there will be no harm in learning a *deterministic* policy! It might only be *less effective*!
 - + Why we should do it then?
 - It could give us some *benefits*, especially when we have *continuous action-space*

Learning Deterministic Policy

With *continuous action-space*, policy is a density *function*

- With *discrete action-space*, we can show policy by a finite vector¹



- Unlike *discrete action-space*, we *cannot* do this with *continuous actions*
 ↳ We should learn a *function* from *state feature*, e.g.,

$$\pi_{\theta}(a | s) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{(a - \text{DNN}(\mathbf{x}(s) | \theta))^2}{2\sigma^2} \right\}$$

- ↳ We then sample from this learned *density function*, e.g.,
 ↳ Draw a sample from Gaussian distribution with *mean* $\text{DNN}(\mathbf{x}(s) | \theta)$ and variance σ^2

¹We have seen this in Assignment 3

Learning Deterministic Policy

There are several things that could go wrong

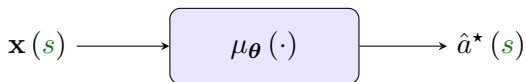
- What if the generated sample is out of accepted range?
 - ↳ Our sample from Gaussian distribution is **extremely large**
 - ↳ In **sensitive control settings**, this could harm the system
- What if we only try a few samples?
 - ↳ We don't see the probabilities as opposed to **continuous actions**
 - ↳ We then **cannot** really **reject** too many samples

With **continuous actions**, we usually prefer to learn a **deterministic policy**: its main feature is that it can be represented by a **single action a^***

$$\pi_{\theta}(a|s) = \begin{cases} 1 & a = a^* \\ 0 & a \neq a^* \end{cases}$$

Deterministic Policy Network

Considering a **deterministic** policy, we **only** need to learn an estimate of **optimal action** in each state: we can revise our policy network into a **deterministic policy network**



Deterministic Policy Network

Deterministic policy network maps a **state**-features into a **single action** and can be realized by a DNN with input being the state feature representation and a **single output**

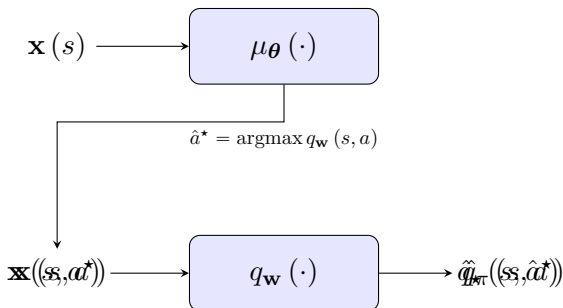
Deterministic Policy Network: *Training*

- + How should we train *these networks*? Similar to other *policy networks*?
- We can look at them as a *special form of policy networks* and do the same thing, *yes!* *However*, it turns out that in this special case, *there are better ways to do it!* Especially as we always implement them *actor-critic*
- + So, we should start all over *again?!*
- Not really! We should basically use what we learned for *DQL*

Deterministic Policy Network: *Training*

Recall the property of *optimal policy*: it gives *maximum value and action-value*

If we have a Q-network that estimates *optimal action-value*, we can say



The output satisfies

$$\hat{q}_{\star}(s, \hat{a}^{\star}) = \max_a \hat{q}_{\star}(s, a)$$

Deterministic Policy Network: Training

So, in *actor-critic* form with a Q-network, we could train the *deterministic policy network* as

$$\theta^* = \operatorname{argmax}_{\theta} Q_{\mathbf{w}}(s, \mu_{\theta}(s))$$

which we can solve using *gradient ascent* by updating as

$$\begin{aligned} \theta &\leftarrow \theta + \alpha \nabla_{\theta} Q_{\mathbf{w}}(s, \mu_{\theta}(s)) \\ &\leftarrow \theta + \underbrace{\alpha \frac{\partial}{\partial a} Q_{\mathbf{w}}(s, a) \big|_{a=\mu_{\theta}(s)}}_{\text{backprop over Q-Net}} \underbrace{\nabla \mu_{\theta}(s)}_{\text{backprop over policy}} \end{aligned}$$

Moral of Story

As long as we have an estimator of *optimal action-value function*, we can train *deterministic policy network* very easily!

Deterministic Policy Gradient

- + But how can we can find such an estimator?
- Well! We have done this *before!*
- + You mean in *DQL?!?*
- *Exactly!* In *Q-learning* we use *Bellman's optimality equation* to estimate *optimal action-value* function: we can do the same here

Recall that *Bellman's optimality equation* indicate that

$$q_{\star}(S_t, A_t) = R_{t+1} + \gamma \mathbb{E}_{S_{t+1} \sim p(\cdot | S_t, A_t)} \left\{ \max_a q_{\star}(S_{t+1}, a) \right\}$$

and if we know the action $a^{\star} = \operatorname{argmax}_a q_{\star}(S_{t+1}, a)$, we could write

$$q_{\star}(S_t, A_t) = R_{t+1} + \gamma \mathbb{E}_{S_{t+1} \sim p(\cdot | S_t, A_t)} \{ q_{\star}(S_{t+1}, a^{\star}) \}$$

Deterministic Policy Gradient

If we use our *deterministic* policy network in one time step we sample

$$S_t, \mu_{\theta}(S_{t+1}) \xrightarrow{R_{t+1}} S_{t+1}$$

We can then sample an estimator of *optimal action-value* at $a = \mu_{\theta}(S_{t+1})$

$$\hat{Q}_t = R_{t+1} + \gamma Q_{\mathbf{w}}(S_{t+1}, \mu_{\theta}(S_{t+1}))$$

Once we are over with *sample trajectory*: we update *Q-network* to minimize loss

$$\mathcal{L}(\mathbf{w}) = \frac{1}{T} \sum_{t=0}^{T-1} \left(Q_{\mathbf{w}}(S_t, \mu_{\theta}(S_t)) - \hat{Q}_t \right)^2$$

And the life is *much easier* as compared to *TRPO and PPO* 😊

Deterministic Policy Gradient

We do very well know how to do this

$$\begin{aligned}
 \mathbf{w} &\leftarrow \mathbf{w} - \beta \nabla \mathcal{L}(\mathbf{w}) \\
 &\leftarrow \mathbf{w} + \beta \text{mean} \left[\left(\hat{Q}_t - Q_{\mathbf{w}}(S_t, \mu_{\boldsymbol{\theta}}(S_t)) \right) \nabla Q_{\mathbf{w}}(S_t, \mu_{\boldsymbol{\theta}}(S_t)) \right] \\
 &\leftarrow \mathbf{w} + \beta \text{mean} [\Delta_t \nabla Q_{\mathbf{w}}(S_t, \mu_{\boldsymbol{\theta}}(S_t))]
 \end{aligned}$$

where $\Delta_t = \hat{Q}_t - Q_{\mathbf{w}}(S_t, \mu_{\boldsymbol{\theta}}(S_t))$ is the TD error

Alternating between the two update rules, we end up with a

Deterministic Policy Gradient \equiv DPG

algorithm: there are various DPG algorithms; we take a look into the famous one

A Basic DPG Algorithm

We can use these updates to write a simple *online* DPG algorithm

DPG_v1():

- 1: Initiate with θ and \mathbf{w} , as well as factor $\alpha < 1$ and learning rate β
- 2: Initiate some *initial state* S_0 and draw *action* A_0 as $A_0 \leftarrow \mu_{\theta}(S_0)$
- 3: **while** *interacting* **do**
- 4: Sample a *time step* $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$
- 5: Draw the next optimal action as $A_{t+1} \leftarrow \mu_{\theta}(S_{t+1})$
- 6: Compute $\Delta = R_{t+1} + \gamma Q_{\mathbf{w}}(S_{t+1}, A_{t+1}) - Q_{\mathbf{w}}(S_t, A_t)$
- 7: Update *value network* as $\mathbf{w} \leftarrow \mathbf{w} + \beta \Delta \nabla Q_{\mathbf{w}}(S_t, A_t)$
- 8: Update *policy network* as $\theta \leftarrow \theta + \alpha \frac{\partial}{\partial \mathbf{a}} Q_{\mathbf{w}}(S_t, \mathbf{a})|_{\mathbf{a}=A_t} \nabla \mu_{\theta}(S_t)$
- 9: Go for next state $S_t \leftarrow S_{t+1}$
- 10: **if** S_t is terminal **then**
- 11: Draw a new *random state* S_0 and $A_0 \leftarrow \mu_{\theta}(S_0)$
- 12: **end if**
- 13: **end while**

Basic DPG: Practical Challenges

Using our knowledge, we can easily detect **challenges** of this basic algorithm

- Lack of **exploration** → **ϵ -greedy improvement**
 - ↳ We follow **blindly** the **deterministic policy network**
 - ↳ We do not give any chance for **exploration**
 - ↳ This can quickly stick us to a **bad locally-optimal deterministic policy**
- **High-variance** gradient estimators → **experience replay**
 - ↳ It's **online** and hence update the networks with **single** time step **samples**
 - ↳ We need more samples to compute **better estimators**
 - ↳ We would like to have **independent samples**
- Variation of **training labels** → **target network**
 - ↳ Each time we update, we **change** the label in the **training batch**
 - ↳ This can severely deteriorate the **convergence of algorithm**

DPG: ϵ -Greedy Improvement

To have sufficient **exploration** of environment: we can follow **ϵ -greedy approach**

- + But how does it work here? You said we have **continuous actions**!
- Well! We can add **continuous randomness** to our policy

Say we get $A_t \leftarrow \mu_{\theta}(S_t)$ at time t : then we replace our action with

$$A_t \leftarrow A_t + \sqrt{\epsilon} Z_t$$

where Z_t is **random noise** with **mean zero** and **variance one**

Classical choice of Z_t is **zero-mean unit-variance Gaussian** variable, i.e.,

$$Z_t \sim \mathcal{N}(0, 1)$$

Note that the noise term $\sqrt{\epsilon} Z_t$ is then **zero-mean** with **variance ϵ**

DPG: ϵ -Greedy Improvement

- + But what if after adding $\sqrt{\epsilon}Z_t$, the **action** gets out of its **allowed range**?
For instance, we get $A_t = 5$ and $\sqrt{\epsilon}Z_t = 3$, but we should have **all actions between 2 and 6**
- That's a valid question! We usually **clip** the action in this case

To avoid **out-of-range actions**, we replace apply ϵ -greedy approach as

$$A_t \leftarrow \text{Clip}(\mu_{\theta}(S_t) + \sqrt{\epsilon}Z_t, a_{\min}, a_{\max})$$

where a_{\min} and a_{\max} are minimum and maximum **allowed actions** and

$$\text{Clip}(x, a_{\min}, a_{\max}) = \begin{cases} a_{\min} & x < a_{\min} \\ x & a_{\min} \leq x \leq a_{\max} \\ a_{\max} & x > a_{\max} \end{cases}$$

DPG: Replay Buffer

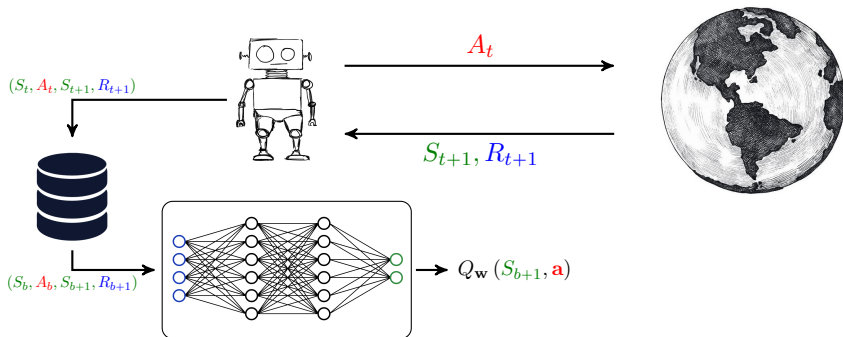
To reduce estimator's variance and enhance sample-efficiency, we can use *experience replay* as in *DQL*

- + Wait a moment! But we talked about the fact that *experience replay* can increase *estimator's variance* in PGMs due to *importance sampling* argument! Now, we just ignore all those discussions?!
- With *stochastic* policy yes! But here we have a *deterministic* policy

Deterministic policy returns *only one action*

- For each choice of θ , our policy chooses *only one action*
 - ↳ If we change θ we only change this action
- Policy update *does not change* the probability of *all actions*
 - ↳ This is in fact why *Q-learning* does not suffer from *high estimate variance*

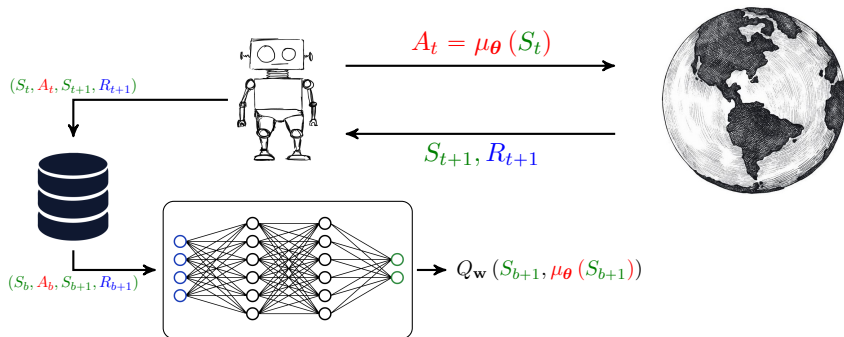
Recall: *Experience Replay in DQL*



We update DQN by randomly sampled mini-batches using **TD error**

$$\Delta_b \leftarrow R_{b+1} + \gamma \max_m Q_w(S_{b+1}, \mathbf{a}^m) - Q_w(S_b, \mathbf{a}_b)$$

DPG: Experience Replay



We now use the *deterministic* policy network to compute *TD error*

$$\Delta_b \leftarrow R_{b+1} + \gamma Q_w(S_{b+1}, \mu_{\theta}(S_{b+1})) - Q_w(S_b, A_b)$$

DPG: Target Network

The last thing to handle is to keep training dataset fixed for a while

- After each **mini-batch**, we change both **policy** and **Q-network**
 - ↳ We update \mathbf{w} and θ
- If we use the same networks to compute **the estimate**

$$\hat{Q}_b = R_{b+1} + \gamma Q_{\mathbf{w}}(S_{b+1}, \mu_{\theta}(S_{b+1}))$$

then our next iteration runs over a **different** dataset

↳ This can cause our training loop to **diverge**

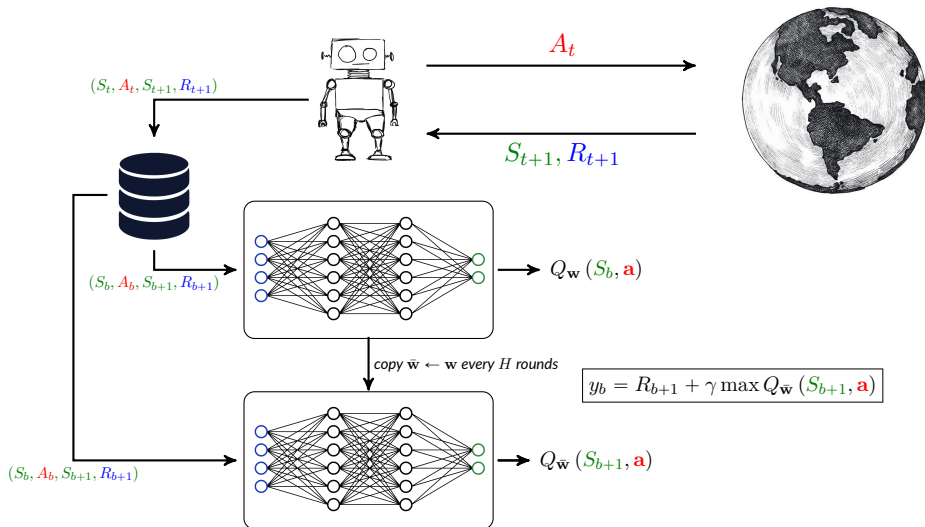
We have dealt with this in **DQL** using **target network**

Target Network in DPG

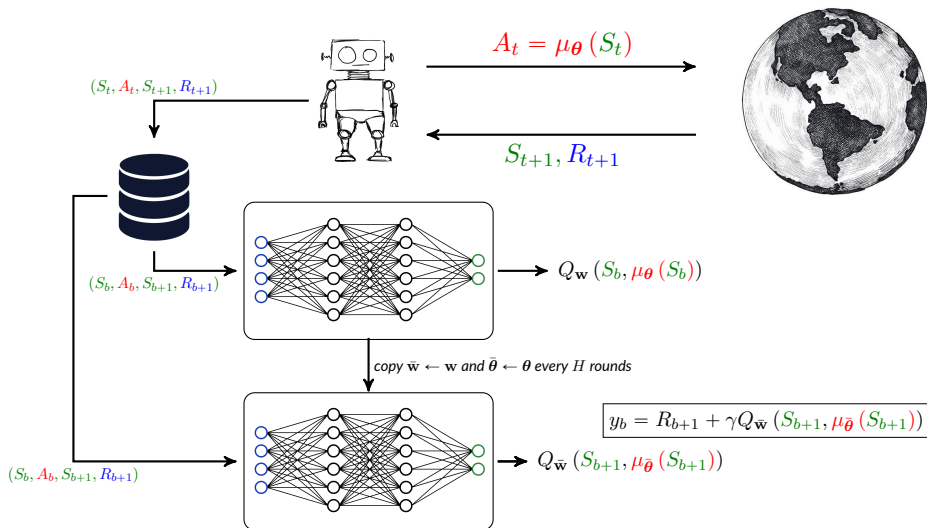
Copy **DQN** and **policy** network into the exactly same target networks

- Use these target networks to compute the **estimates**
- Update them every multiple iterations by new copies of **online** networks

Recall: Target Network in DQL



DPG: Target Networks



DDPG: Deep DPG Algorithm

DDPG() :

- 1: Initiate with $\theta = \bar{\theta}$ and $\mathbf{w} = \bar{\mathbf{w}}$, as well as factor $\alpha < 1$ and learning rate β
- 2: Initiate *state* S_0 and draw $A_0 \leftarrow \text{Clip}(\mu_{\theta}(S_0) + \sqrt{\epsilon}\mathcal{N}(0, 1))$
- 3: **while** *interacting* **do**
- 4: Sample a *time step* $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$ and save in *replay buffer*
- 5: **for** multiple iterations **do**
- 6: Sample $S_b, A_b \xrightarrow{R_{b+1}} S_{b+1}$ from *replay buffer*
- 7: Draw $A_{b+1} \leftarrow \text{Clip}(\mu_{\bar{\theta}}(S_{b+1}) + \sqrt{\epsilon}\mathcal{N}(0, 1))$
- 8: Compute $\Delta = R_{b+1} + \gamma Q_{\bar{\mathbf{w}}}(S_{b+1}, A_{b+1}) - Q_{\mathbf{w}}(S_b, A_b)$
- 9: Update *value network* as $\mathbf{w} \leftarrow \mathbf{w} + \beta \Delta \nabla Q_{\mathbf{w}}(S_t, A_t)$
- 10: Update *policy network* as $\theta \leftarrow \theta + \alpha \frac{\partial}{\partial \mathbf{a}} Q_{\mathbf{w}}(S_t, \mathbf{a})|_{\mathbf{a}=A_t} \nabla \mu_{\theta}(S_t)$
- 11: **if** H iterations passed **then**
- 12: Copy $\bar{\theta} \leftarrow \theta$ and $\bar{\mathbf{w}} \leftarrow \mathbf{w}$
- 13: **end if**
- 14: **end for**
- 15: **end while**

DDPG: Overestimate Issue

DDPG has been the key DPG algorithm and widely used for

- Dealing with continuous action spaces
- Enabling **off-policy** learning with a **deterministic** version of PGM

DDPG has shown a key issue; namely, **overestimate** of **values**

Value Overestimate of DDPG

Action-values estimated by DDPG can have **significant biases**

- + We had it also in **DQL** and you said “it’s not a big deal in general”! So, do we care about it here?!
- Well! Here is more important! Because, we use those estimates to update the policy and large biases would explode **TD error**!
- + So, shall we do **double DQL** then?!
- Pretty much yes!

TD3: Twin Delayed DDPG

Value overestimate has been addressed in the extended version of DDPG

Twin Delayed DDPG \equiv *TD3*

In TD3, we add three extra tricks to DDPG

- 1 We use *double DQN* to *suppress* undesired *bias*
 - ↳ Remember that *bias* was mainly coming out of *max operator*
 - ↳ We came with a *remedy* called *double Q-learning*
- 2 We *delay* the policy update, i.e., update the policy *less frequent*
 - ↳ We update *DQN* after each mini-batch, but
 - ↳ We update *policy network* *once* every *couple of mini-batches*
- 3 We add *extra noise* to actions when we use them in *target networks*
 - ↳ This way we make the *estimation error* somehow *independent*
 - ↳ This leads to *less bias*

From DPG to PGM

Both ideas of learning *deterministic* or *stochastic* policy have pros and cons

- For *deterministic* policy we could say
 - ↳ We can efficiently use *off-policy* learning
 - ↳ We can get better sample efficiency
 - ↳ *But* we get less chance of being optimal
 - ↳ We do not search among possible random optimal policies
 - For *stochastic* policy we could say
 - ↳ We search among much larger set of policies
 - ↳ We can converge to a better policy
 - ↳ *But* we have troubles with sample efficiency
 - ↳ We cannot easily learn *off-policy* due to limits of *importance sampling*
- + Is there any way to get good things of both worlds?
- Soft actor-critic approaches actually do this

Recall: Information Content and Entropy

To understand the idea behind soft actor-critic, let's recap some definitions

Information Content

The information content of random variable $X \sim p(x)$ is

$$i(X) = \log \frac{1}{p(X)}$$

The information contents have some interesting properties

- It's always **non-negative**, since $0 \leq p(x) \leq 1$
- The **less likely** outcome $X = x$ is, the **more** will be its **information content**
 - ↳ Think about it! You will find it very intuitive

Recall: Information Content and Entropy

Entropy

For random variable $X \sim p(x)$, entropy is its average **information content**, i.e.,

$$H_p(X) = \mathbb{E}_p\{i(X)\} = \mathbb{E}_p\left\{\log \frac{1}{p(X)}\right\} = \int_x p(x) \log \frac{1}{p(x)}$$

Entropy quantifies how much **confusion** we have about X

- If X is **highly random**, e.g., uniformly or Gaussian distributed,
↳ Then $H_p(X)$ is very large
- If X is **deterministic**
↳ Then $H_p(X) = 0$

Redefining Value Function

After dealing with both *deterministic* and *stochastic* policies we might formulate the best policy as follows

- It's globally *deterministic*
 - ↳ If one action gives better reward, it should go for it
- It's locally *stochastic*
 - ↳ Among actions with same rewards, it chooses one at random

We could capture both these behaviors by looking into a new metric

Say we play with policy π : at time t , we are interested in

$$\tilde{R}_{t+1} = R_{t+1} + \xi H_{\pi}(A_t | S_t)$$

for some ξ , where $H_{\pi}(A_t | S_t)$ is entropy of action $A_t \sim \pi(\cdot | S_t)$

Redefining Value Function

Say we play with policy π : at time t , we are interested in

$$\tilde{R}_{t+1} = R_{t+1} + \xi H_{\pi}(A_t | S_t)$$

for some ξ , where $H_{\pi}(A_t | S_t)$ is entropy of action $A_t \sim \pi(\cdot | S_t)$

This new **modified reward** incorporates both desires

- Being globally **deterministic**

↳ For actions with larger R_{t+1} , the modified \tilde{R}_{t+1} is also larger

- Being locally **stochastic**

↳ For actions with same R_{t+1} , policy with higher randomness has larger \tilde{R}_{t+1}

Well! This might be a **better reward!**

SAC: Soft Actor-Critic

We can use either DPG or PGM to develop an actor-critic method for this new reward, i.e., we could define

$$\begin{aligned} v_{\pi}(s) &= \mathbb{E}_{\pi} \left\{ \sum_{i=0}^{\infty} \gamma^i \tilde{R}_{t+i+1} | S_t = s \right\} \\ &= \mathbb{E}_{\pi} \left\{ \sum_{i=0}^{\infty} \gamma^i [R_{t+i+1} + \xi H_{\pi}(A_{t+i} | S_{t+i})] | S_t = s \right\} \end{aligned}$$

Interestingly, we end up in both cases with the *same* policy and value *gradients*!

The derived actor critic method is referred to as

Soft Actor-Critic \equiv SAC

DRL Algorithms

Most DRL algorithms used in practice are *actor-critic*

We already discussed all main classes of *actor-critic* approaches

To each class, there are various extensions

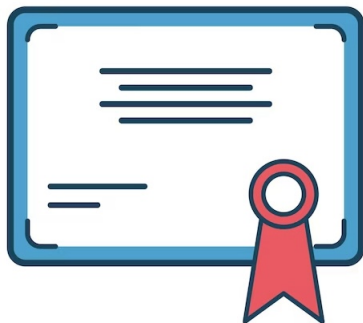
- You are able now to *follow* all those extensions
- If necessary, you could *come up with your* own particular extension!

A rule-of-thumb is

- If you deal with *discrete* actions and have no concern on *sample efficiency*
 - ↳ Use *stochastic-policy actor-critic* approaches
- If you deal with *continuous* actions and/or need *sample efficiency*
 - ↳ Use DPG-like *actor-critic* approaches

OpenAI: Spinning Up in DRL

Congratulations! You are now *Deep RL experts!*



Looking for some mini-projects for further practices? Take a look at [OpenAI page](#)